Artificial Intelligence II

Part 2: Lecture 11

Yalda Mohsenzadeh

Winter 2024

Slides are adapted from Antonio Torralba and Philip Isola (MIT)

Image Synthesis

- Image synthesis
  - Generative Adversarial Networks
- Structured prediction
  - Image-to-image GANs
- Domain mapping

Image classification

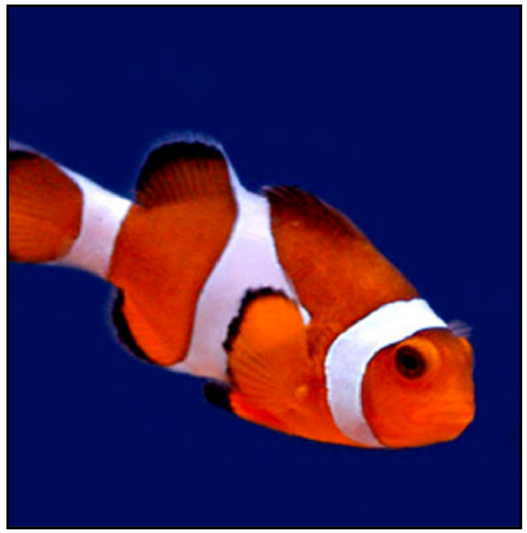

Classifier

“Fish”

image **x**

label y

Image classification

Classifier

“Fish”

image **x**

label y

Image classification

Classifier

“Fish”

image **x**

label y

Image classification

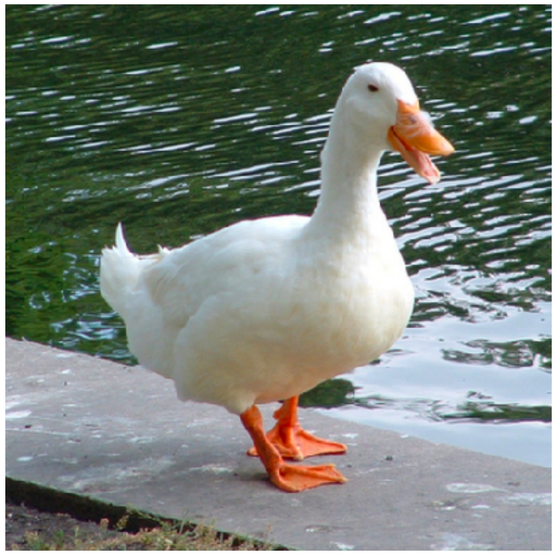

image **x**

Classifier

"Duck"

label y

Image synthesis

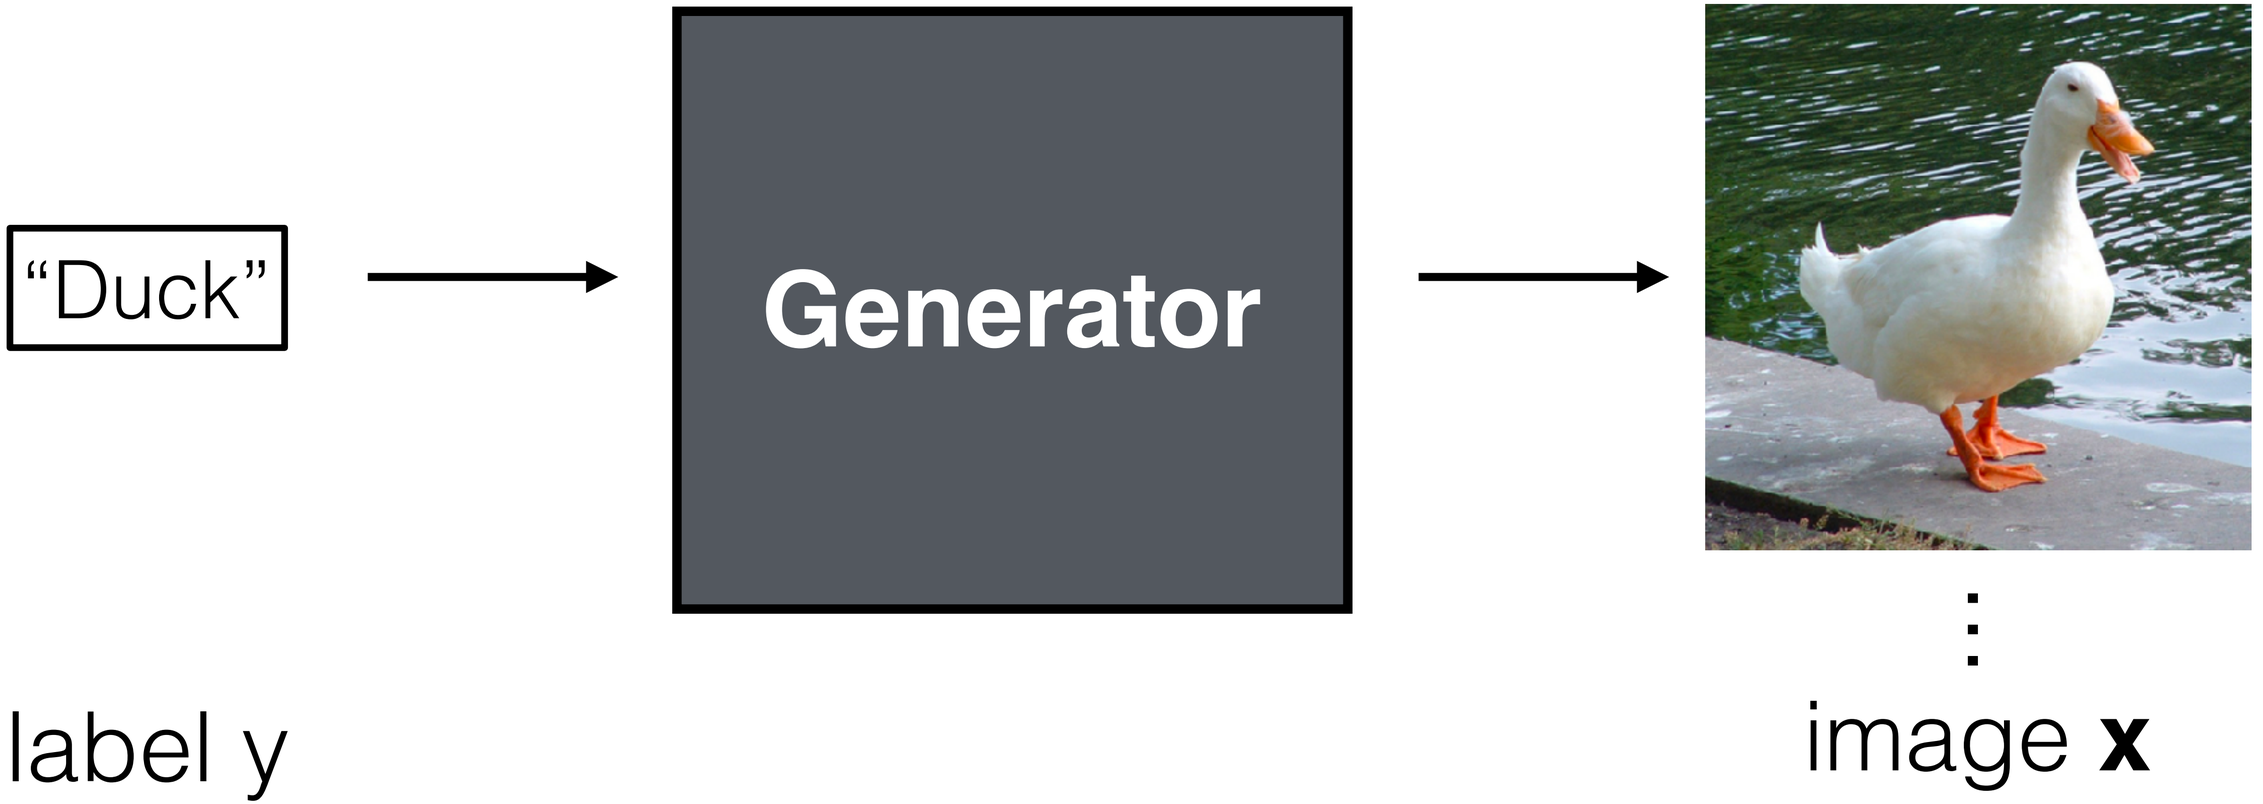

Image synthesis

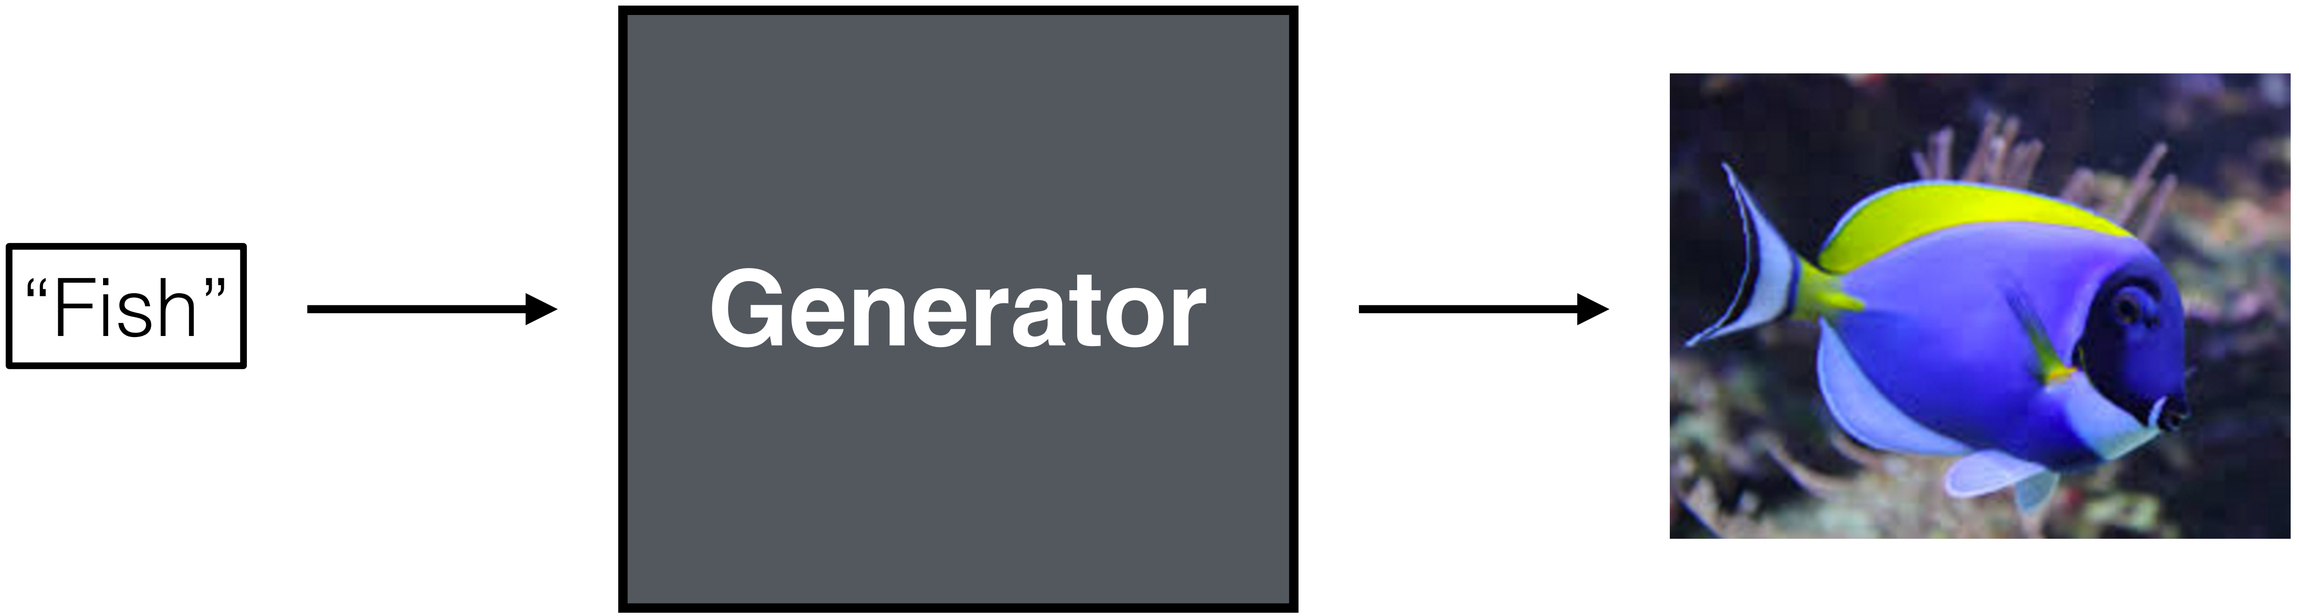

label y

image **x**

Image translation

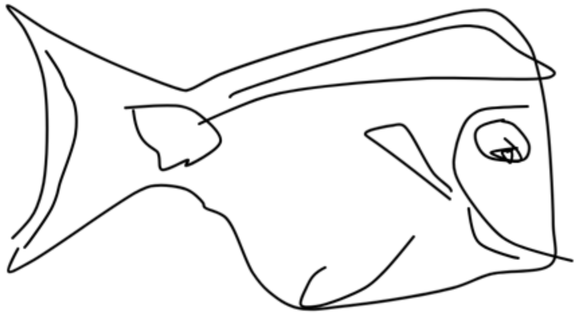

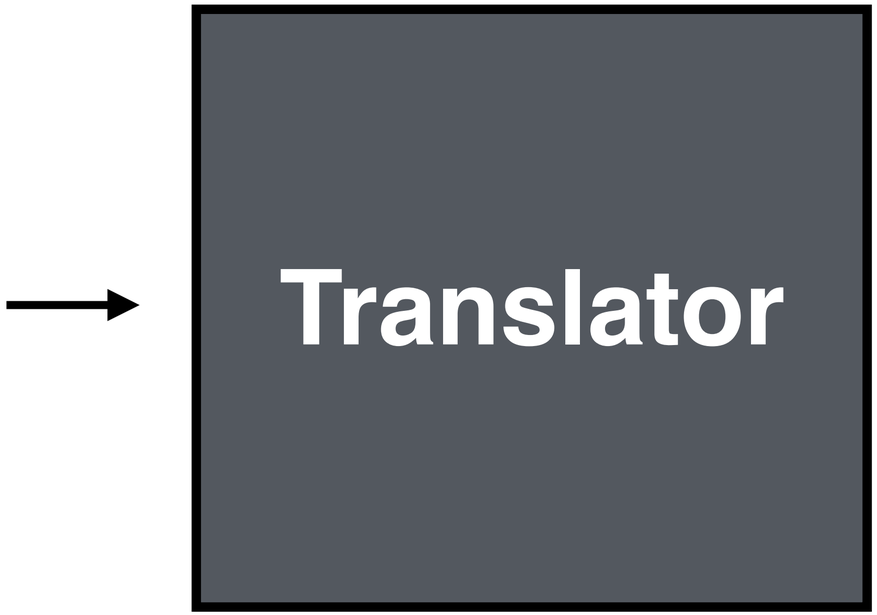

User sketch

Photo

Image translation

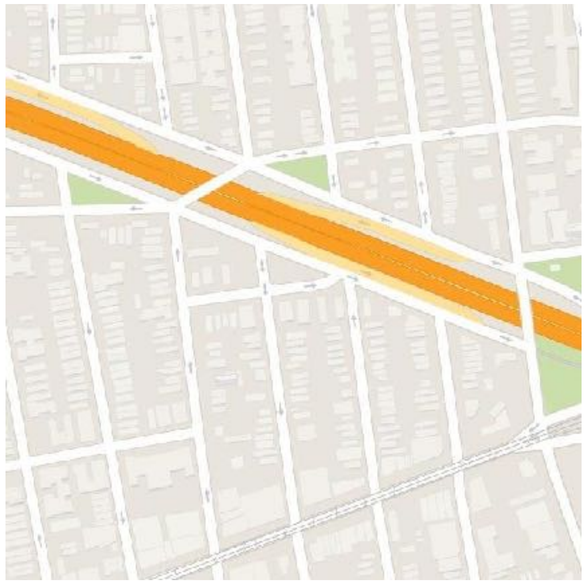

Translator

Google Map

Satellite photo

Image translation

BW image

Translator

Color image

Image synthesis via **generative modeling**

X is high-dimensional!

Model of high-dimensional structured data $P(\mathbf{X}|\mathbf{Y}=\mathbf{y})$

In vision, this is usually what we are interested in!

What can you do with generative models?

1. Image synthesis

2. Structured prediction

3. Domain mapping

1. **Image synthesis**
2. Structured prediction
3. Domain mapping

[Images: https://ganbreeder.app/]

Image synthesis

Procedural graphics

[Anders Scheil]

Ayliean @Ayliean · Nov 17

Made up a set of rules and rolled some dice to decide how this plant would grow. I never did get that five of a kind, as expected, but I was still hopeful! 🌱🍀🌿

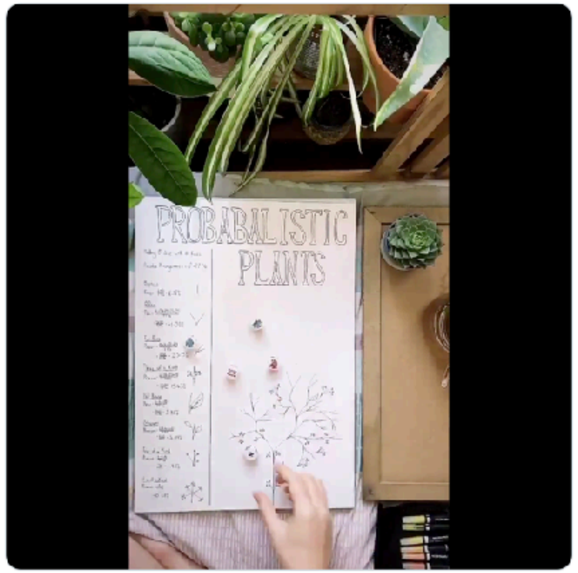

52 1.1K 4.5K

Image synthesis from “noise”

G

$z \sim p(z)$

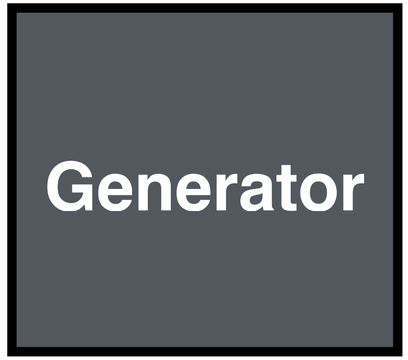

$x = G(z)$

Sampler

$$G : \mathcal{Z} \rightarrow \mathcal{X}$$

$$z \sim p(z)$$

$$x = G(z)$$

Learning a generative model

Input samples

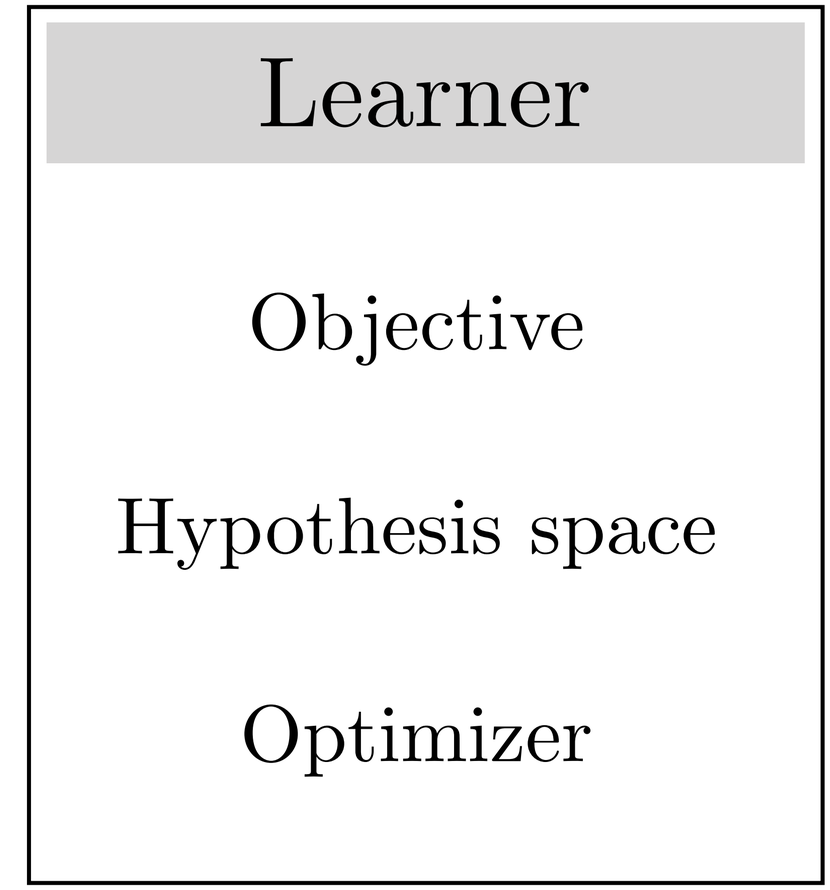

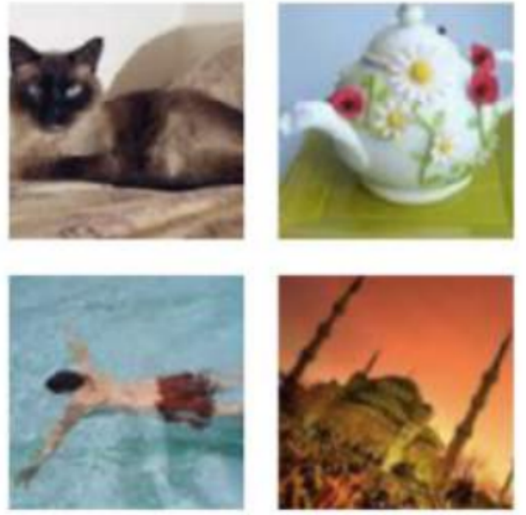

Generated samples

[figs modified from: http://introtodeeplearning.com/materials/2019_6S191_L4.pdf]

Learning a density model

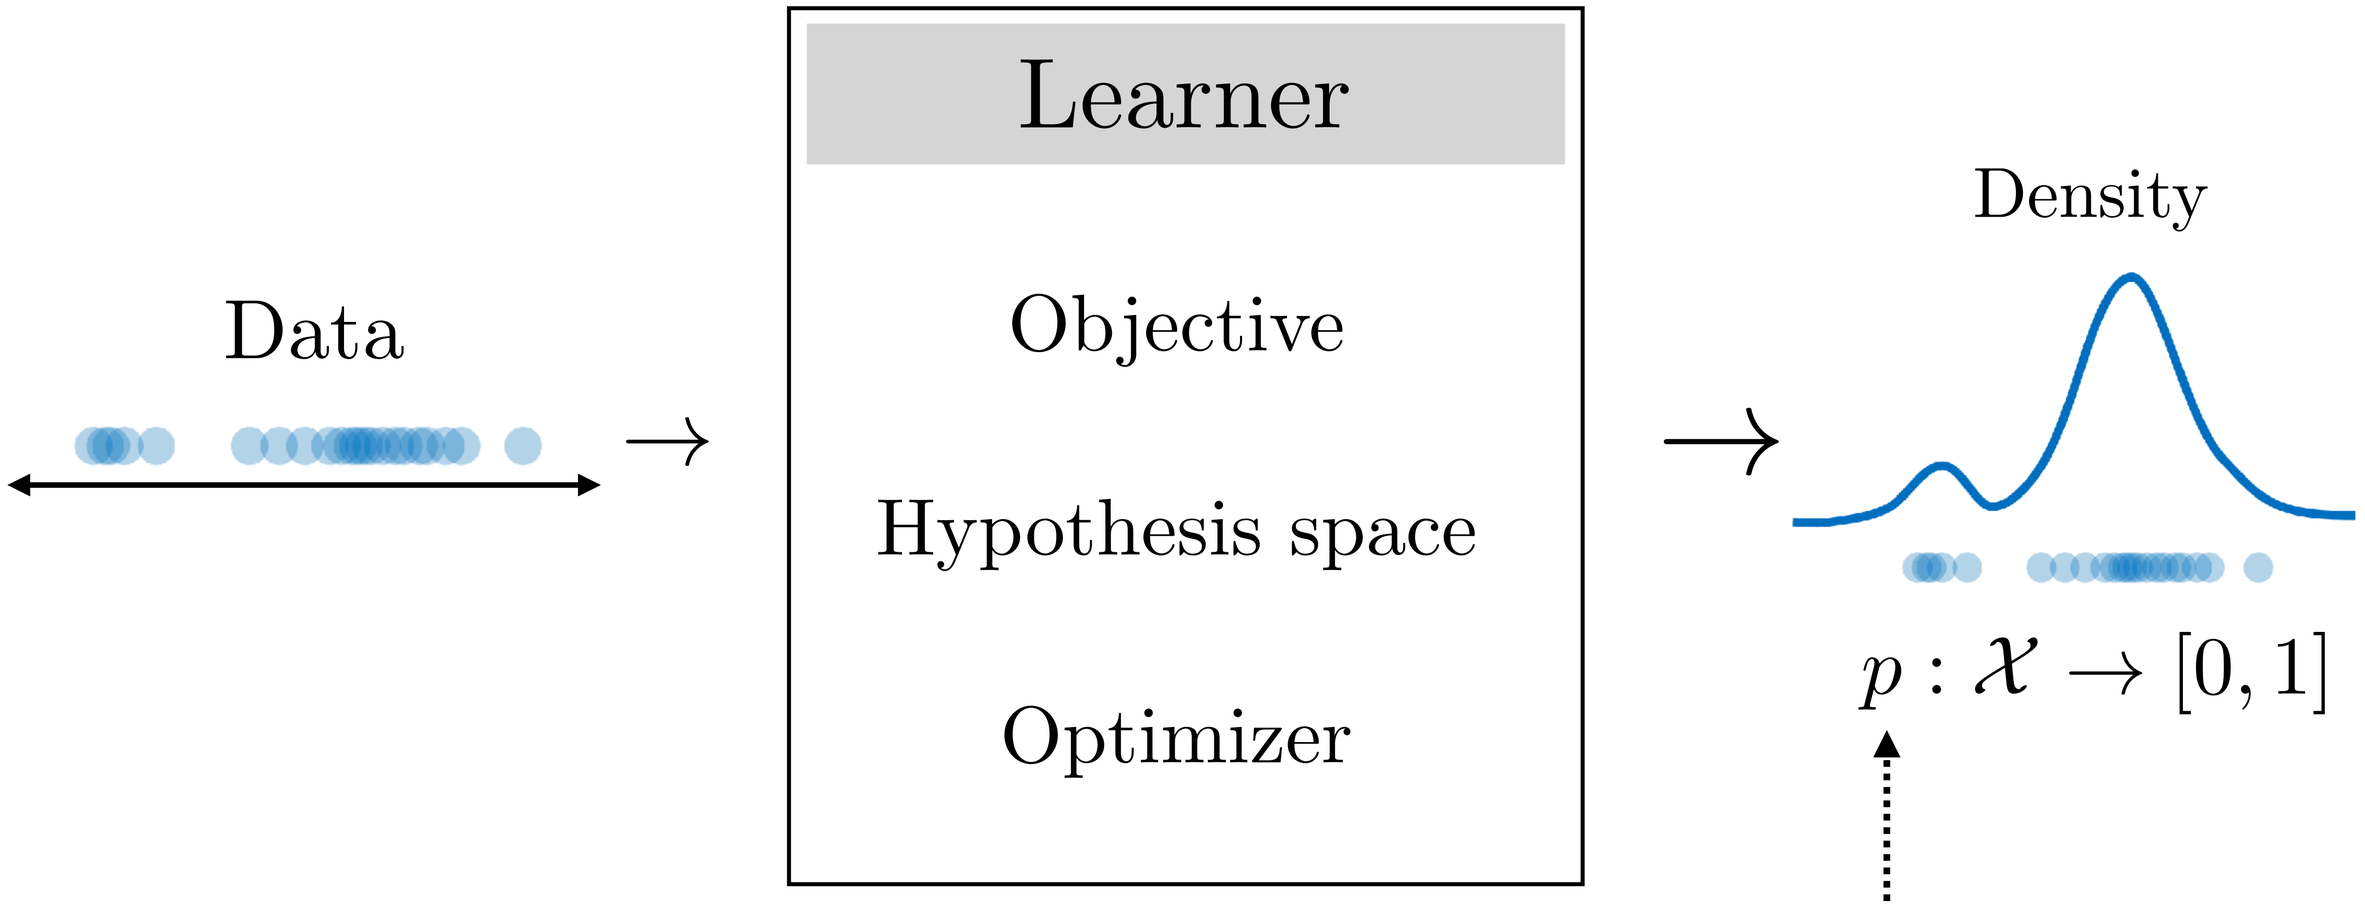

[figs modified from: http://introtodeeplearning.com/materials/2019_6S191_L4.pdf]

Case study #1: Fitting a Gaussian to data

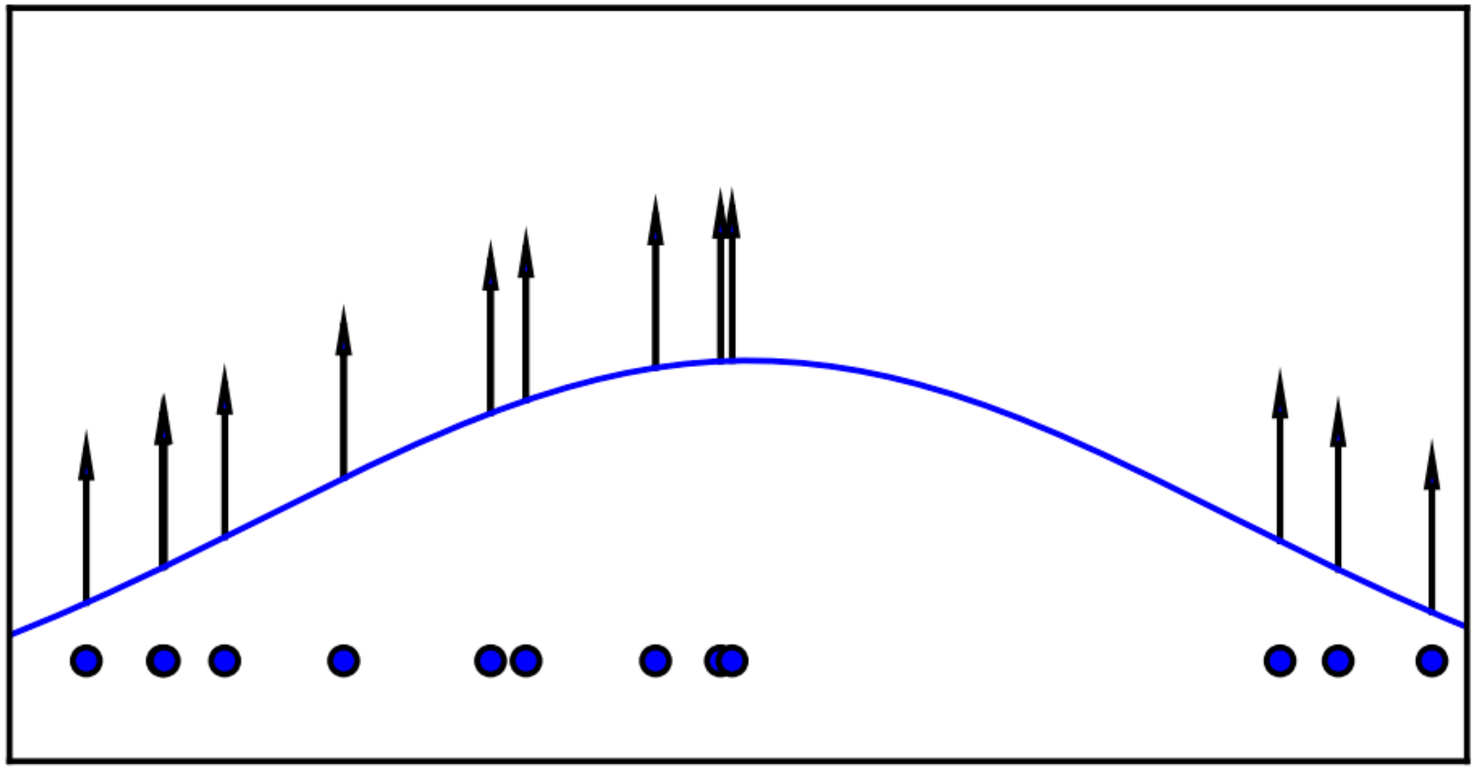

fig from [Goodfellow, 2016]

Max likelihood objective

$$\max_{\theta} \mathbb{E}_{x \sim p_{\text{data}}} [\log p_\theta(x)]$$

Considering only Gaussian fits

$$p_\theta(x) = \mathcal{N}(x; \mu, \sigma)$$

$$\theta = [\mu, \sigma]$$

Closed form optimum:

$$\mu = \frac{1}{N} \sum_{i=1}^{N} x_i \qquad \sigma^2 = \frac{1}{N} \sum_{i=1}^{N} (x_i - \mu)^2$$

Case study #1: Fitting a Gaussian to data

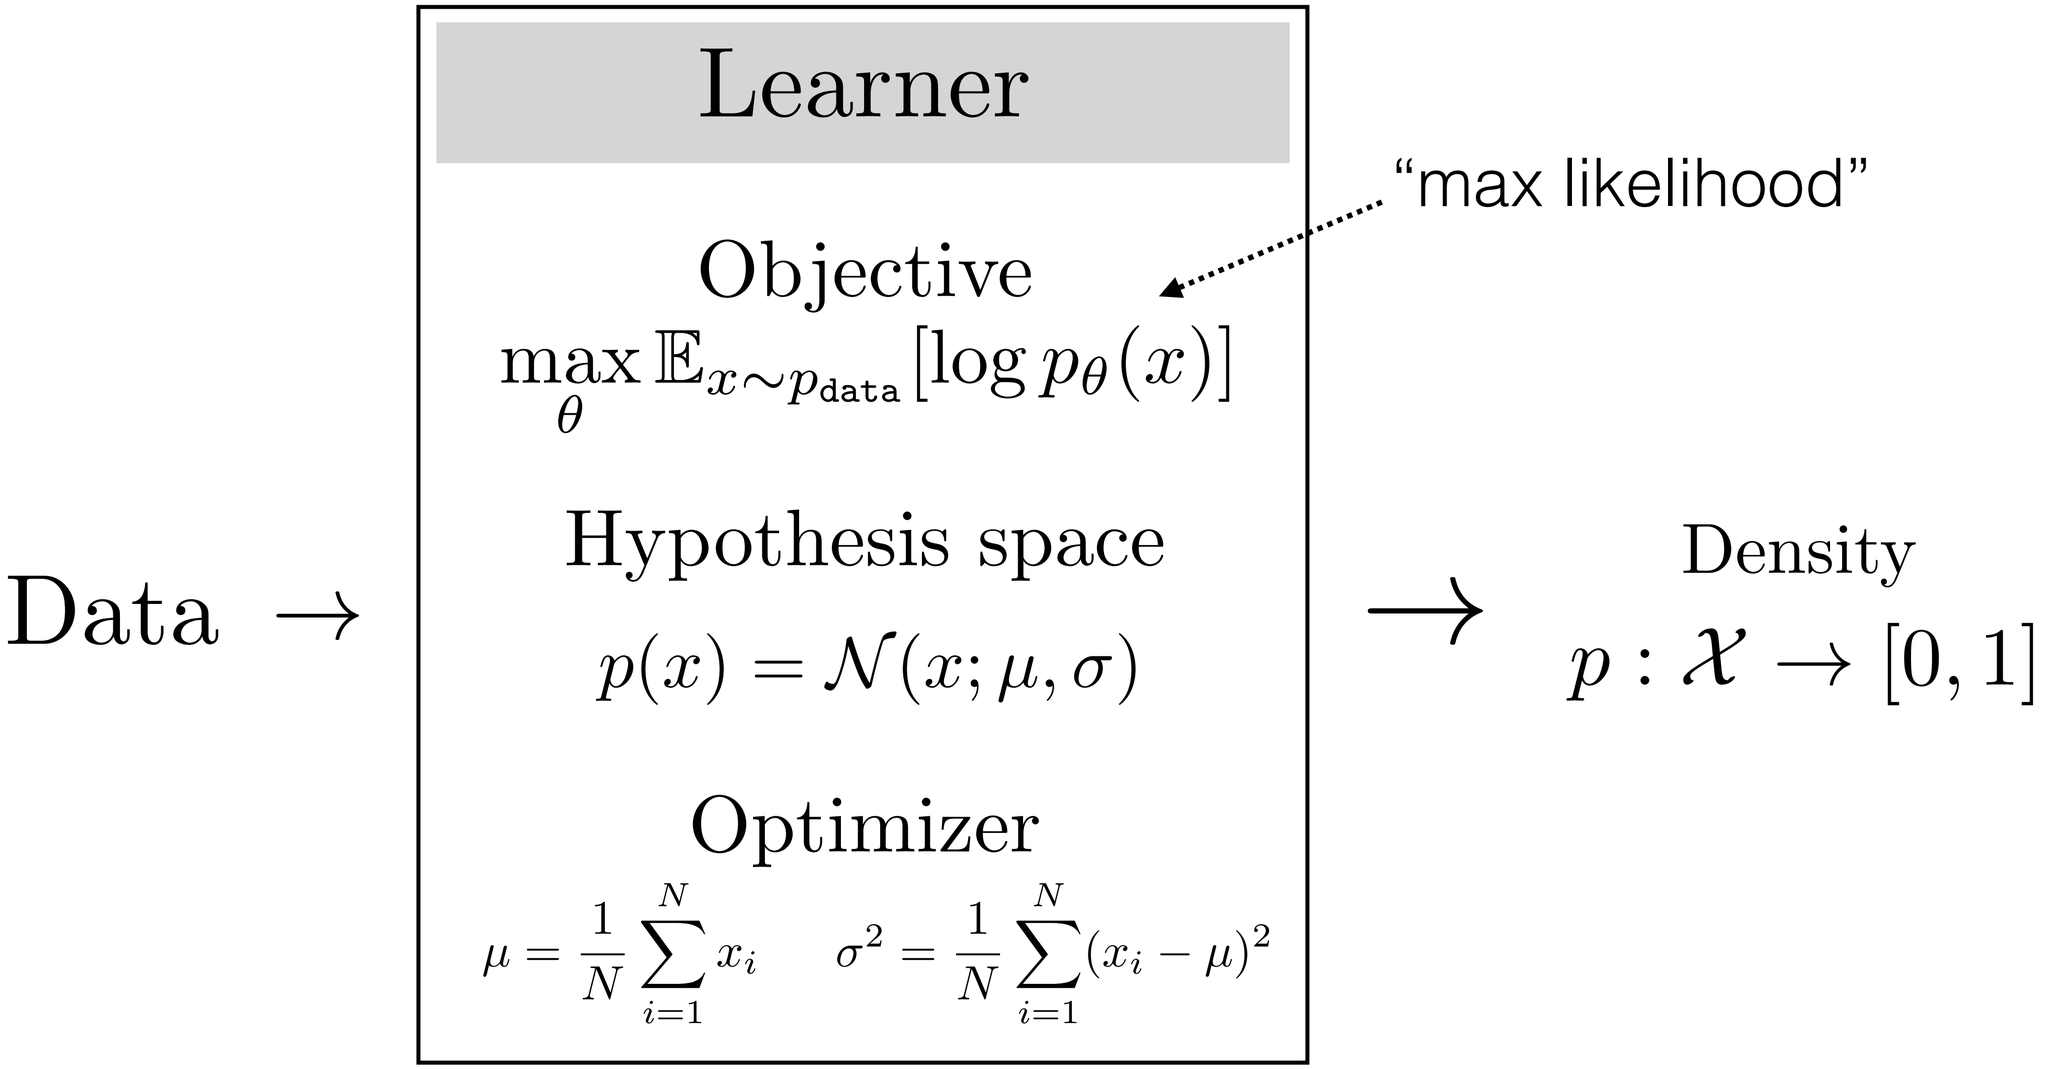

Case study #2: learning a deep generative model

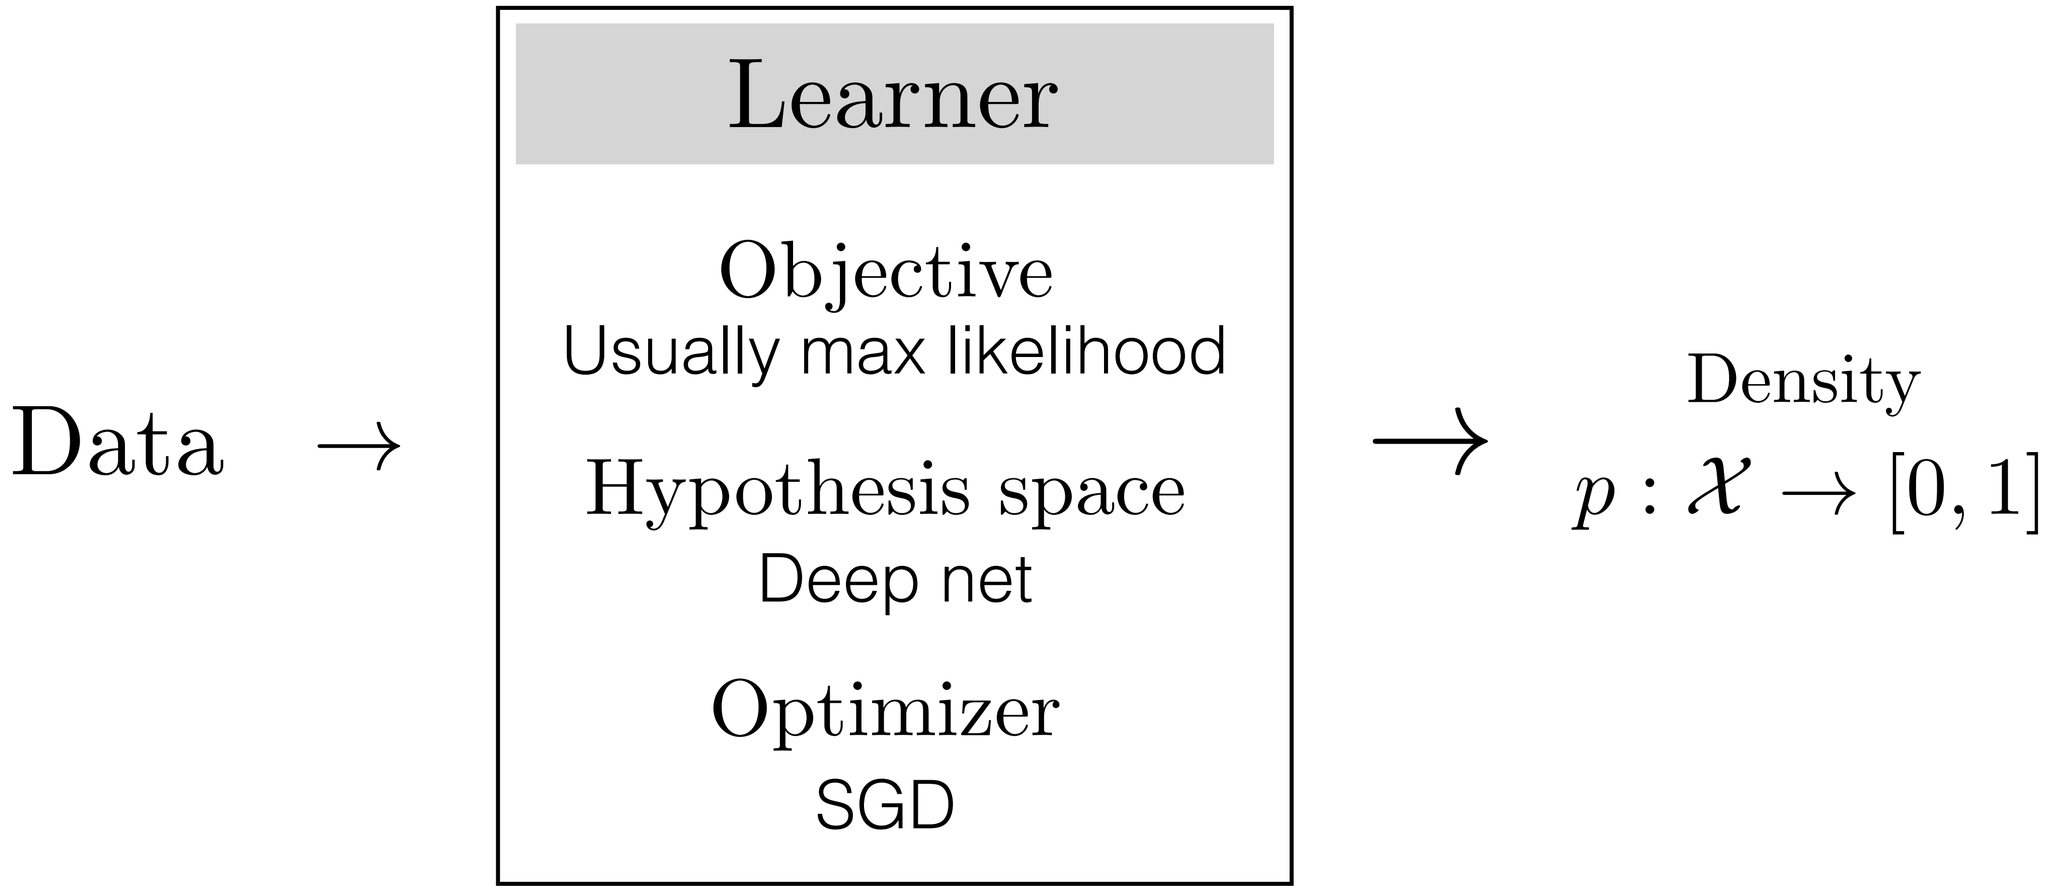

Case study #2: learning a deep generative model

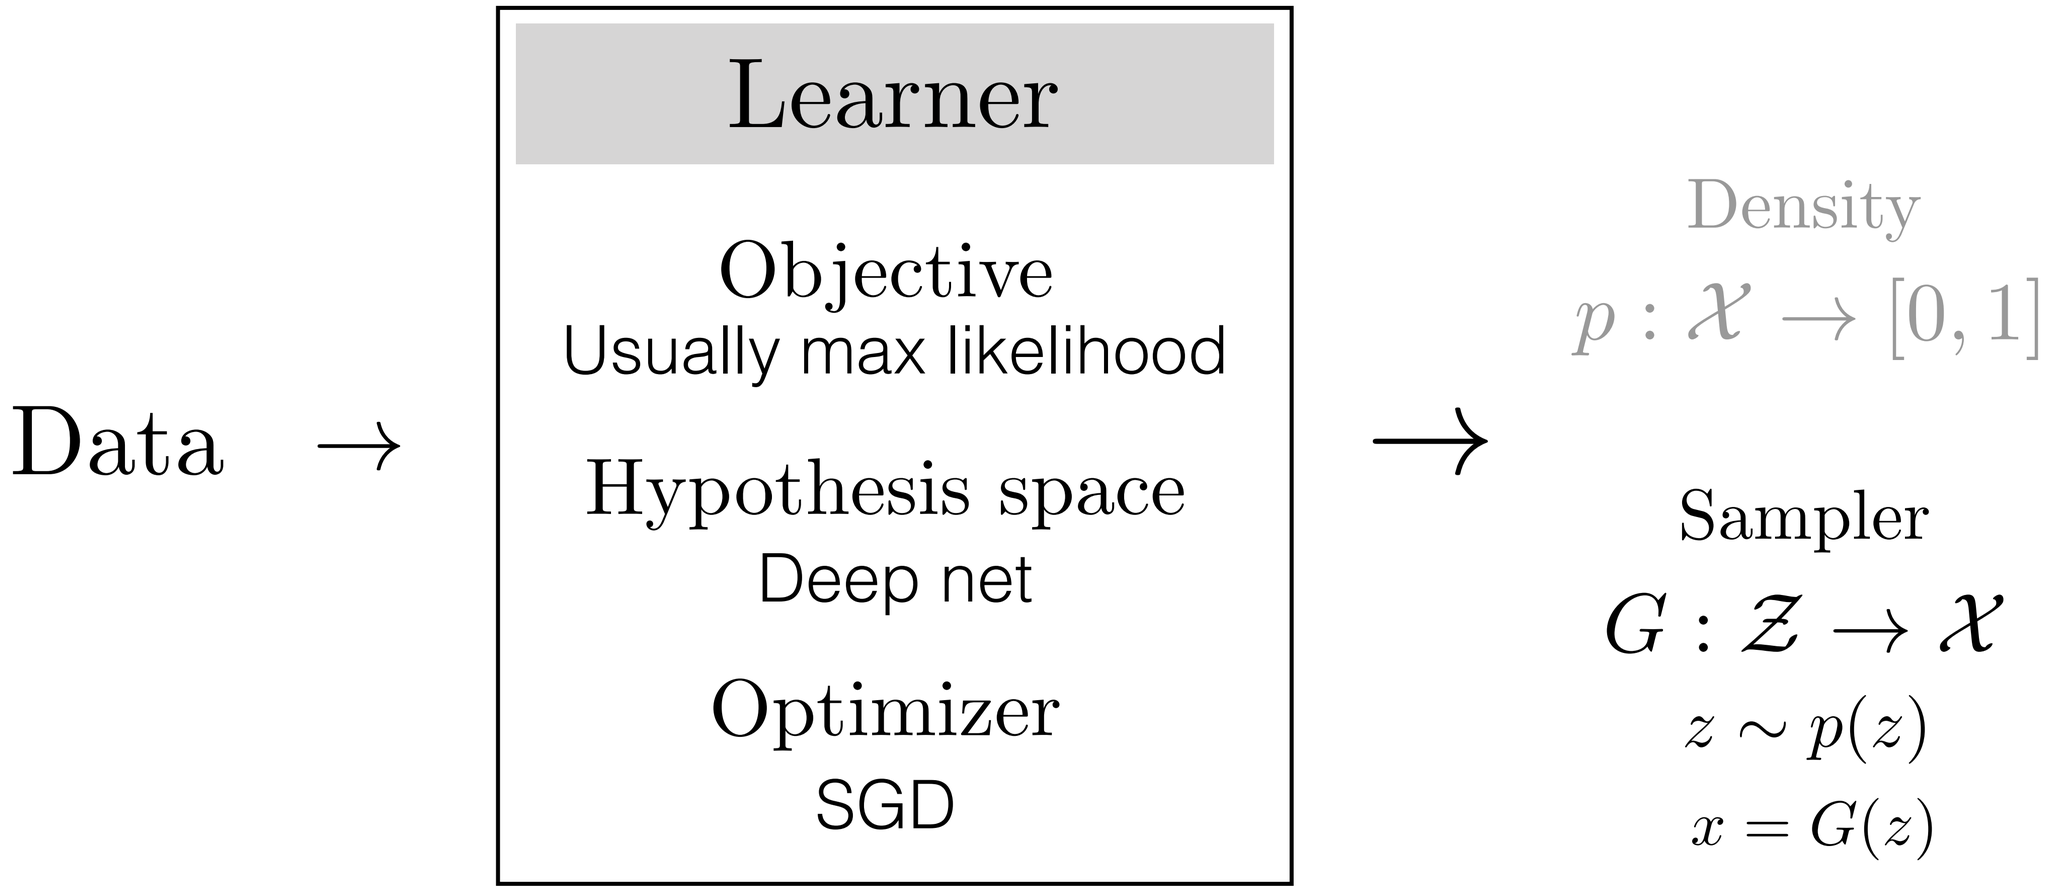

Models that provide a sampler but no density are called **implicit generative models**

Deep generative models are distribution transformers

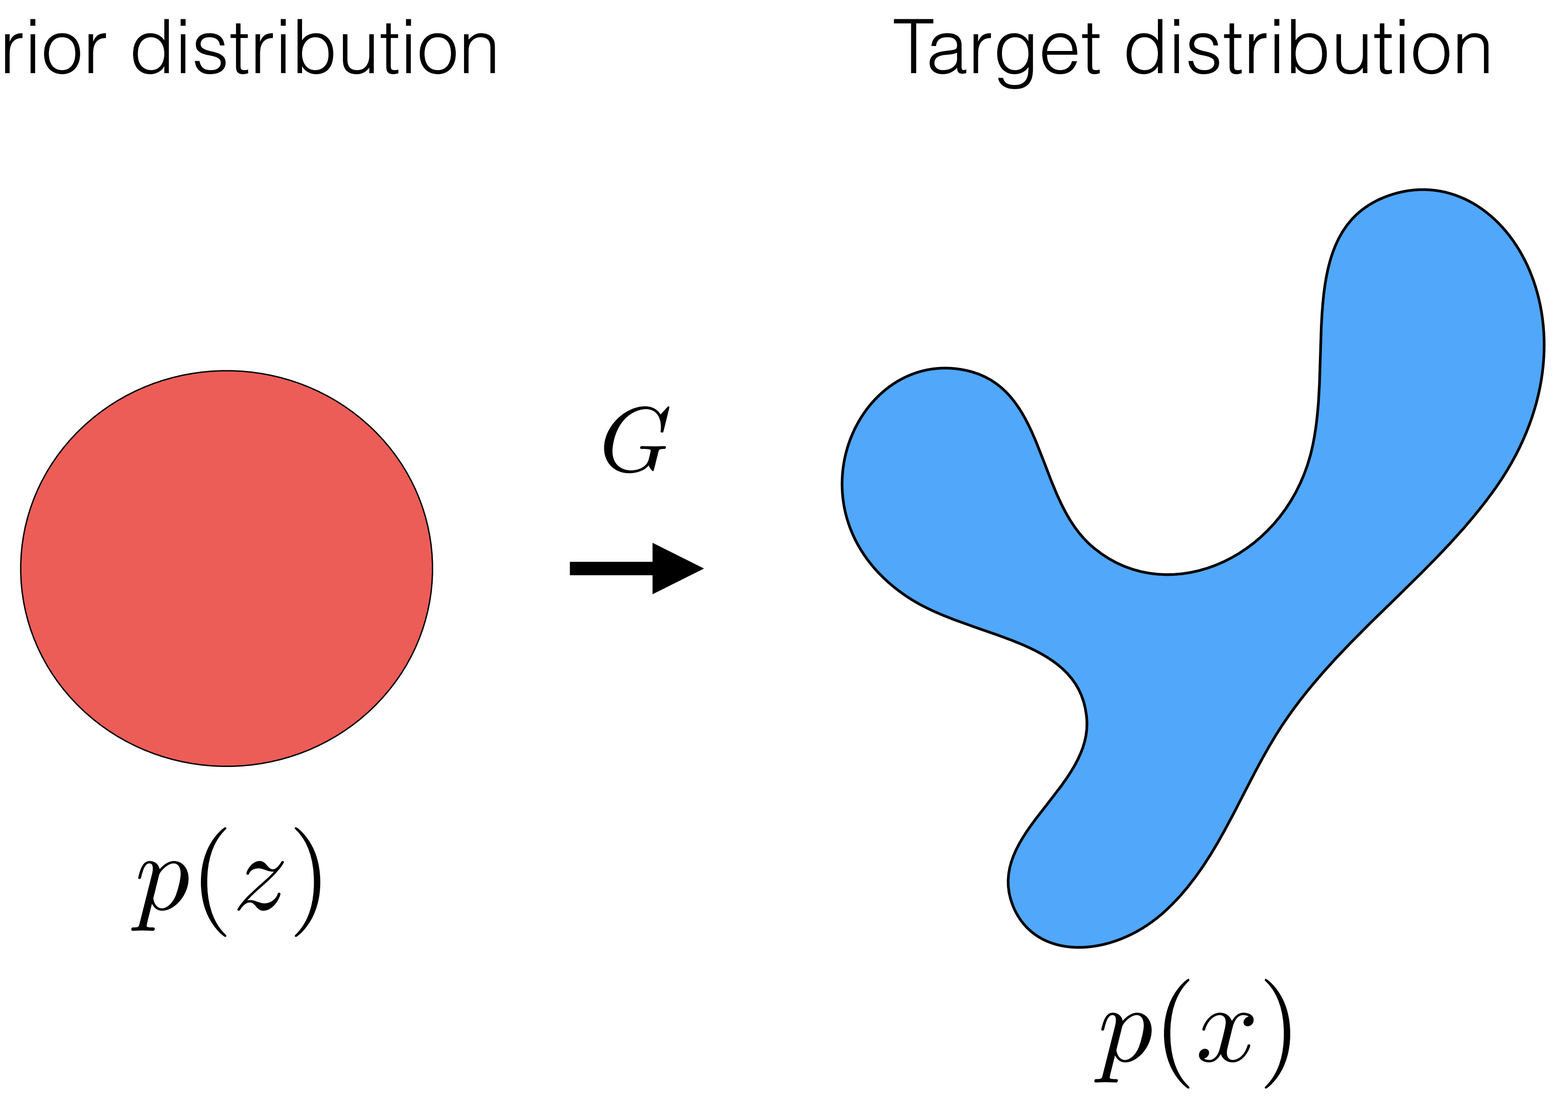

Deep generative models are distribution transformers

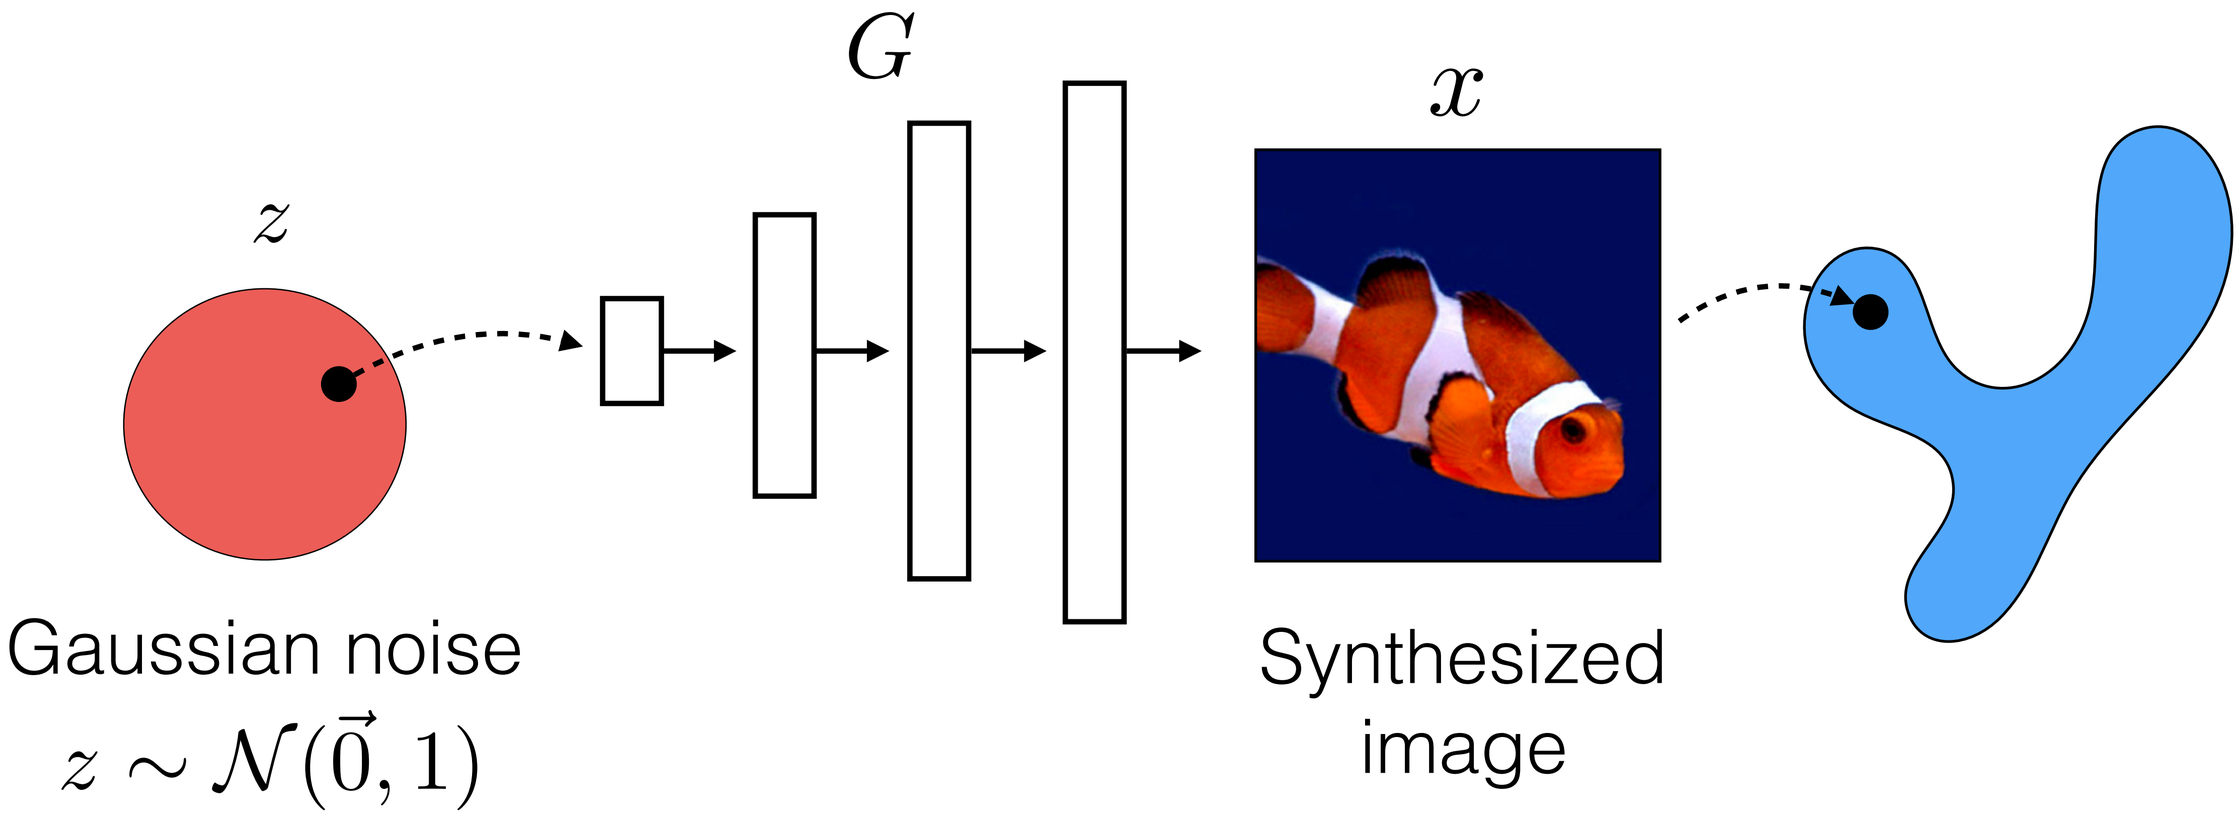

Deep generative models are distribution transformers

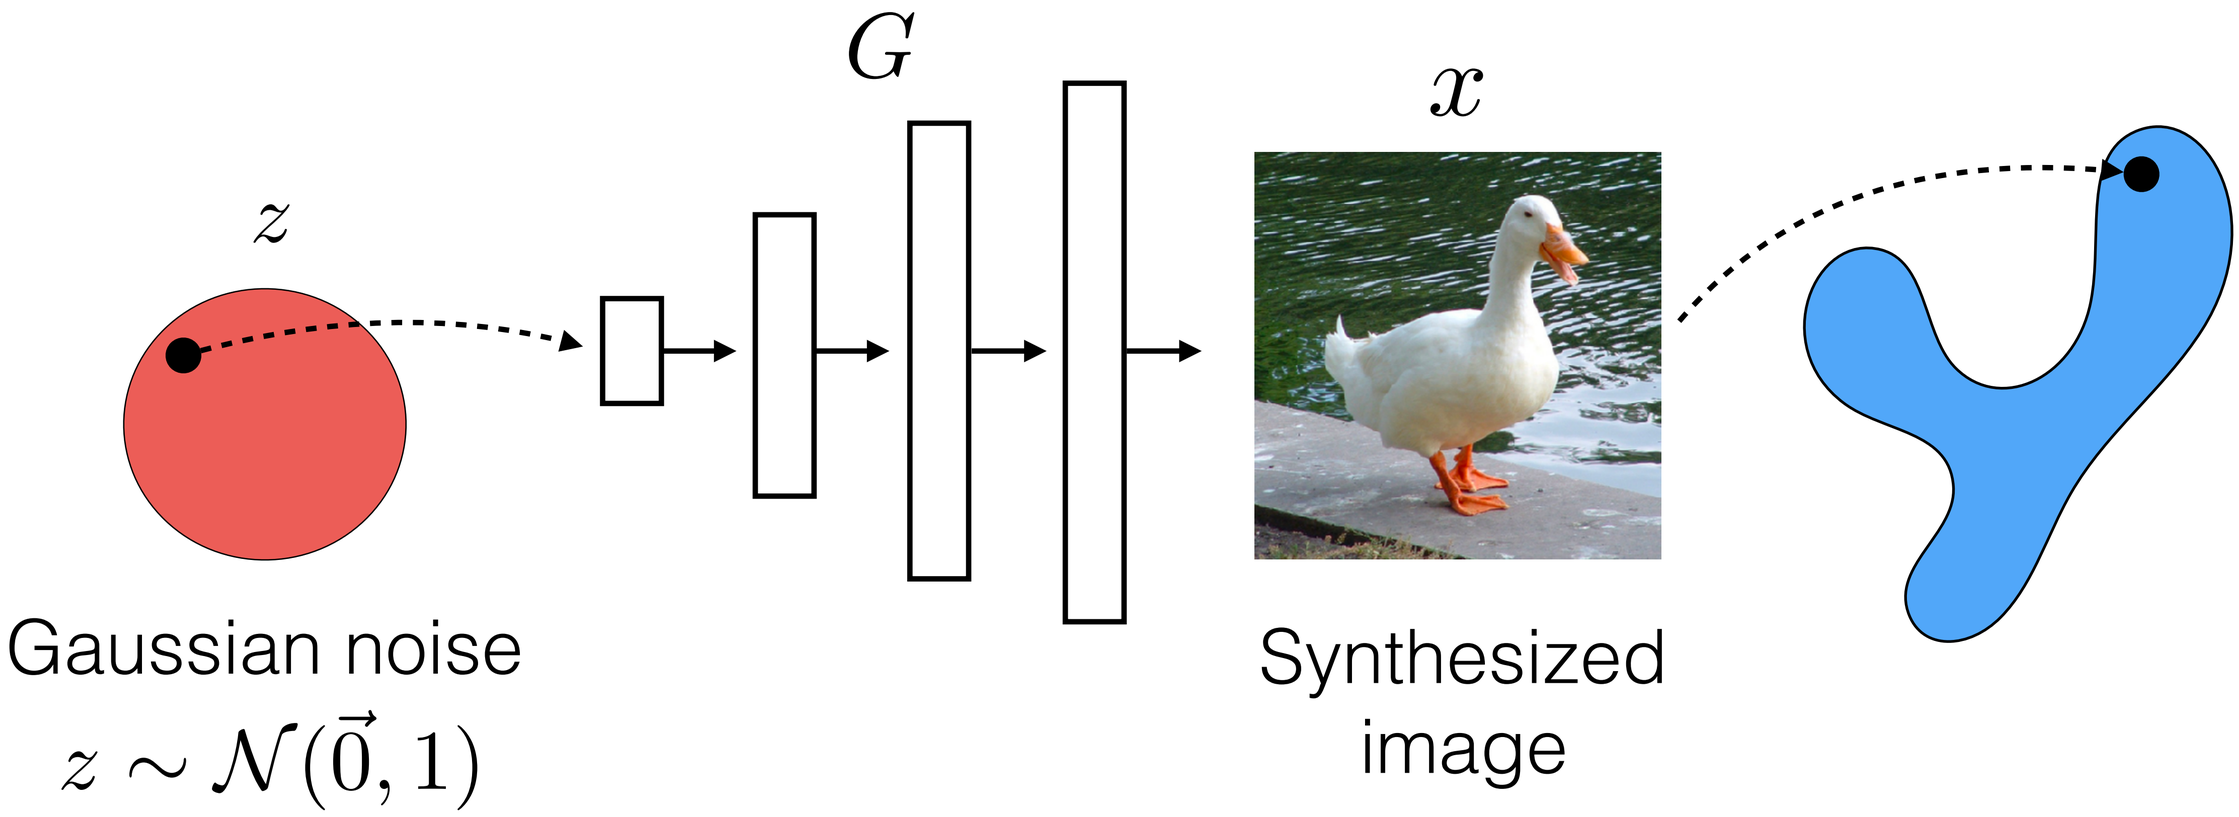

Generative Adversarial Networks (GANs)

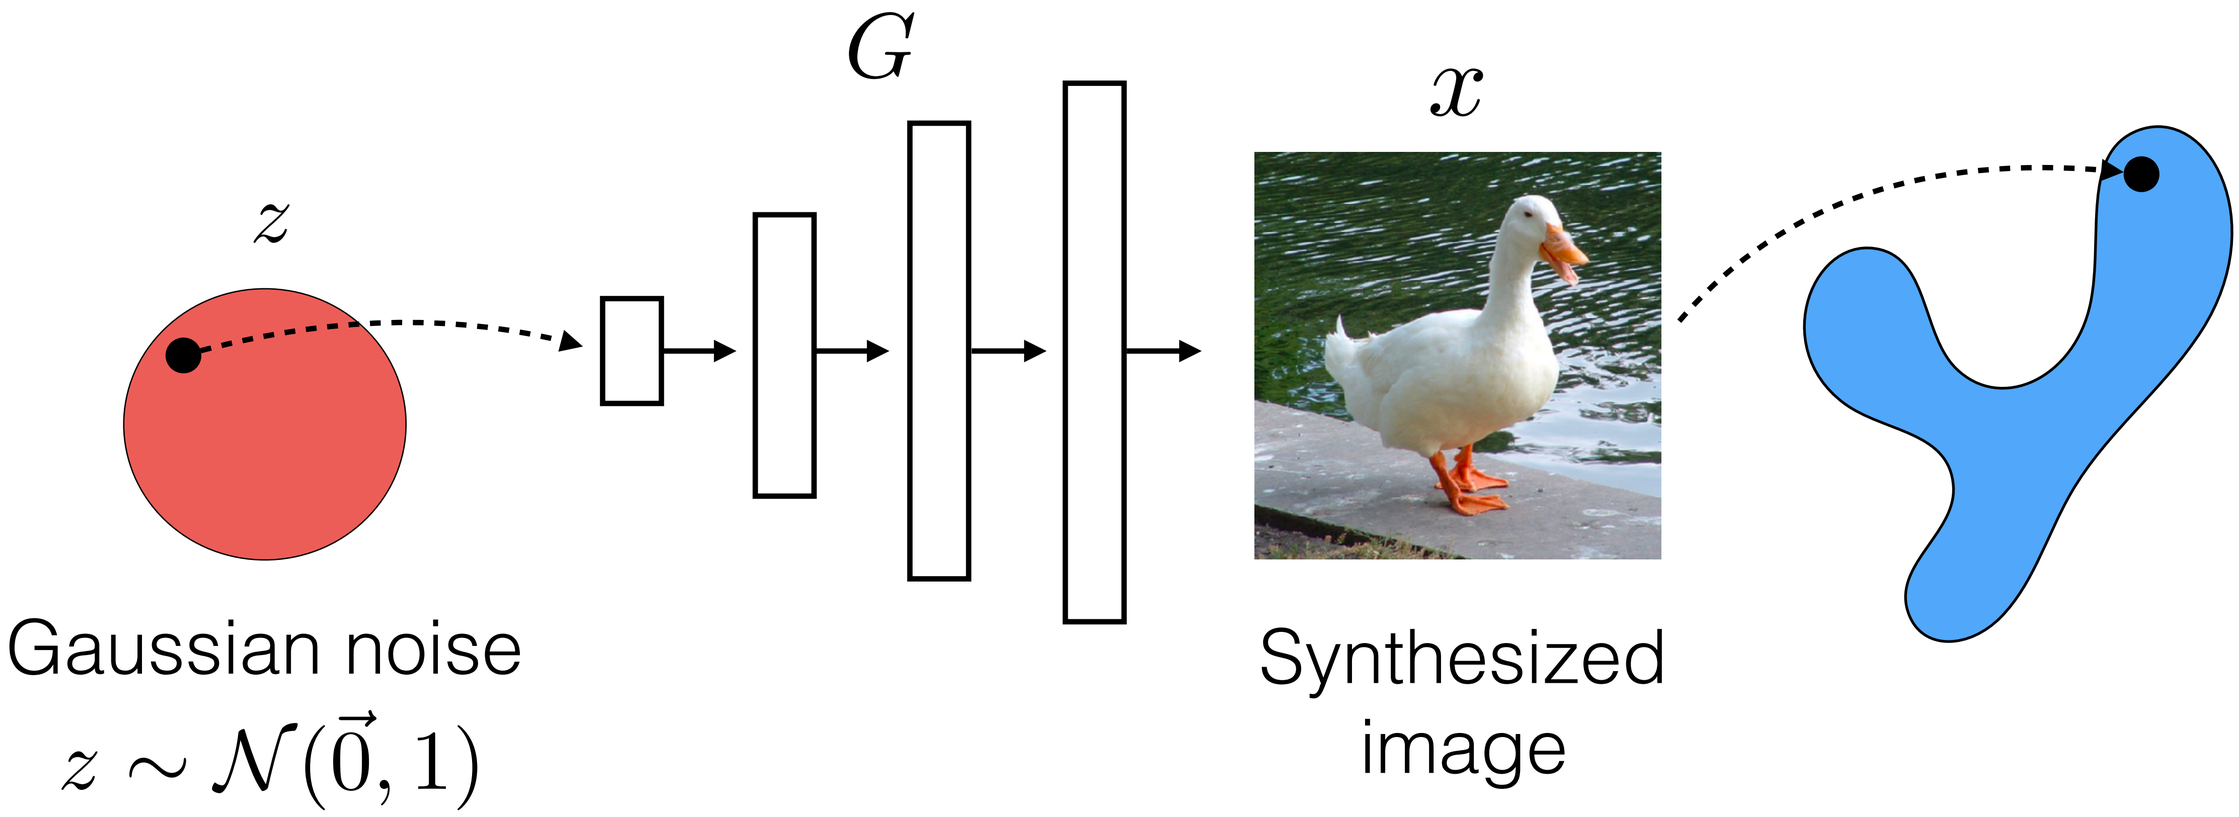

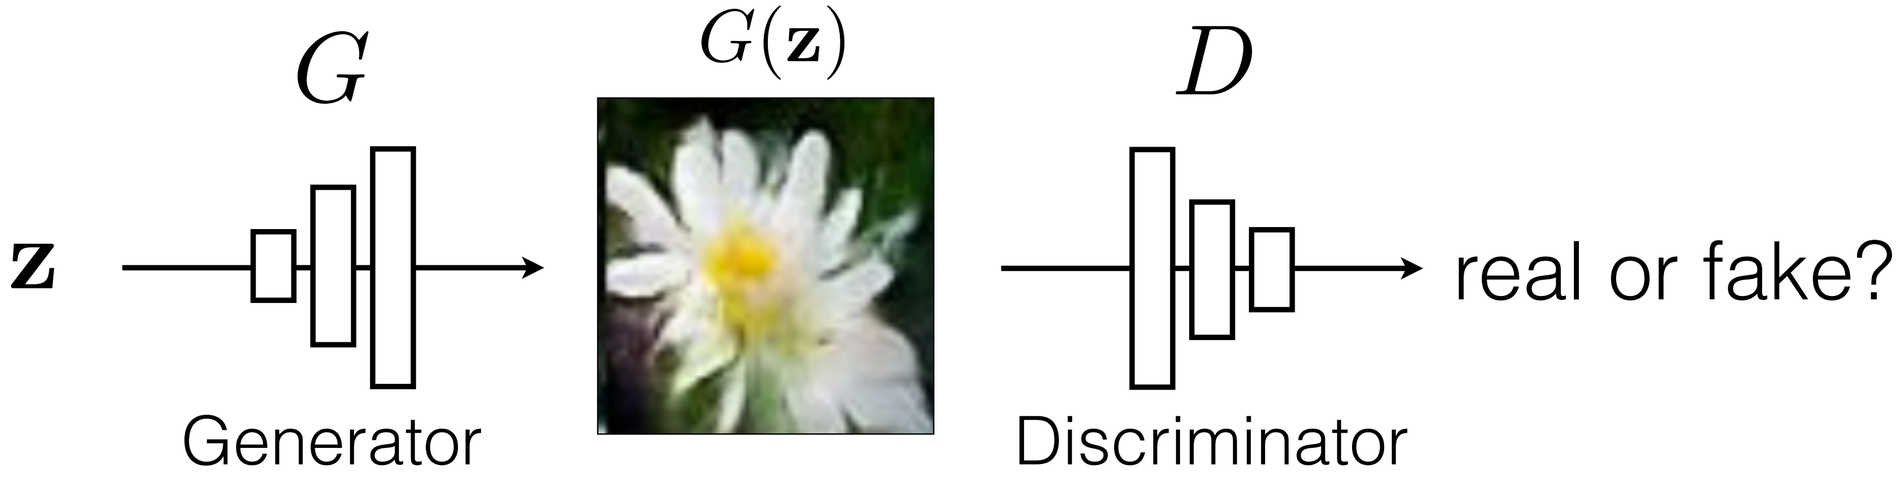

G tries to synthesize fake images that fool **D**

D tries to identify the fakes

[Goodfellow et al., 2014]

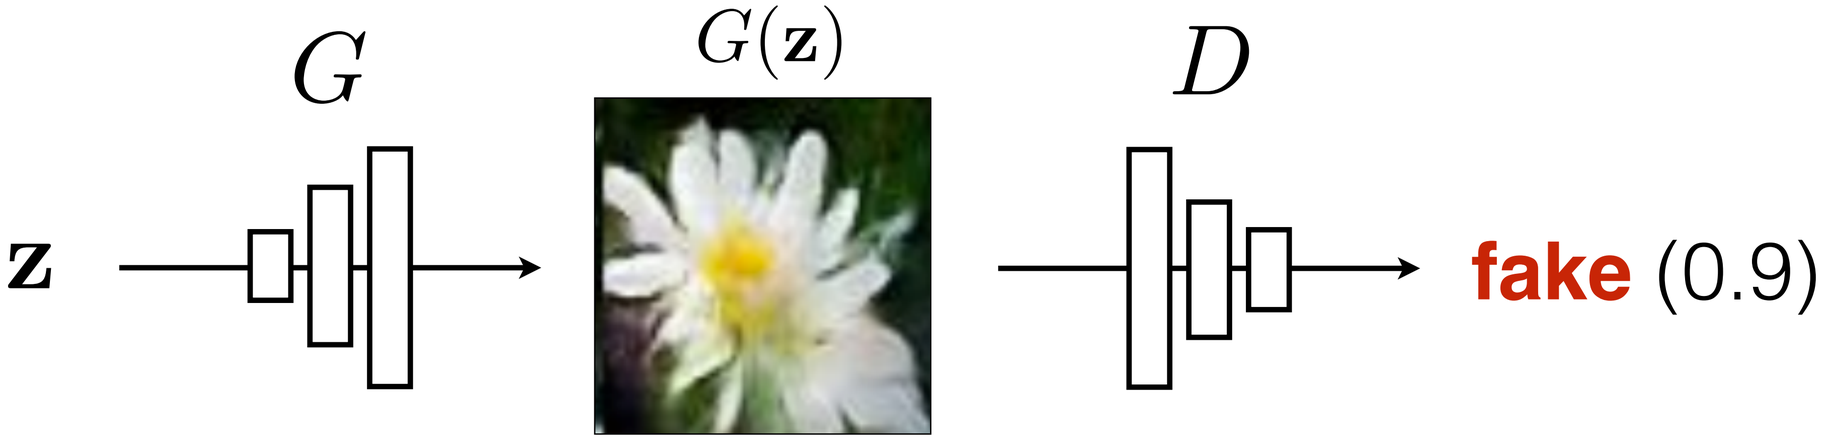

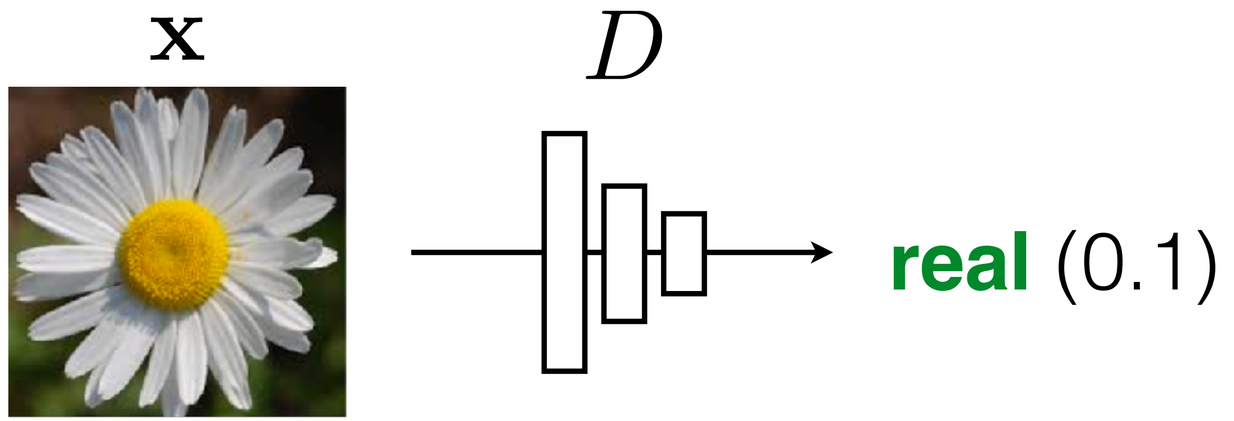

$$\arg\max_D \mathbb{E}_{\mathbf{z},\mathbf{x}}[\ \log D(G(\mathbf{z}))\ +\ \log\left(1 - D(\mathbf{x})\right)\ ]$$

[Goodfellow et al., 2014]

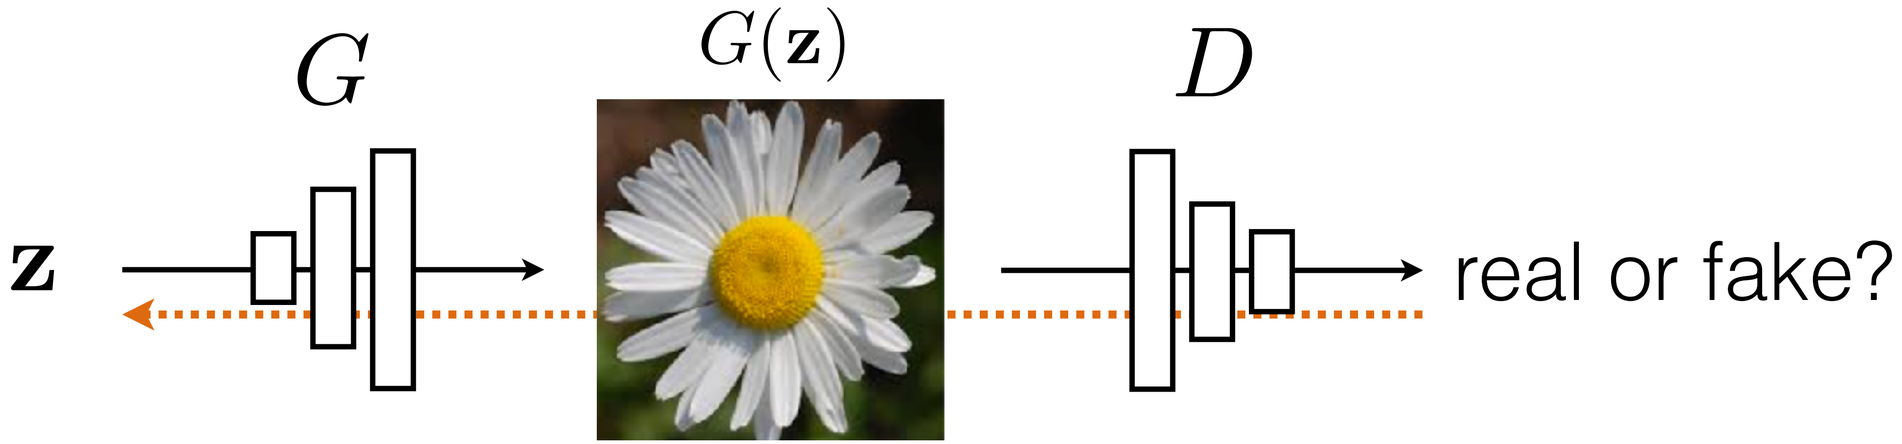

G tries to synthesize fake images that ***fool*** **D**:

$$\arg\min_{G} \mathbb{E}_{\mathbf{z},\mathbf{x}}[\ \log D(G(\mathbf{z}))\ +\ \log(1 - D(\mathbf{x}))\ ]$$

[Goodfellow et al., 2014]

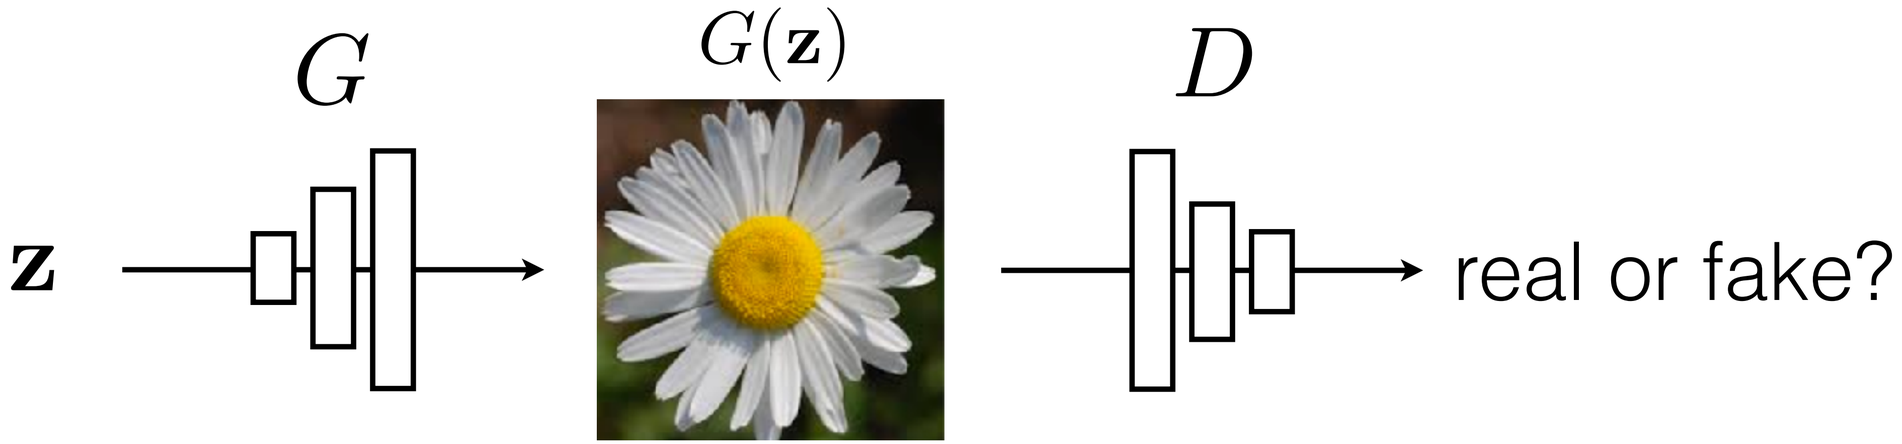

G tries to synthesize fake images that ***fool*** the ***best*** **D**:

$$\arg \min_G \max_D \; \mathbb{E}_{\mathbf{z},\mathbf{x}}[\;\; \log D(G(\mathbf{z})) \;\; + \;\; \log(1 - D(\mathbf{x})) \;\;\;]$$

[Goodfellow et al., 2014]

Training

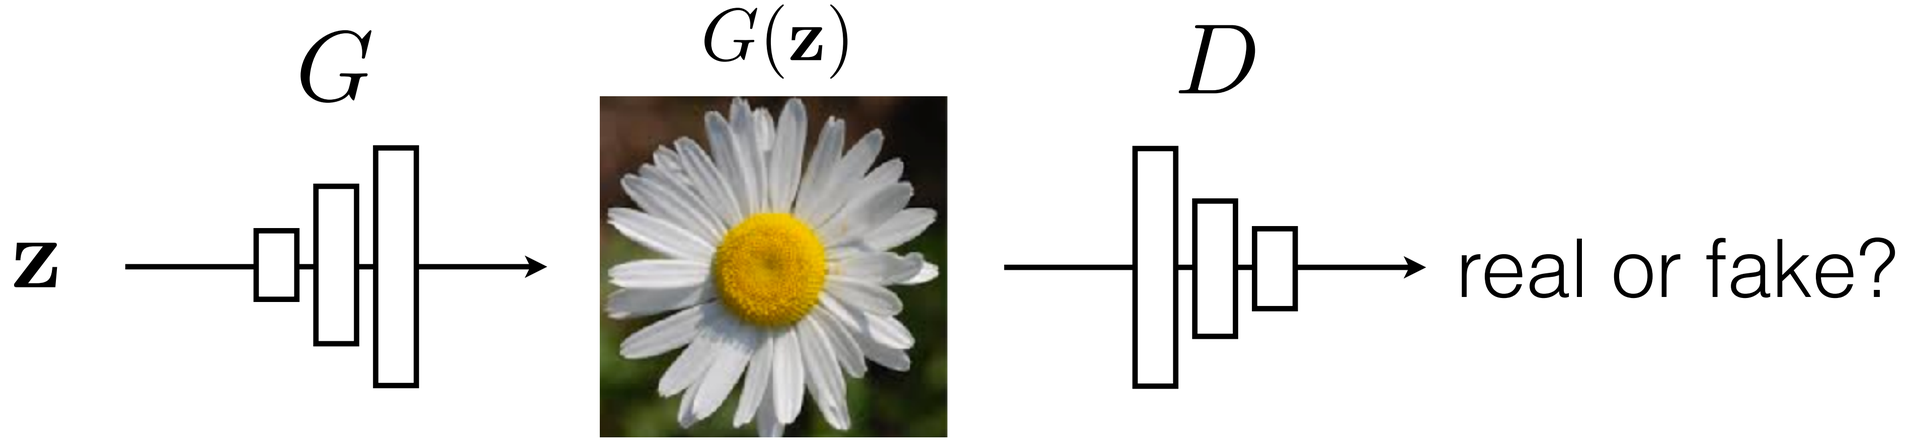

G tries to synthesize fake images that fool **D**

D tries to identify the fakes

- Training: iterate between training D and G with backprop.
- Global optimum when G reproduces data distribution.

[Goodfellow et al., 2014]

Samples from BigGAN

[Brock et al. 2018]

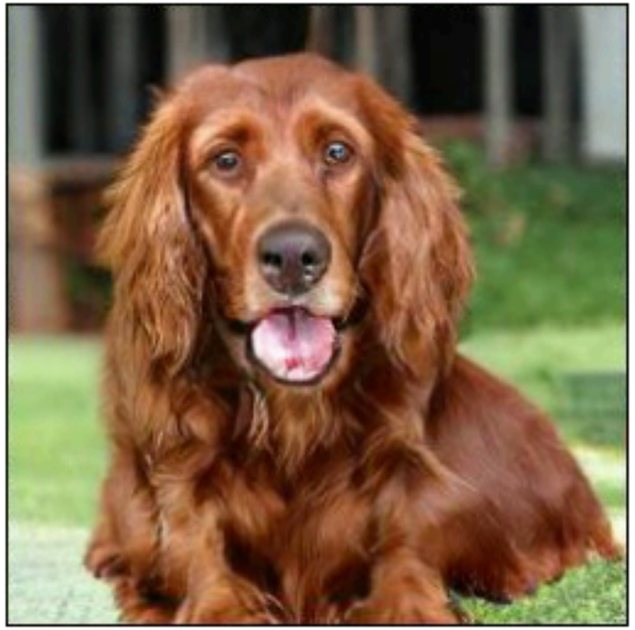

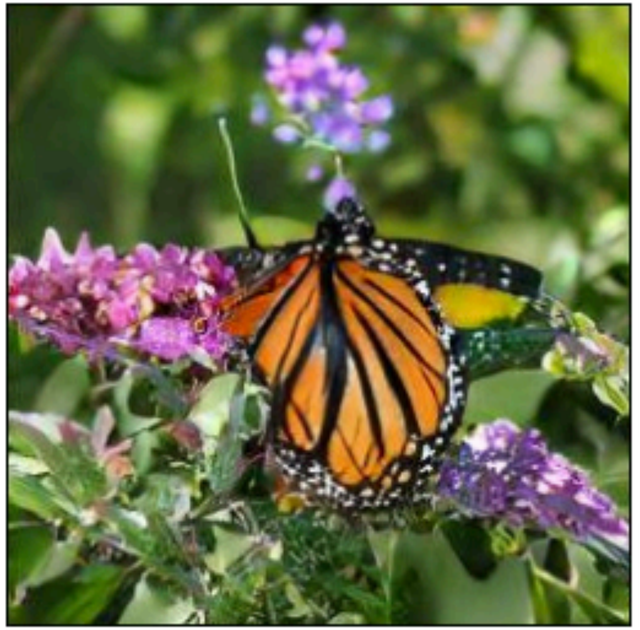


Latent space
(Gaussian)

$\mathbf{z}$

Data space
(Natural image manifold)

$\mathbf{x}$

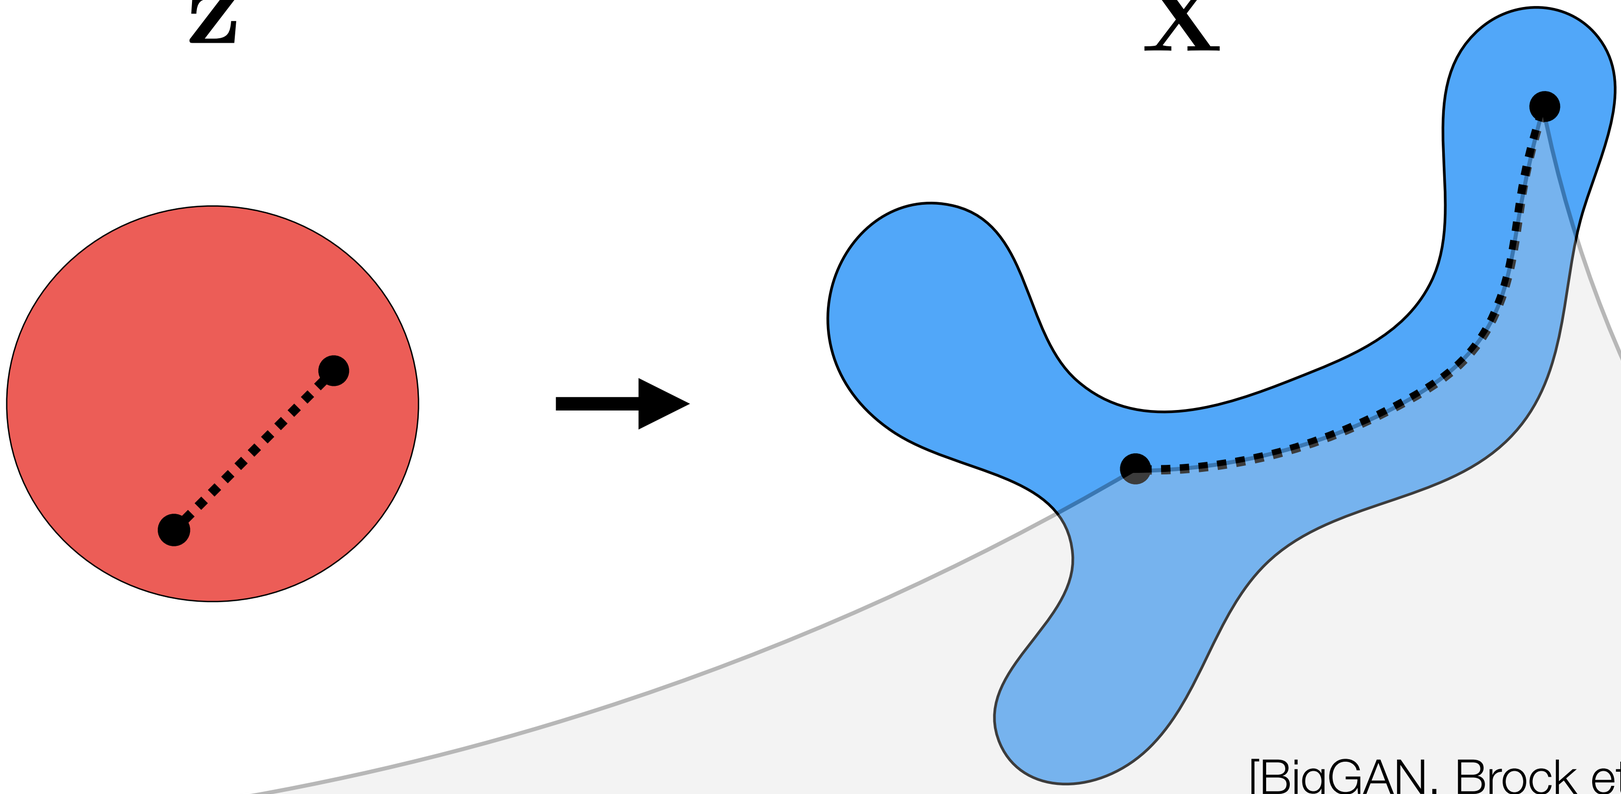

[BigGAN, Brock et al. 2018]

Generative models organize the manifold of natural images

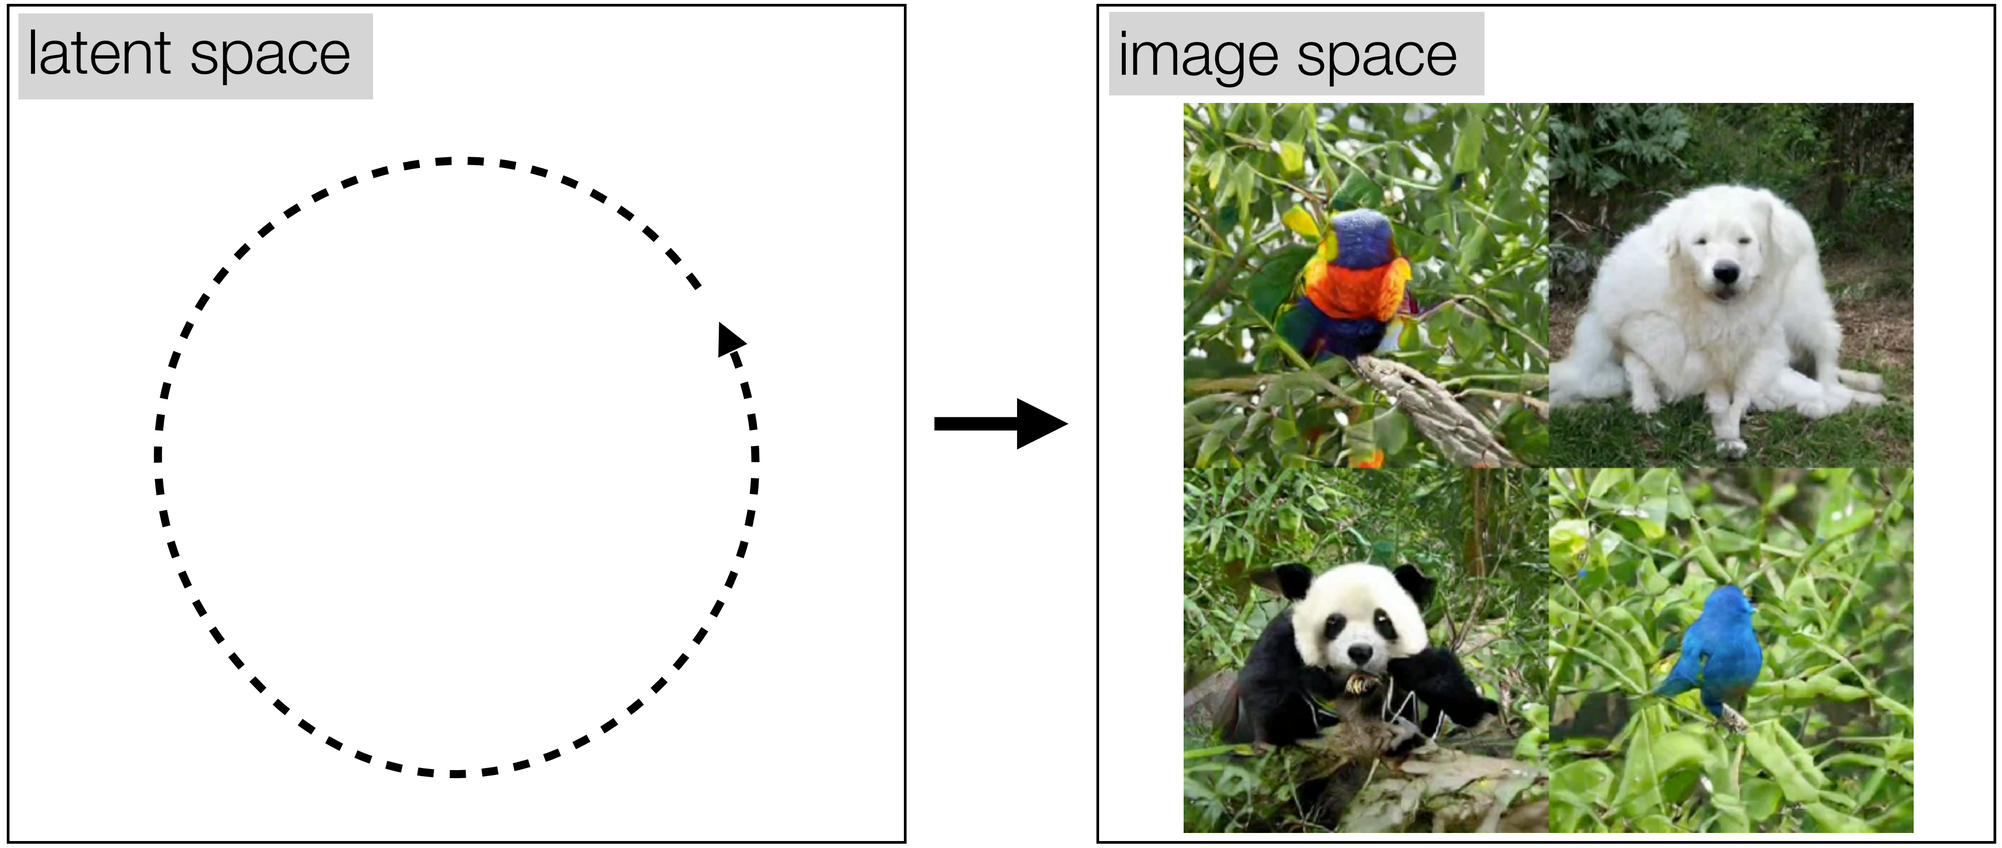

Generative Adversarial Network

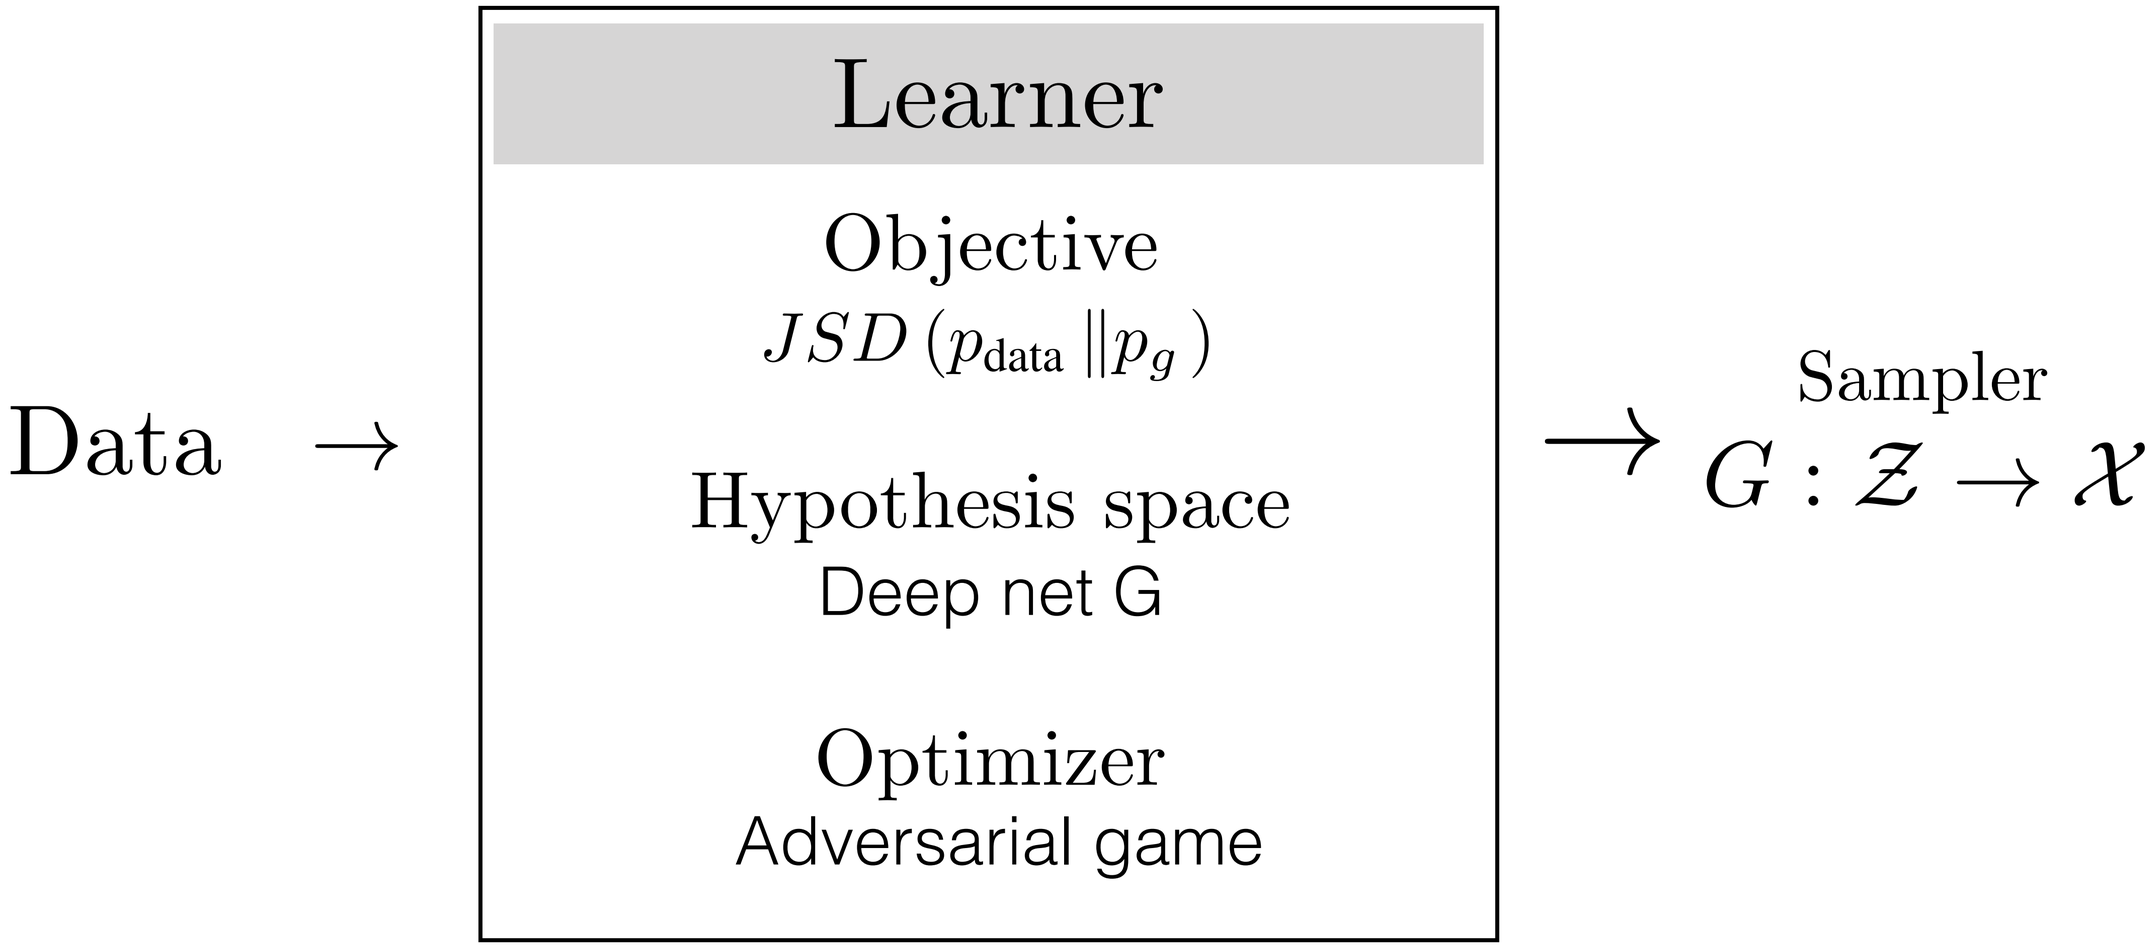

Generative Adversarial Network

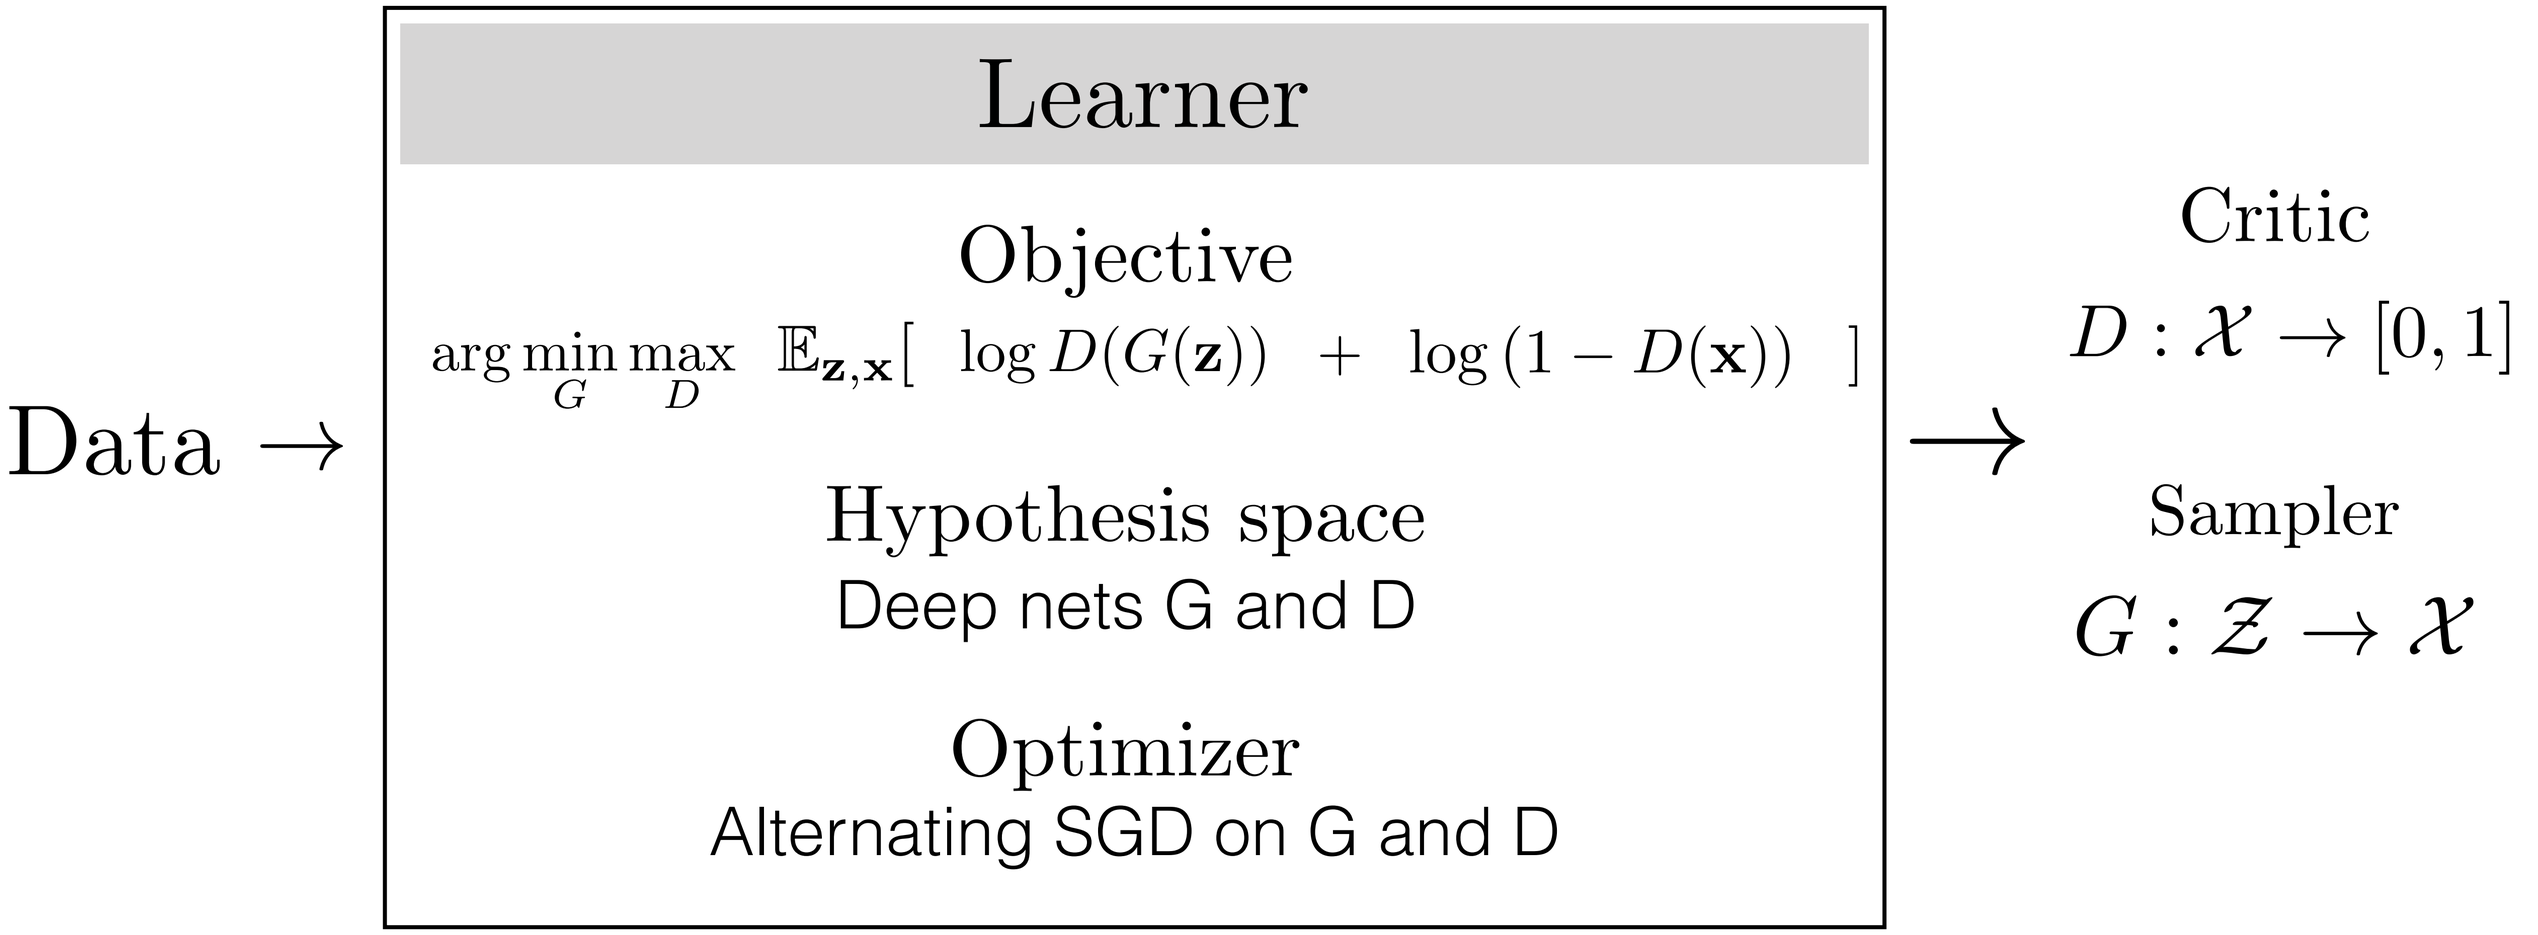

What has driven GAN progress?

Ian Goodfellow @goodfellow_ian · Jan 14

4.5 years of **GAN progress** on face generation. arxiv.org/abs/1406.2661 arxiv.org/abs/1511.06434 arxiv.org/abs/1606.07536 arxiv.org/abs/1710.10196 arxiv.org/abs/1812.04948

Better objectives? optimizers?

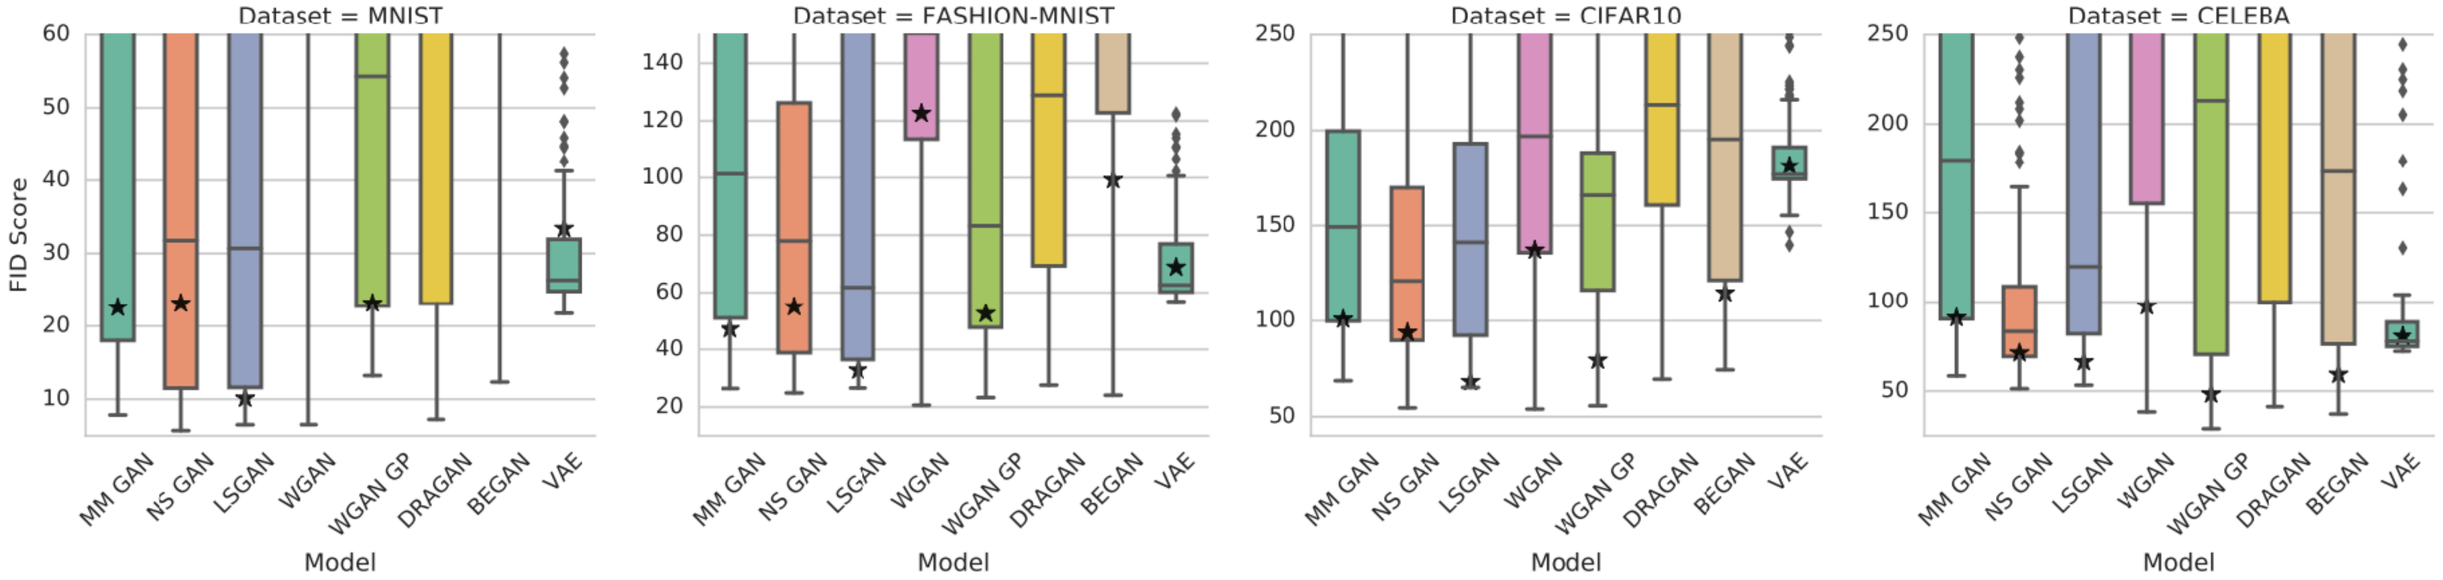

Figure 4: A *wide range* hyperparameter search (100 hyperparameter samples per model). Black stars indicate the performance of suggested hyperparameter settings. We observe that GAN training is extremely sensitive to hyperparameter settings and there is no model which is significantly more stable than others.

["Are all GANs Created Equal?", Lucic*, Kurach*, et al. 2018]

More data?

ACGAN [Odena et al. 2016]

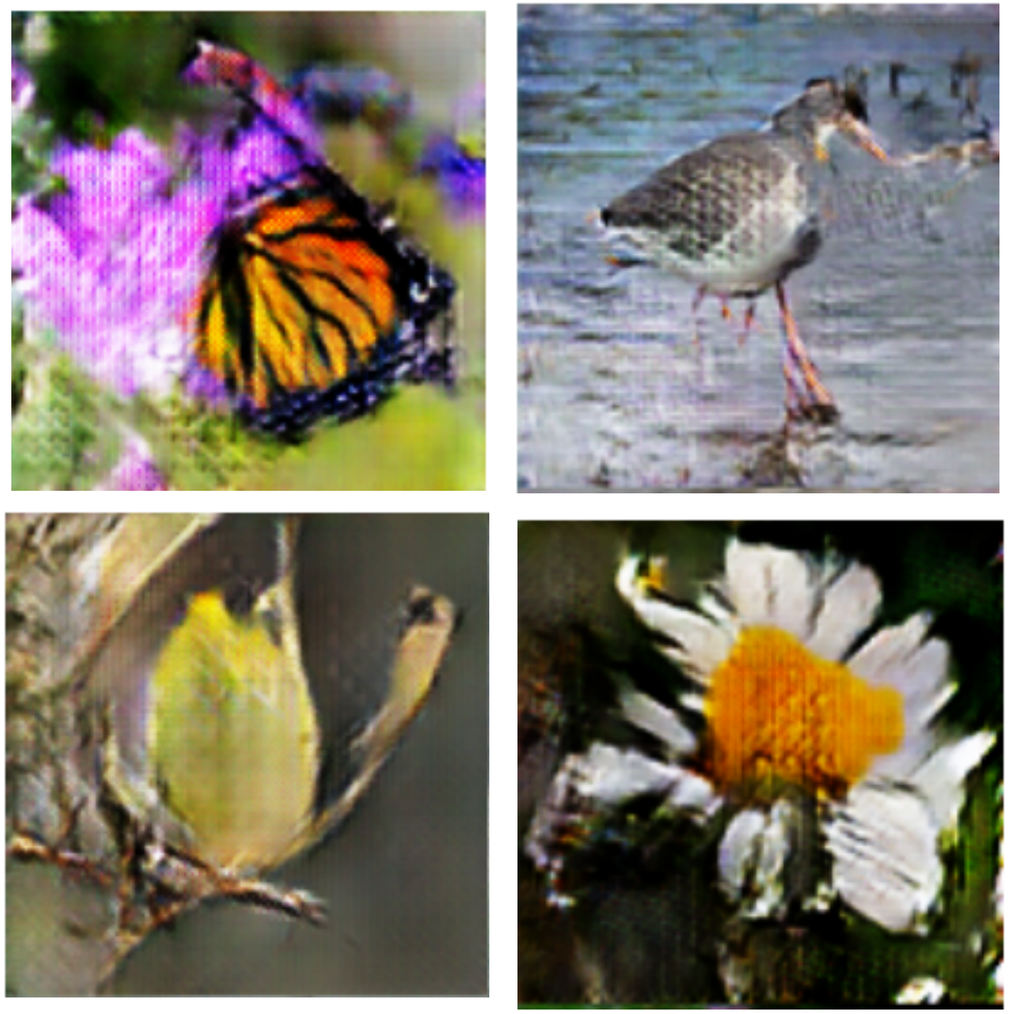

BigGAN [Brock et al. 2018]

Both trained on Imagenet

Architectures

DCGAN
[Radford, Metz, Chintala 2016]

StyleGAN
[Karras, Laine, Aila 2019]

Architectures

DCGAN
[Radford, Metz, Chintala 2016]

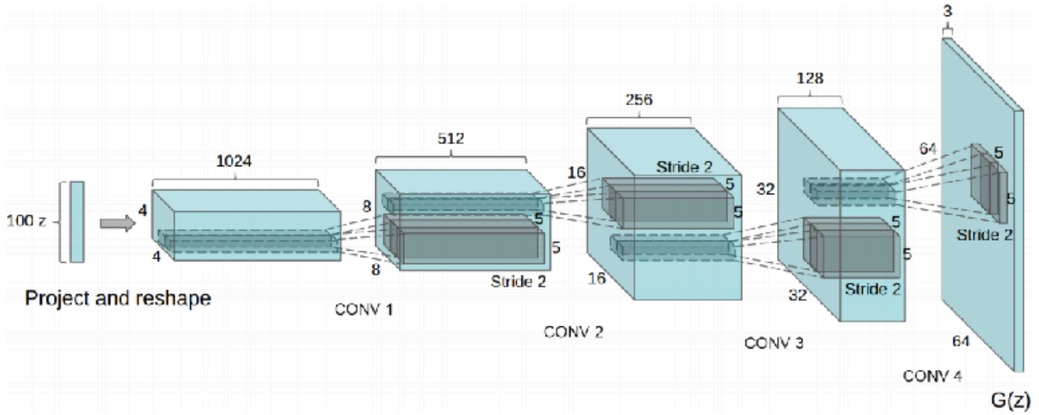

StyleGAN
[Karras, Laine, Aila 2019]

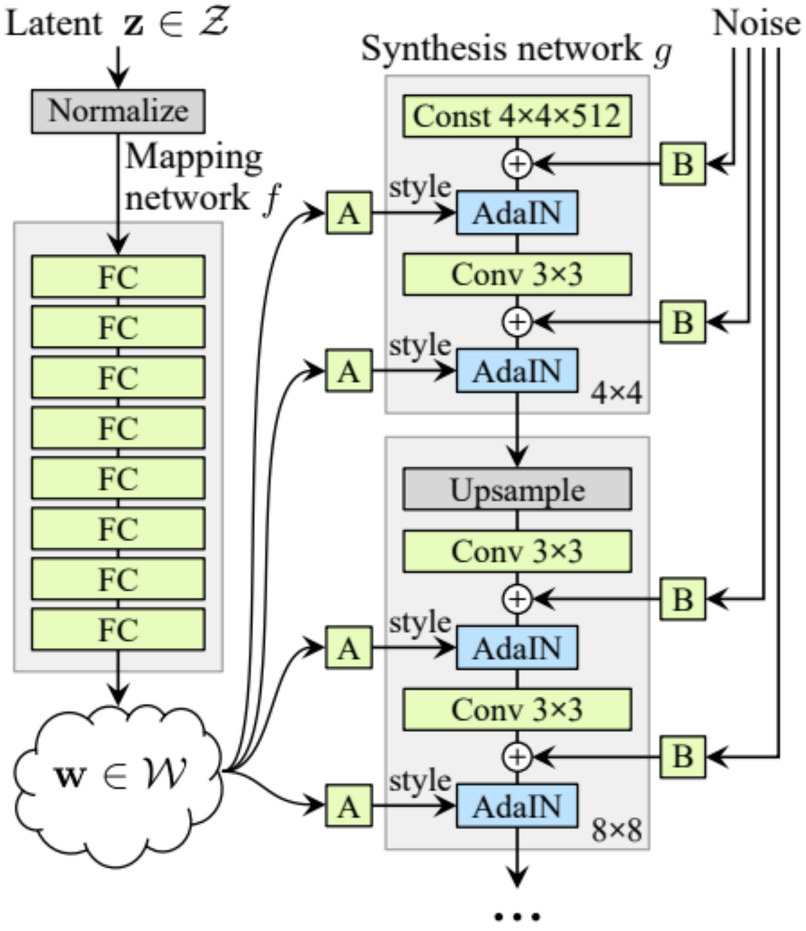

What has driven GAN progress?

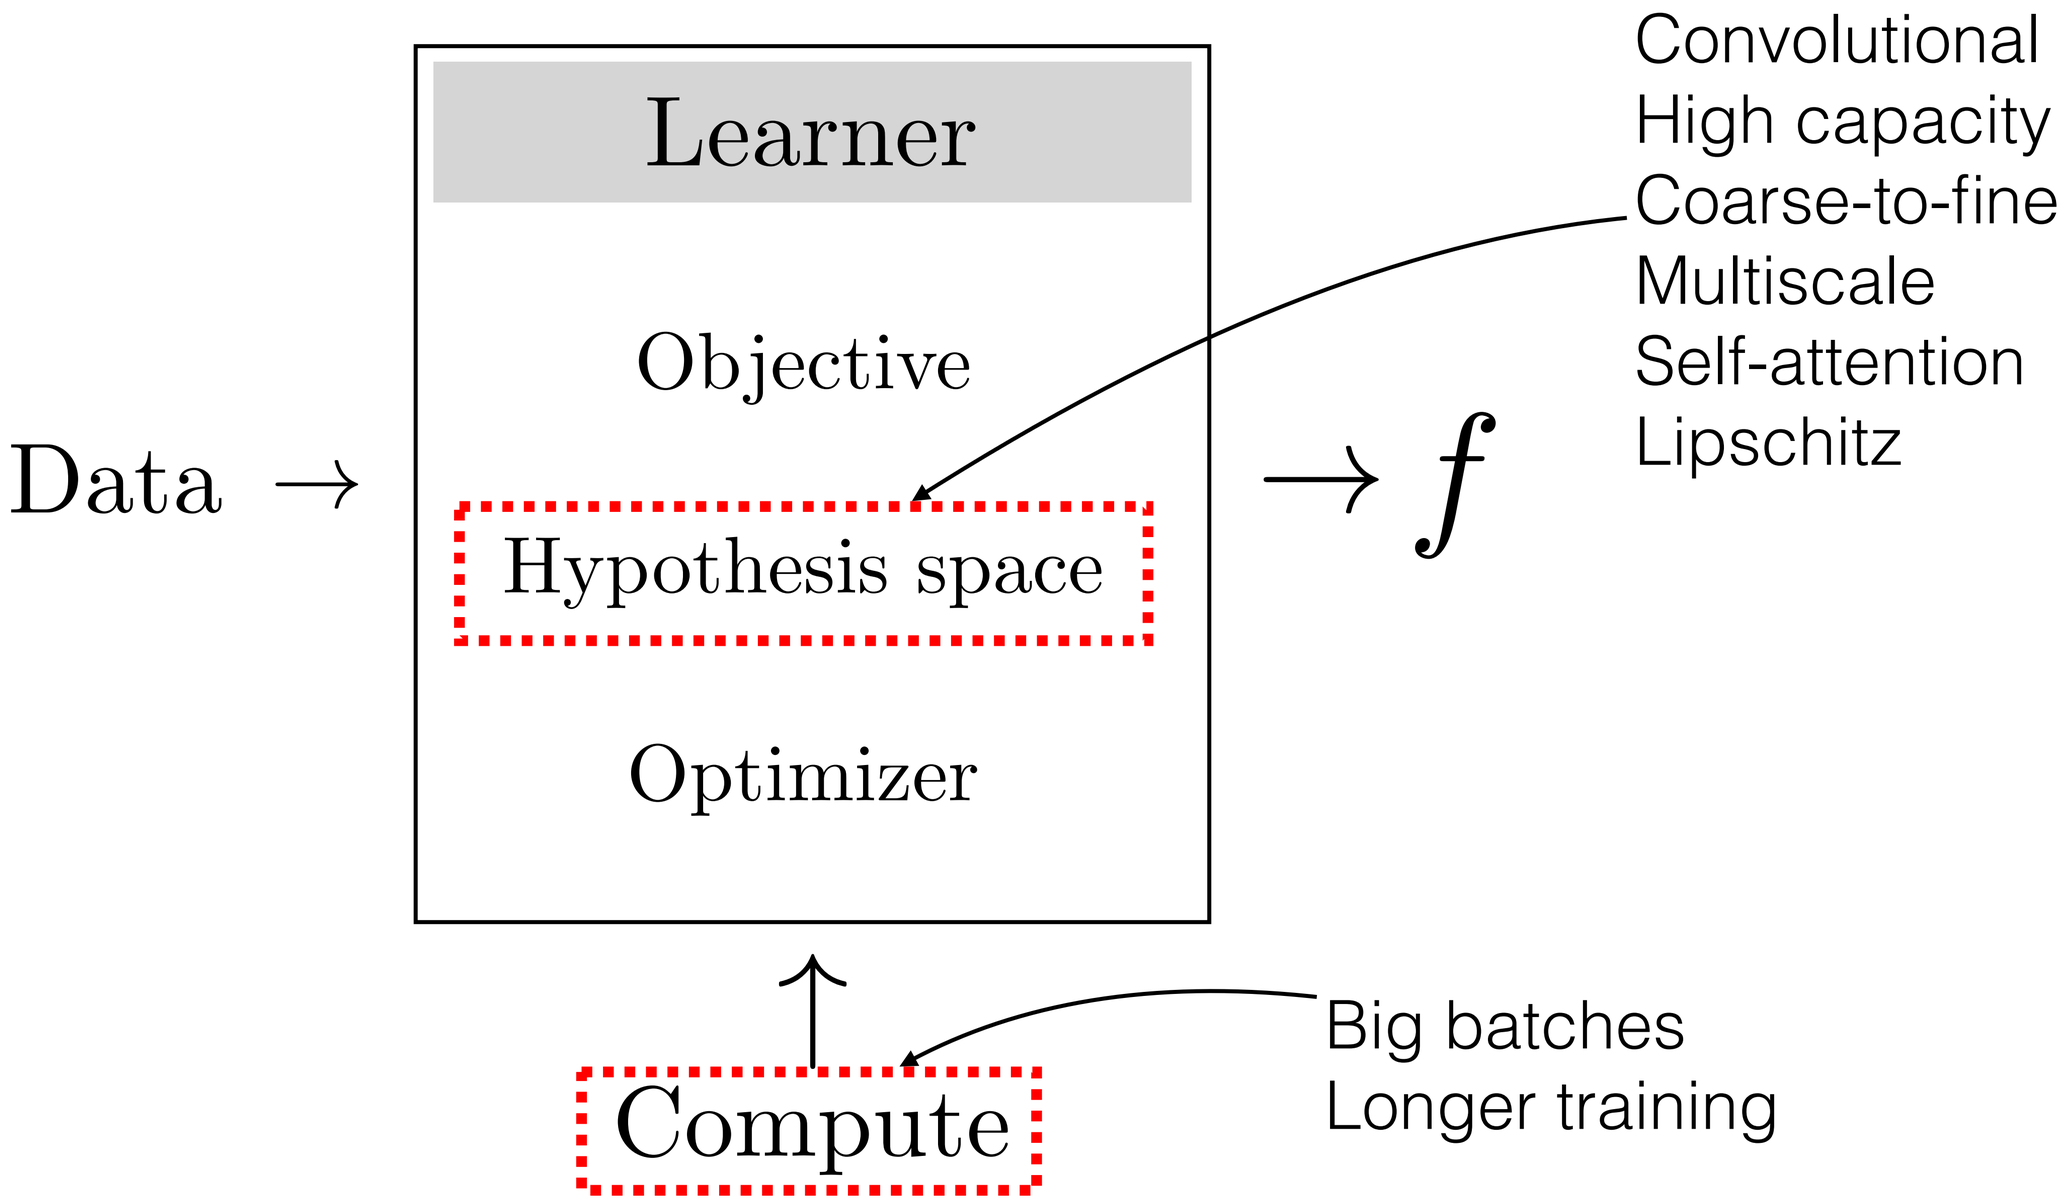

GANs

Pros: Cheap to sample, fast to train, require little data

Cons: No likelihoods, bad coverage (mode collapse), finicky to train (minimax)

Other deep generative models:
Autoregressive models, Normalizing flows, **Energy-based models**

[adapted from slide by David Duvenaud]

1. Image synthesis
2. **Structured prediction**
3. Domain mapping

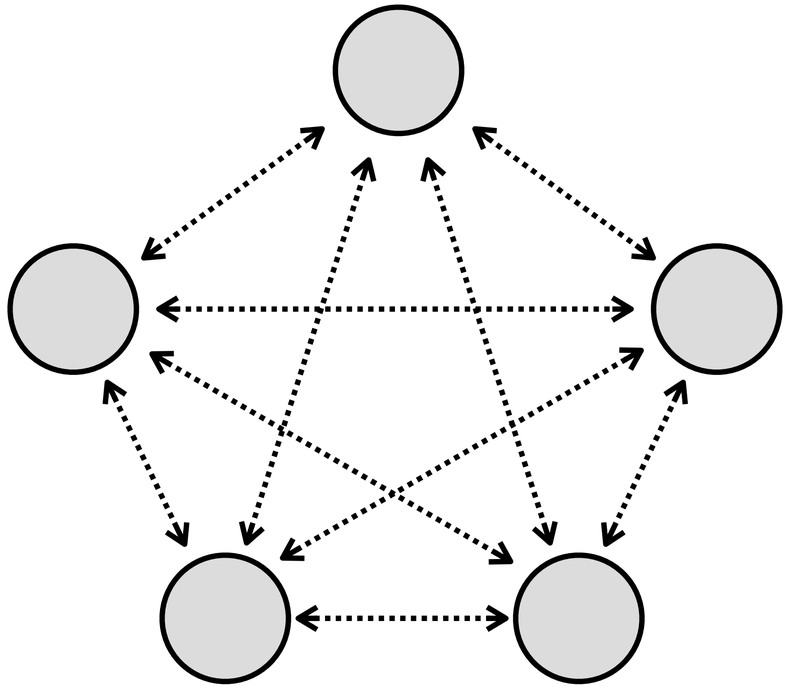

Strutured Prediction

Data prediction problems ("structured prediction")

Semantic segmentation

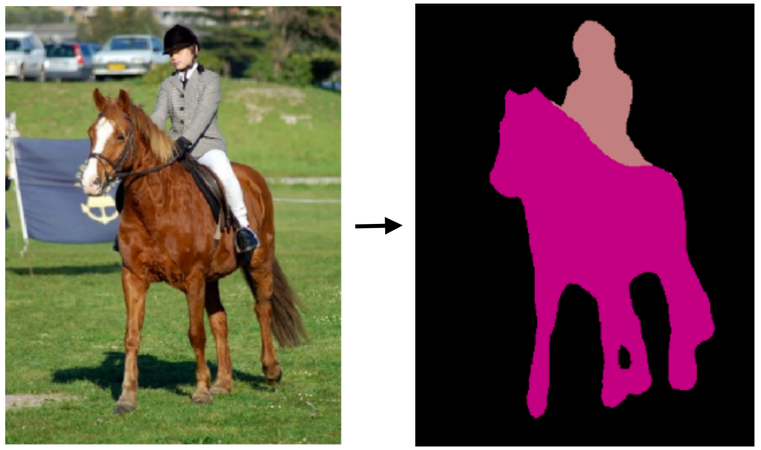

[Long et al. 2015, …]

Edge detection

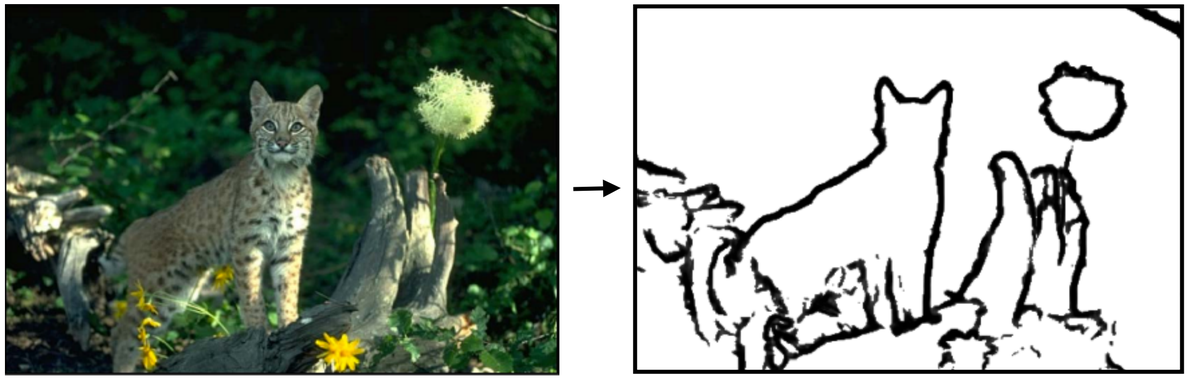

[Xie et al. 2015, …]

Text-to-photo

"this small bird has a pink breast and crown…"

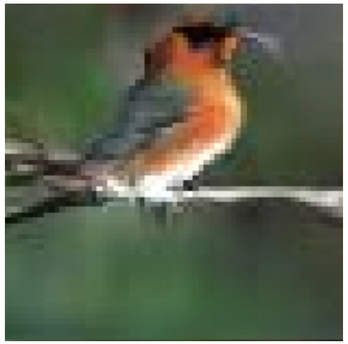

[Reed et al. 2014, …]

Future frame prediction

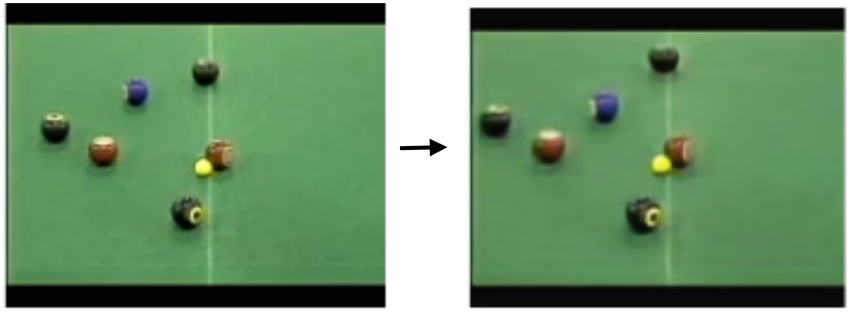

[Mathieu et al. 2016, …]

Structured prediction

X is high-dimensional

Model *joint* distribution of high-dimensional data $P(\mathbf{X}|\mathbf{Y}=\mathbf{y})$

In vision this is usually what we are interested in

Unstructured: $\prod_i p(X_i|\mathbf{Y}=\mathbf{y})$

Deep learning in 2012

Use a **hypothesis space** that can model complex structure
(e.g., a CNN, nearest-neighbor)

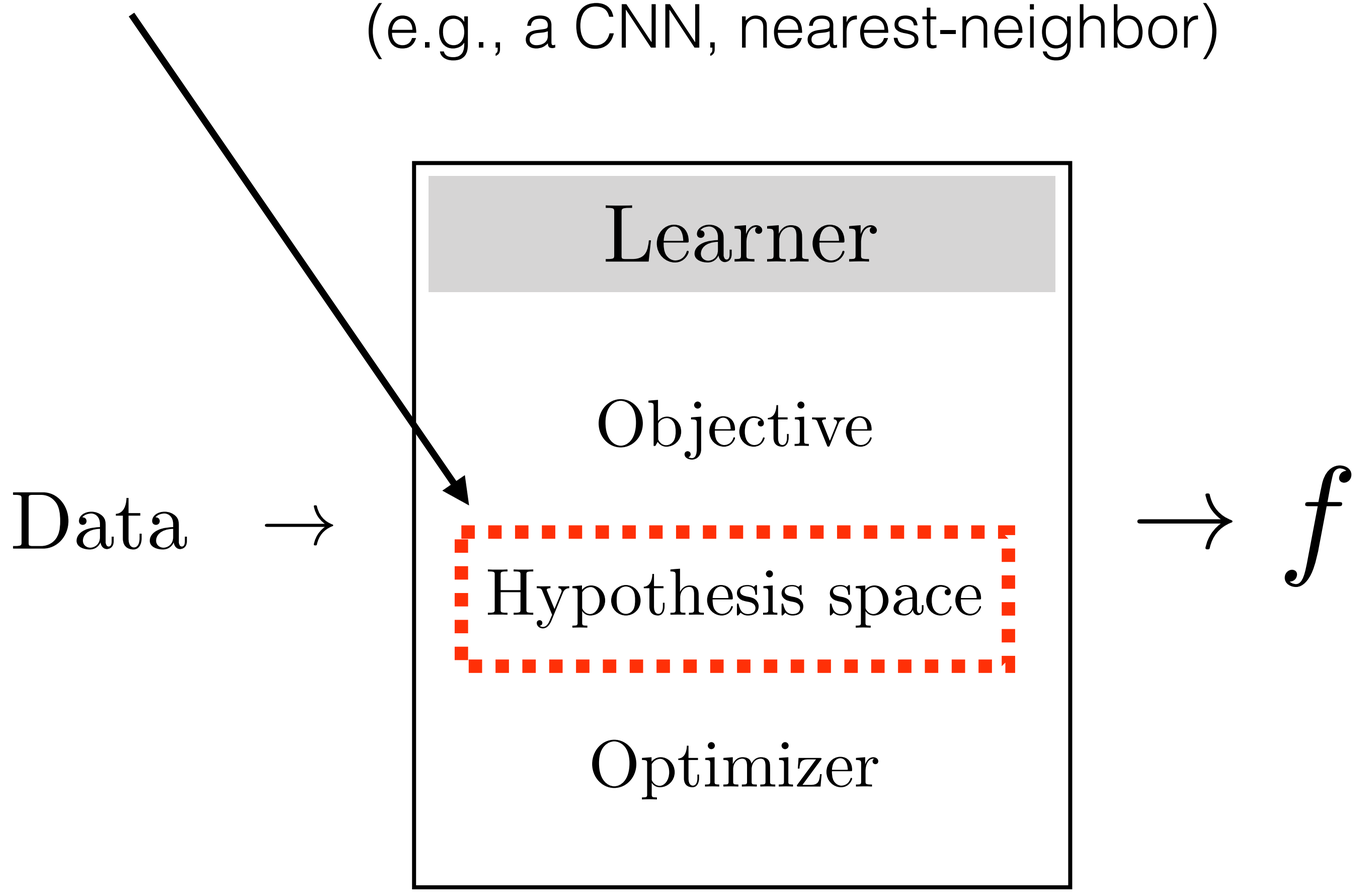

Why deep learning

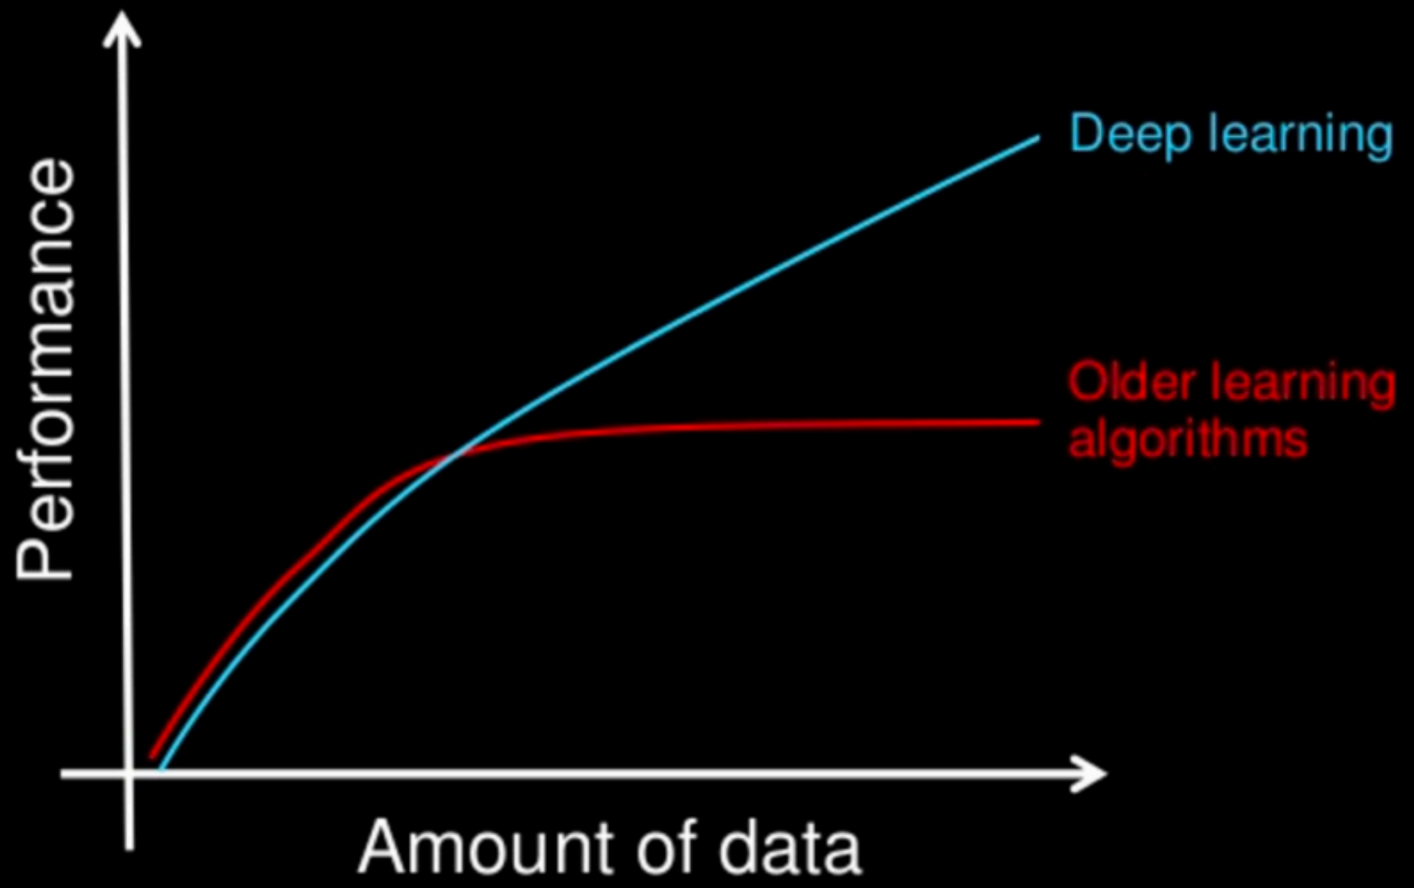

How do data science techniques scale with amount of data?

[Slide credit: Andrew Ng]

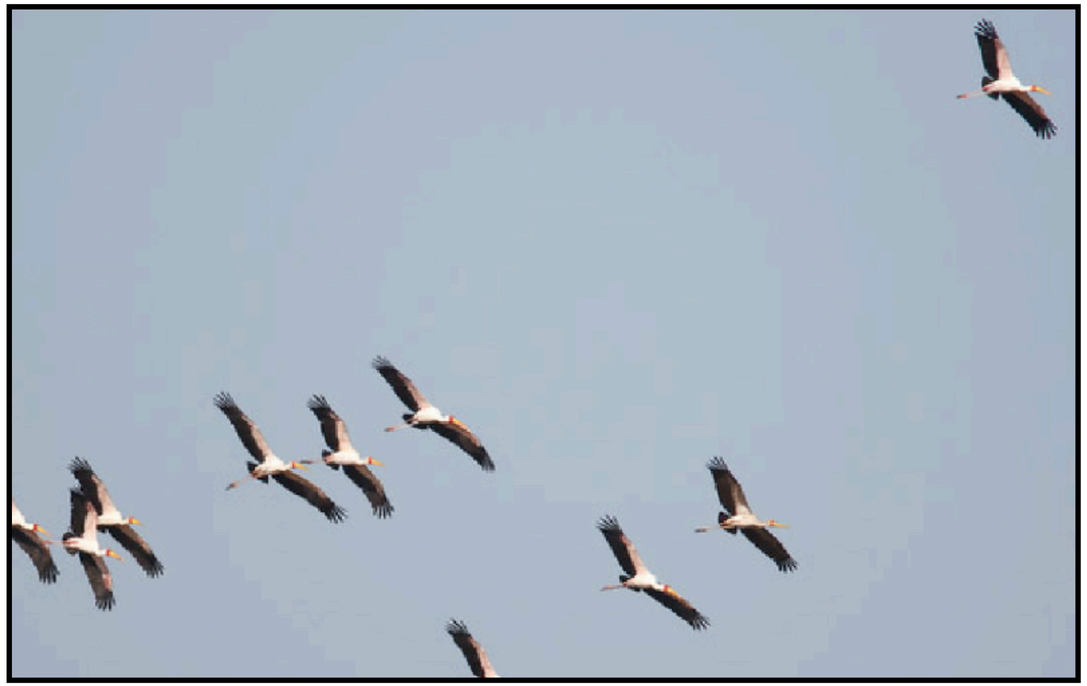

[Photo credit: Fredo Durand]

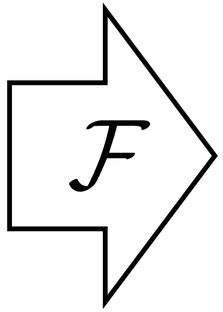

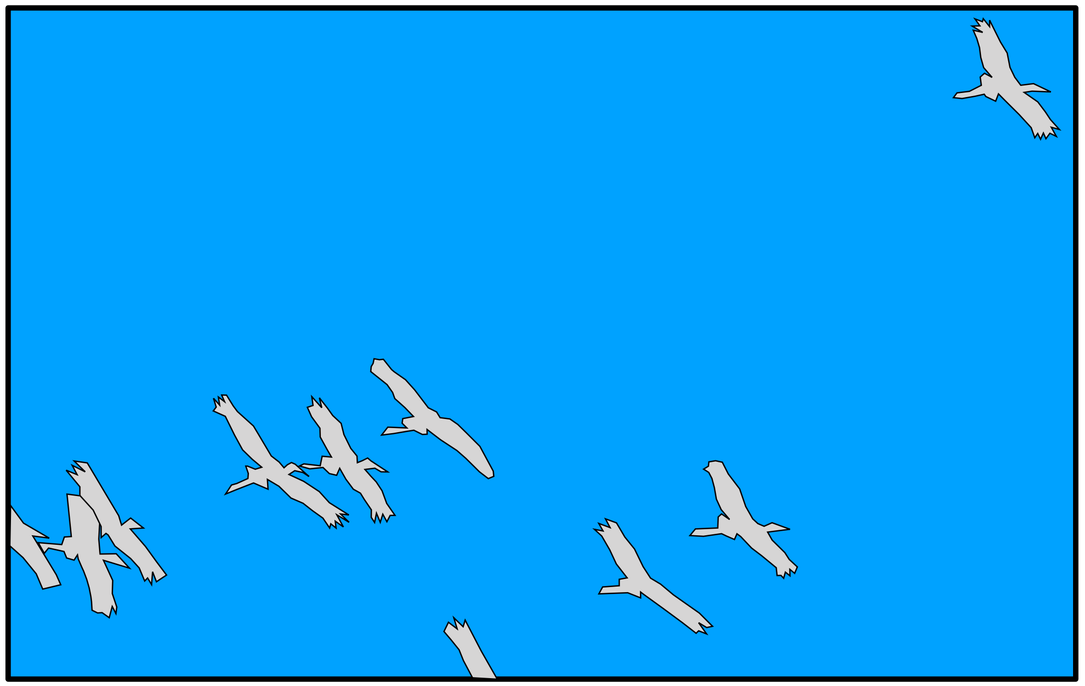

(Colors represent one-hot codes)

$$\arg\min_{\mathcal{F}} \mathbb{E}_{\mathbf{x},\mathbf{y}}[L(\mathcal{F}(\mathbf{x}), \mathbf{y})]$$

Hypothesis space

Objective function
(loss)

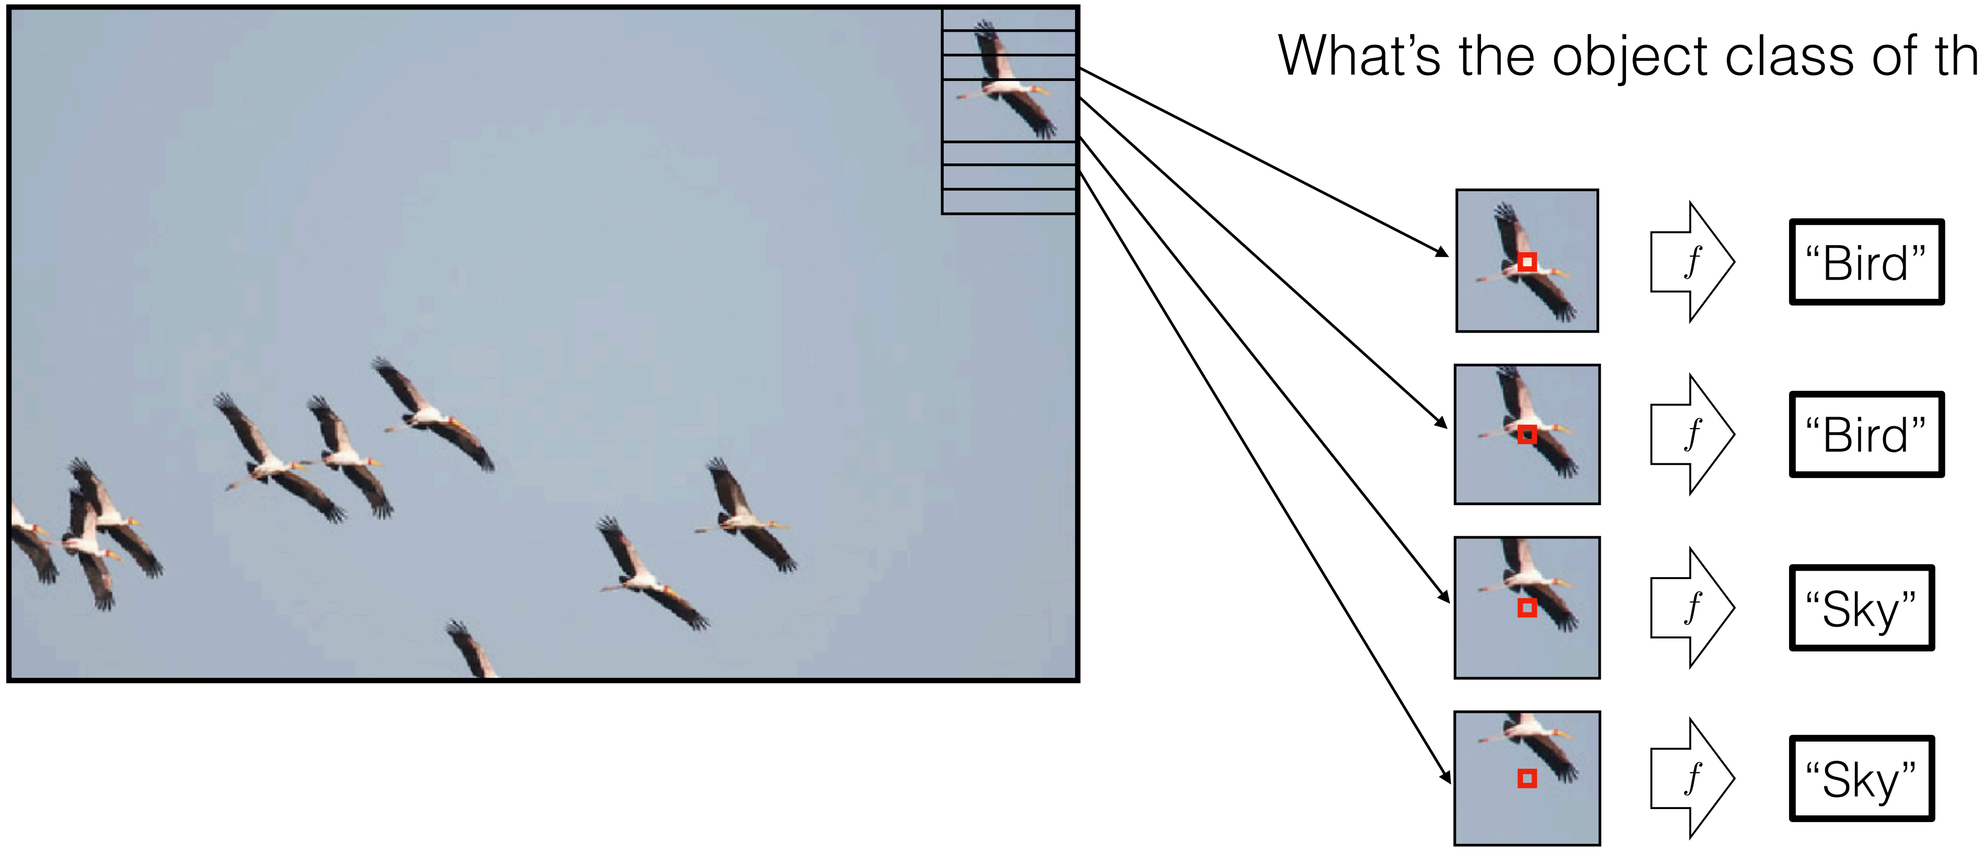

Fully-factored loss: $L(\hat{\mathbf{y}}, \mathbf{y}) = \sum_i \phi_i(\hat{\mathbf{y}}_i, \mathbf{y}_i)$

Semantic Segmentation

Data

$\mathbf{x}$ $\mathbf{y}$

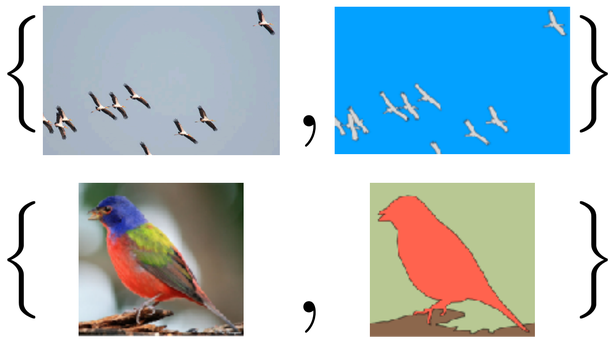

$\vdots$

$\mathbf{x} \in \mathbb{R}^{H \times W \times 3}$

$\mathbf{y} \in \mathbb{R}^{H \times W \times K}$

$\rightarrow$

Learner

Objective

$$f^* = \arg\min_{f \in \mathcal{F}} \sum_{i=1}^{N} H(\mathbf{y}_i, \hat{\mathbf{y}}_i)$$

Hypothesis space

Convolutional neural net

Optimizer

Stochastic gradient descent

$\rightarrow$ f

Sat2Map

Data

$\mathbf{x}$ $\mathbf{y}$

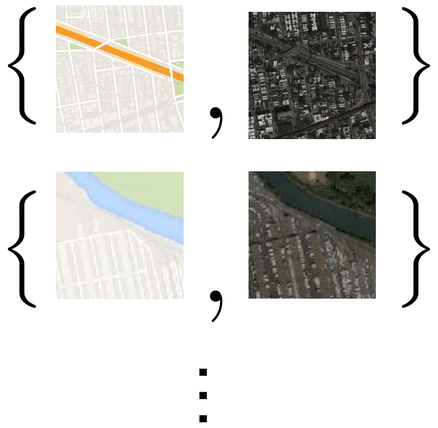

$\mathbf{x} \in \mathbb{R}^{H \times W \times 3}$

$\mathbf{y} \in \mathbb{R}^{H \times W \times 3}$

$\rightarrow$

Learner

Objective

$$\theta^* = \arg\min_{\theta} \sum_{i=1}^{N} (f_\theta(\mathbf{x})_i - y_i)^2$$

Hypothesis space

Convolutional neural net

Optimizer

Stochastic gradient descent

$\rightarrow$ f

Input

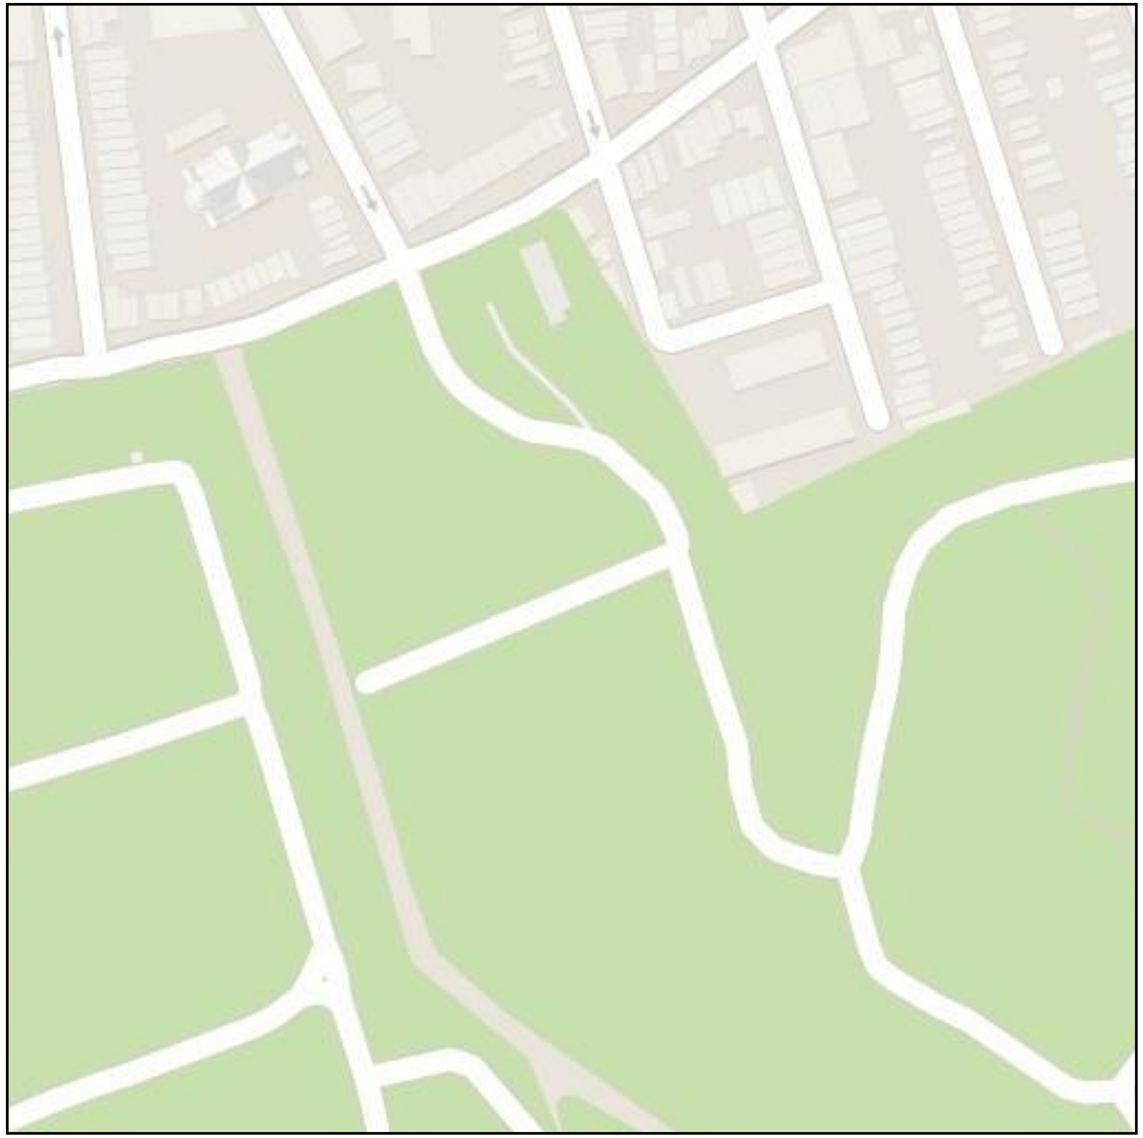

Deep net output

Input

Independent prediction per-pixel

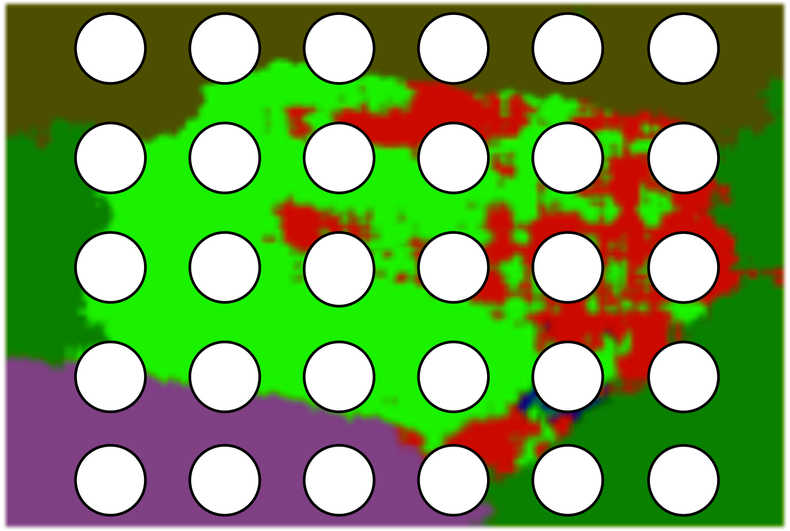

$$\max \prod_i p(y_i|\mathbf{x})$$

Find a configuration of compatible labels

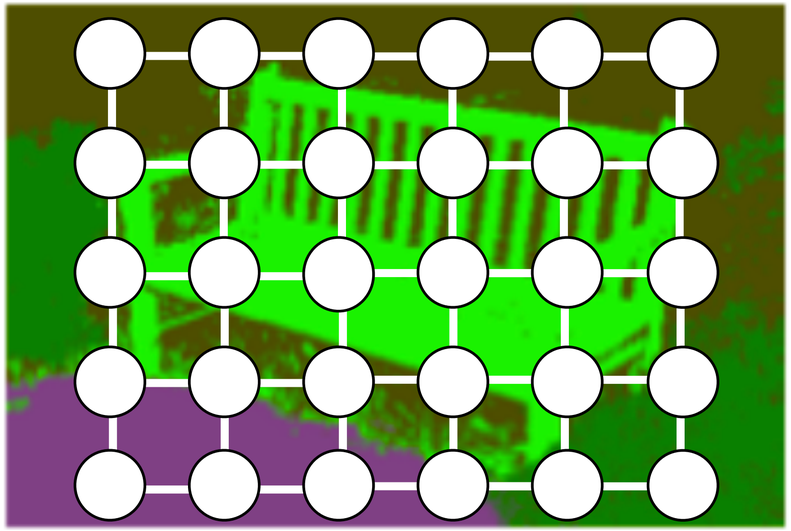

$$\max \frac{1}{Z} \prod_{i,j} p(y_i, y_j|\mathbf{x})$$

["Efficient Inference in Fully Connected CRFs with Gaussian Edge Potentials", Krahenbuhl and Koltun, NeurIPS 2011]

Structured prediction

Use an **objective** that can model structure! (e.g., a graphical model, a GAN, etc)

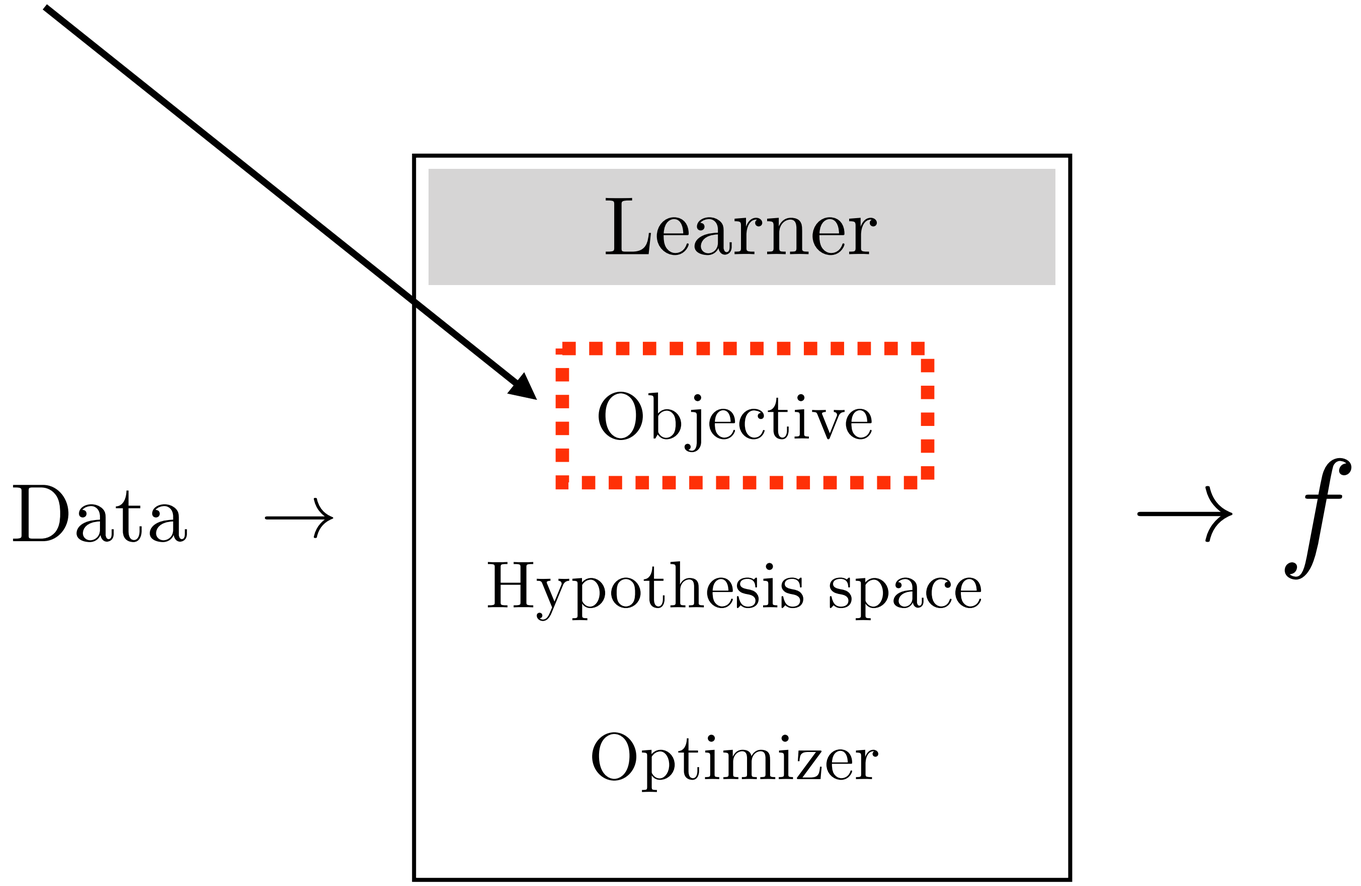

$\mathbf{x}$

G

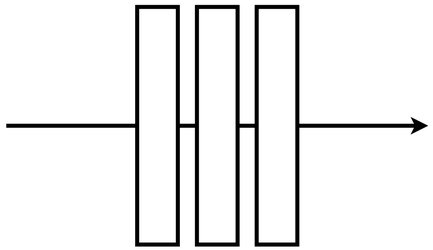

Generator

$G(\mathbf{x})$

G tries to synthesize fake images that fool **D**

D tries to identify the fakes

$$\arg\max_{D} \; \mathbb{E}_{\mathbf{x},\mathbf{y}}[\; \log D(G(\mathbf{x})) \; + \; \log(1 - D(\mathbf{y}))\;]$$

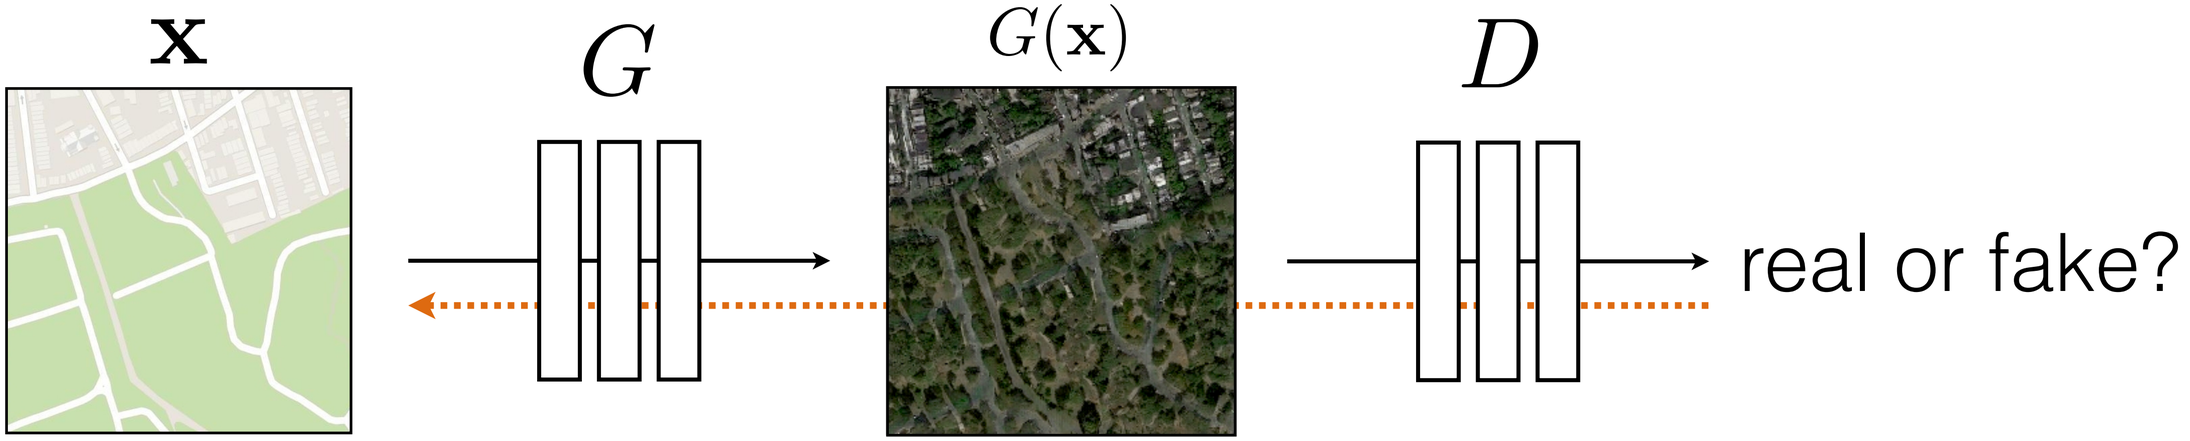

G tries to synthesize fake images that ***fool*** **D**:

$$\arg\min_{G} \mathbb{E}_{\mathbf{x},\mathbf{y}}[\ \log D(G(\mathbf{x})) \quad + \quad \log(1 - D(\mathbf{y}))\ ]$$

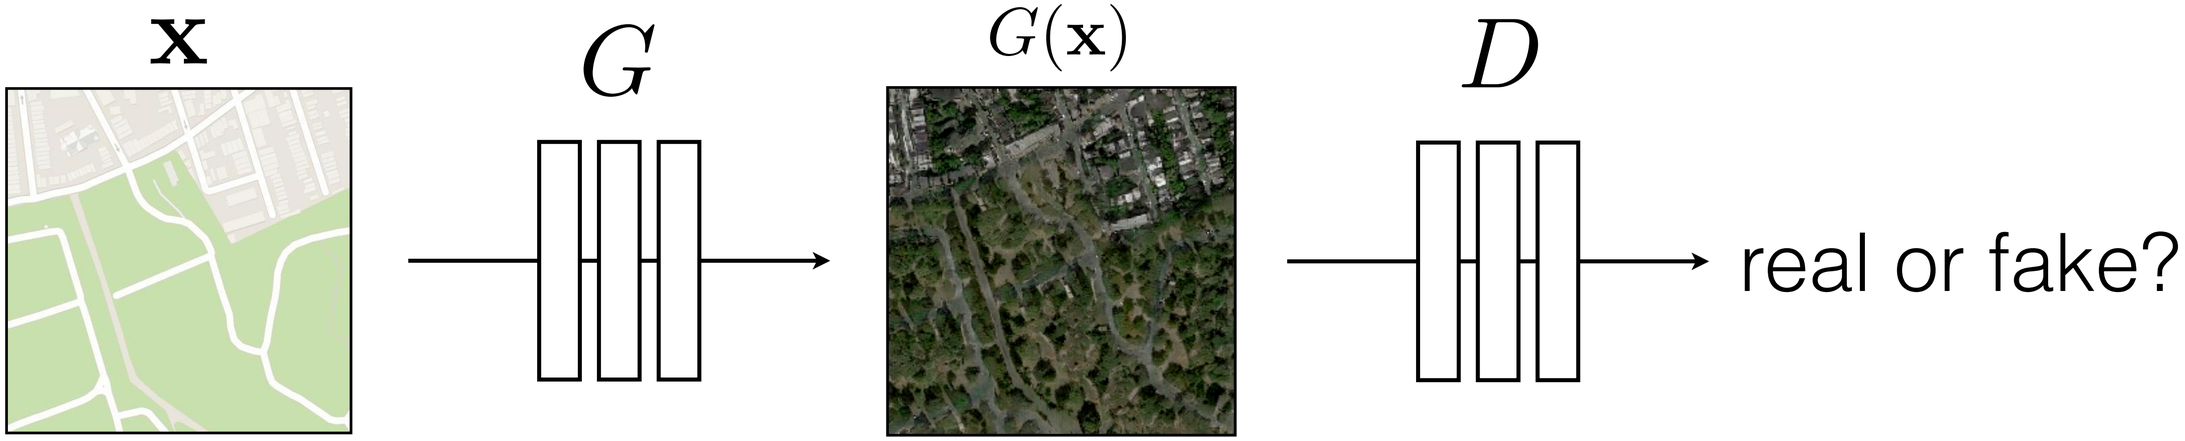

G tries to synthesize fake images that ***fool*** the ***best*** **D**:

$$\arg \min_G \max_D \; \mathbb{E}_{\mathbf{x},\mathbf{y}}[\; \log D(G(\mathbf{x})) \quad + \quad \log(1 - D(\mathbf{y})) \;\;]$$

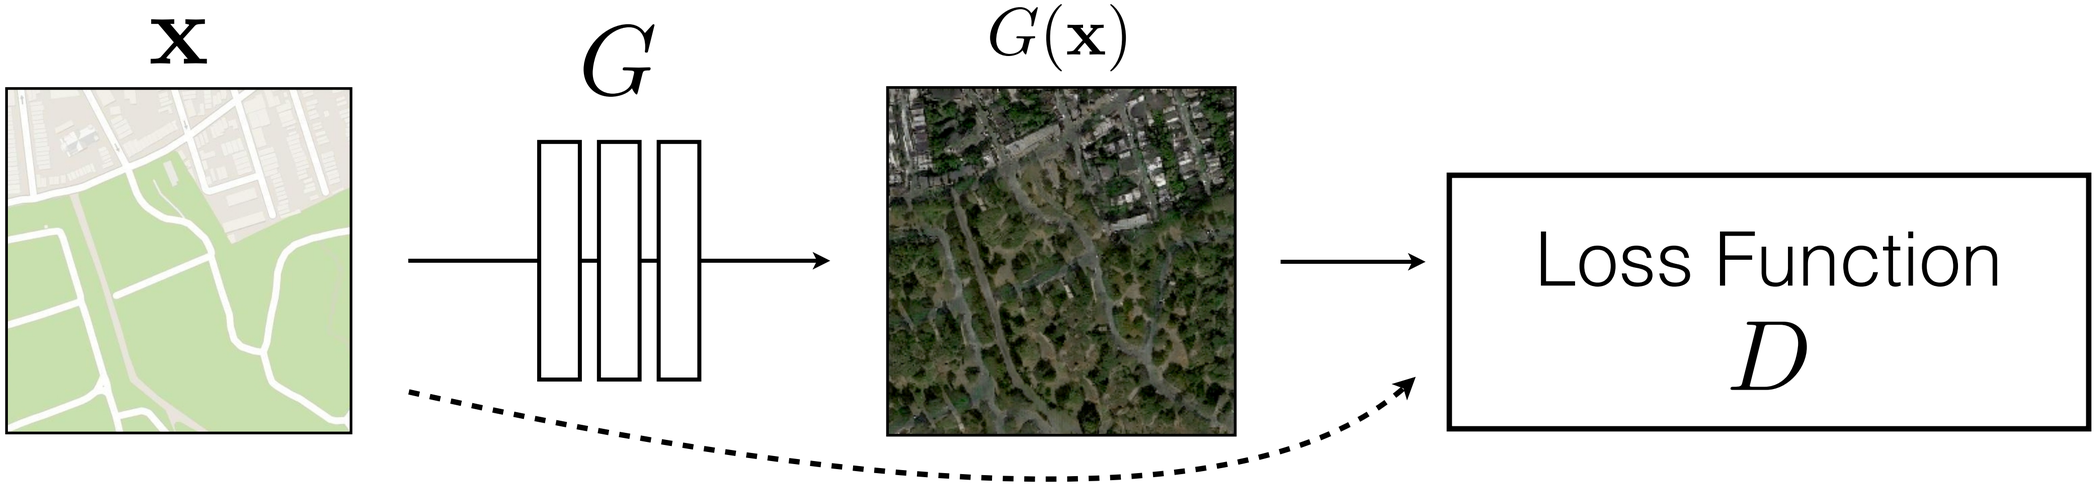

G's perspective: **D** is a loss function.

Rather than being hand-designed, it is *learned* and *highly structured*.

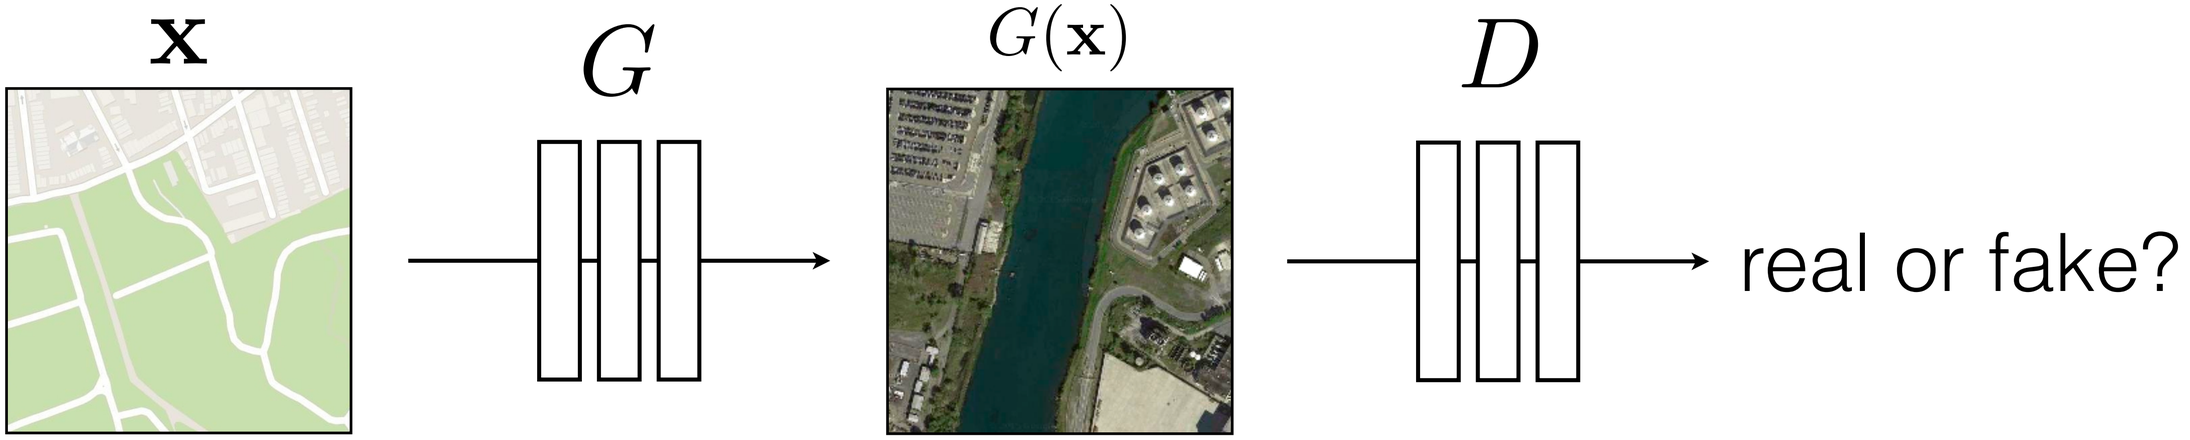

$$\arg\min_G \max_D \; \mathbb{E}_{\mathbf{x},\mathbf{y}}[\; \log D(G(\mathbf{x})) \quad + \quad \log(1 - D(\mathbf{y}))\;]$$

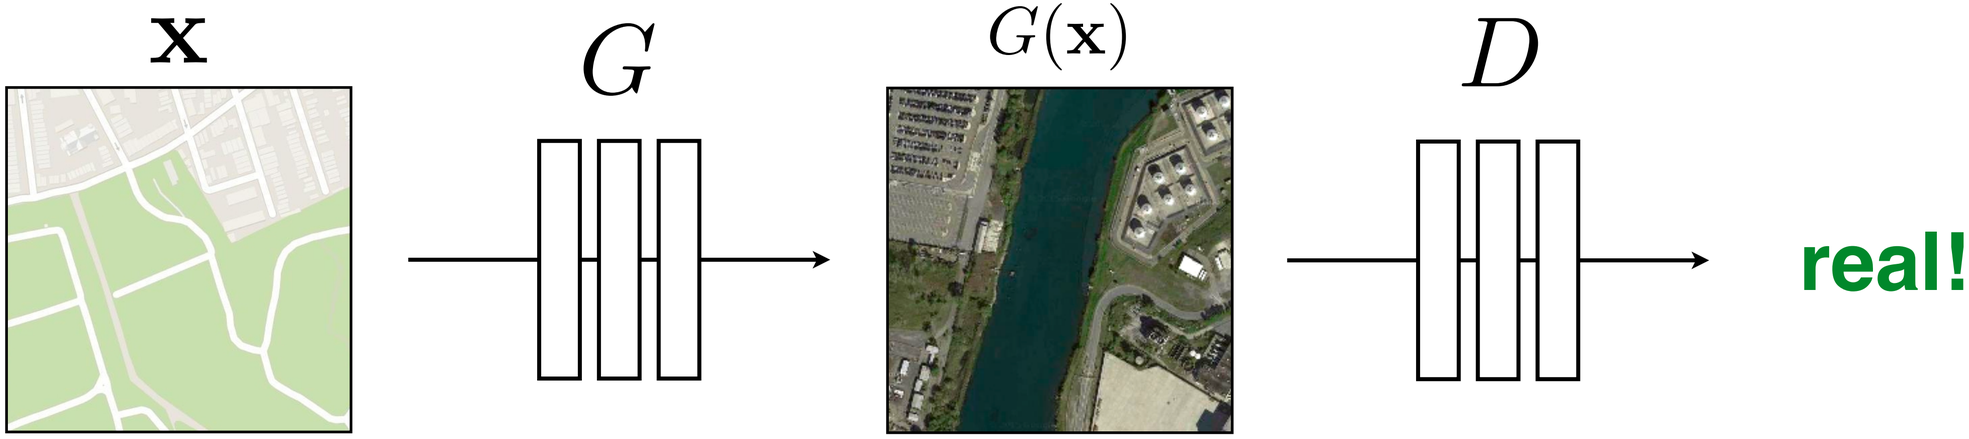

$$\arg\min_G \max_D \; \mathbb{E}_{\mathbf{x},\mathbf{y}}[\; \log D(G(\mathbf{x})) \quad + \quad \log(1 - D(\mathbf{y}))\;]$$

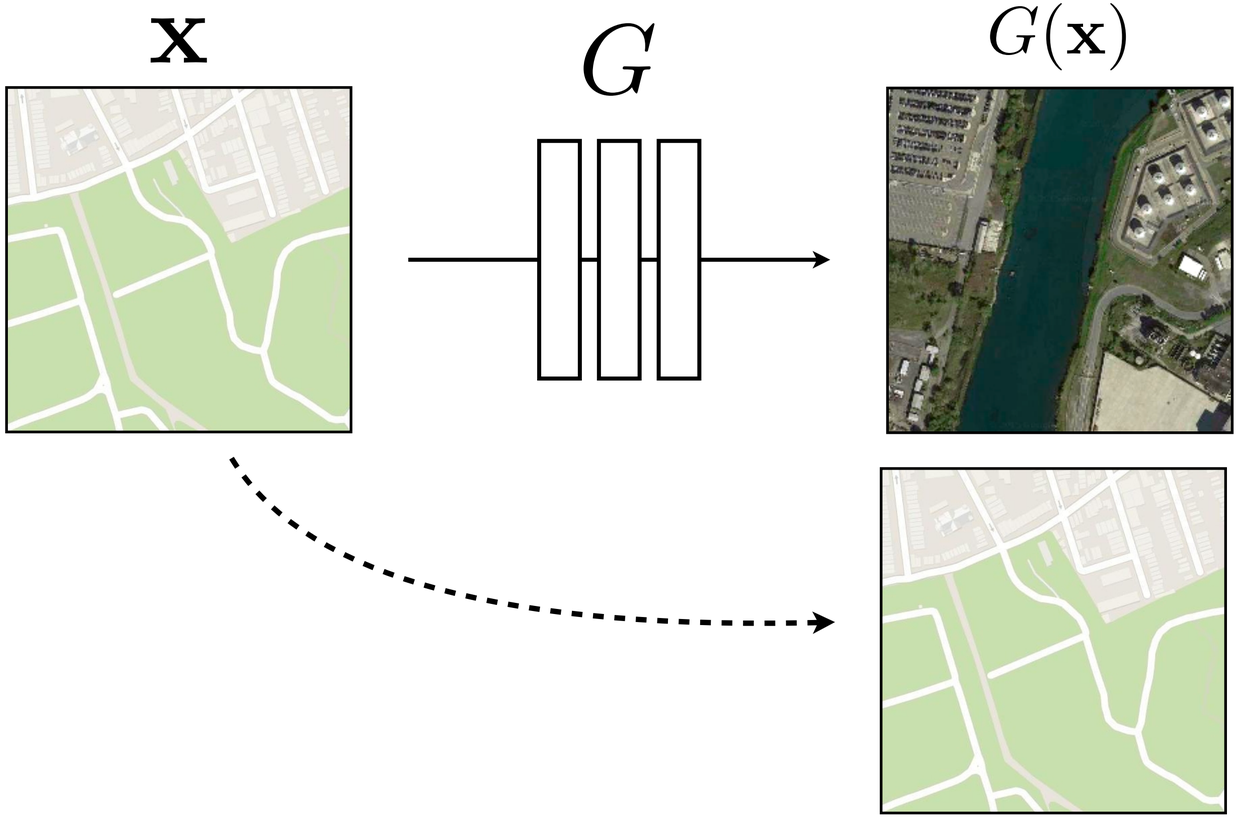

$$\arg\min_G \max_D \; \mathbb{E}_{\mathbf{x},\mathbf{y}}[\; \log D(G(\mathbf{x})) \quad + \quad \log(1 - D(\mathbf{y}))\;]$$

$\mathbf{x}$

G

$G(\mathbf{x})$

D

real or fake *pair* ?

$$\arg\min_G \max_D \; \mathbb{E}_{\mathbf{x},\mathbf{y}}[\; \log D(\mathbf{x}, G(\mathbf{x})) + \log(1 - D(\mathbf{x}, \mathbf{y}))\;]$$

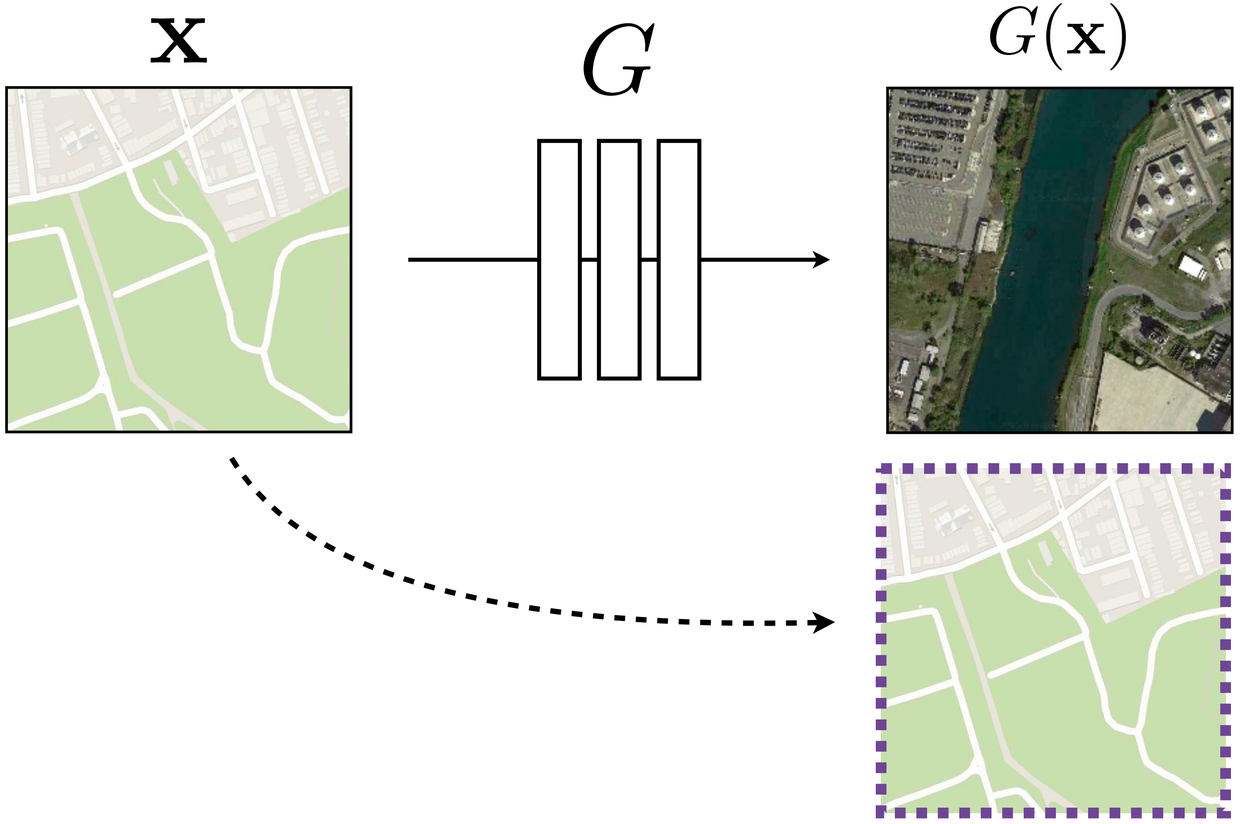

$$\arg\min_G \max_D \; \mathbb{E}_{\mathbf{x},\mathbf{y}}[\; \log D(\mathbf{x}, G(\mathbf{x})) + \log(1 - D(\mathbf{x}, \mathbf{y}))\;]$$

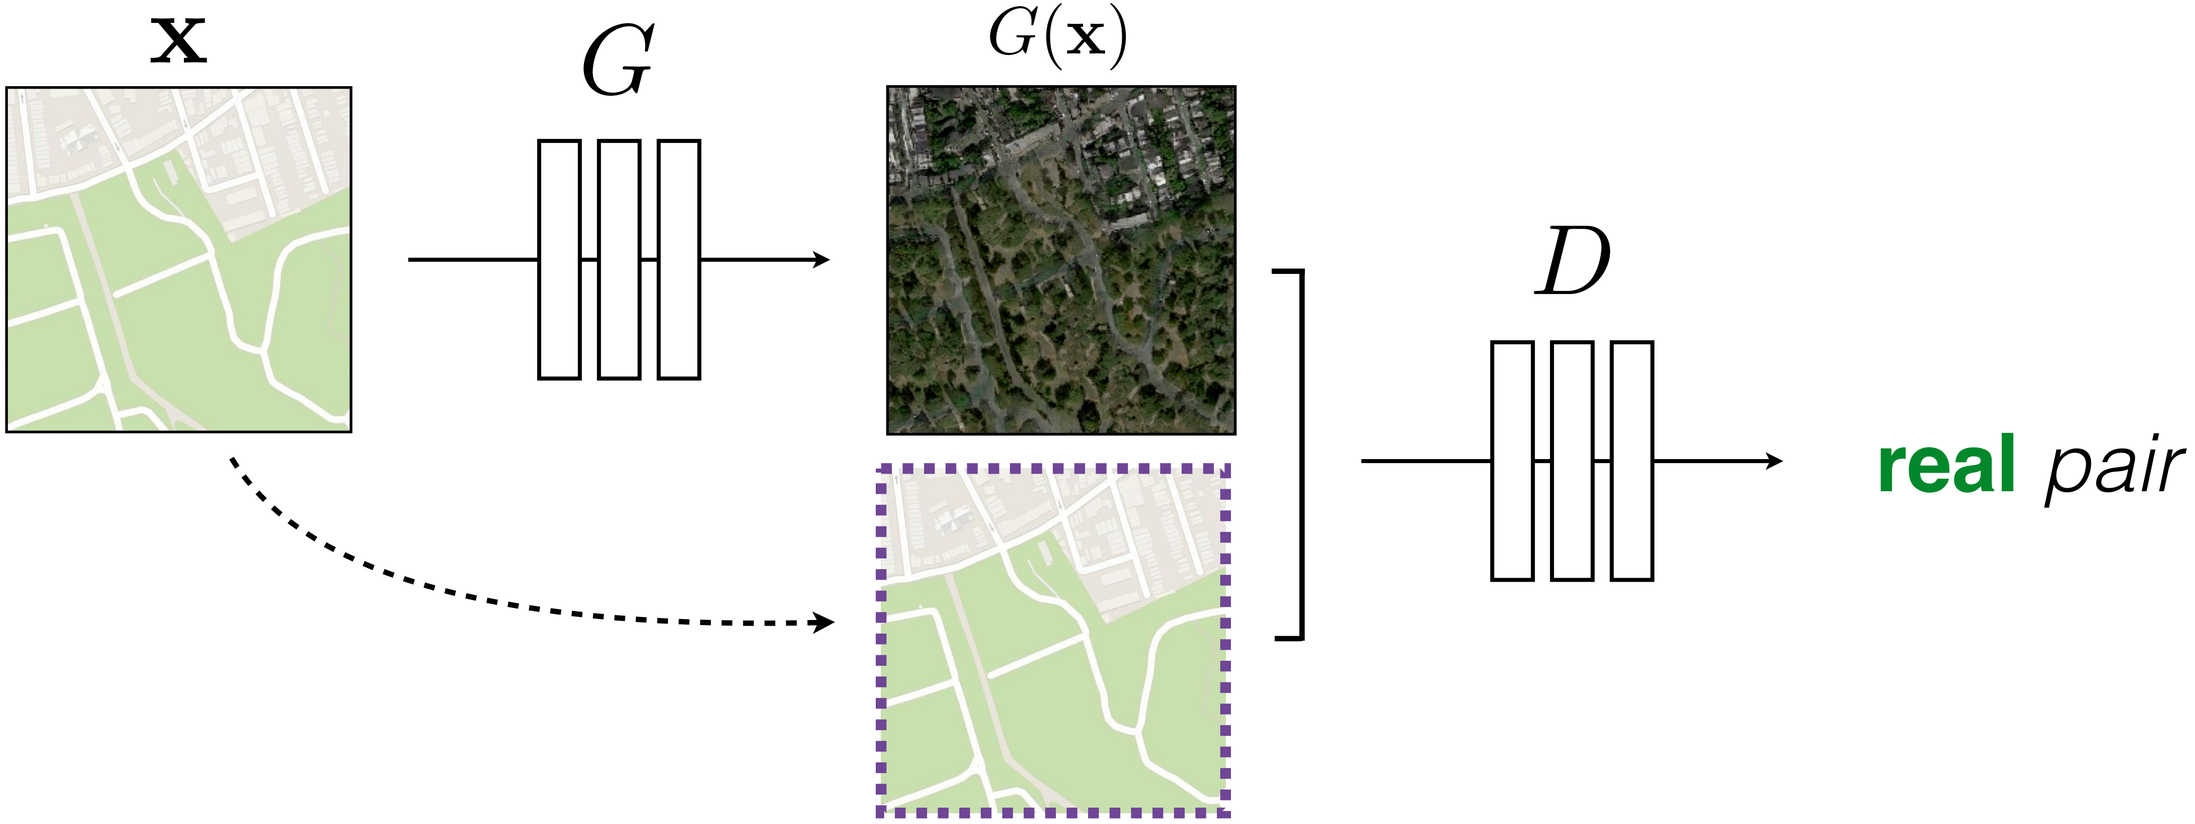

$$\arg\min_G \max_D \; \mathbb{E}_{\mathbf{x},\mathbf{y}}[\; \log D(\mathbf{x}, G(\mathbf{x})) + \log(1 - D(\mathbf{x}, \mathbf{y}))\;]$$

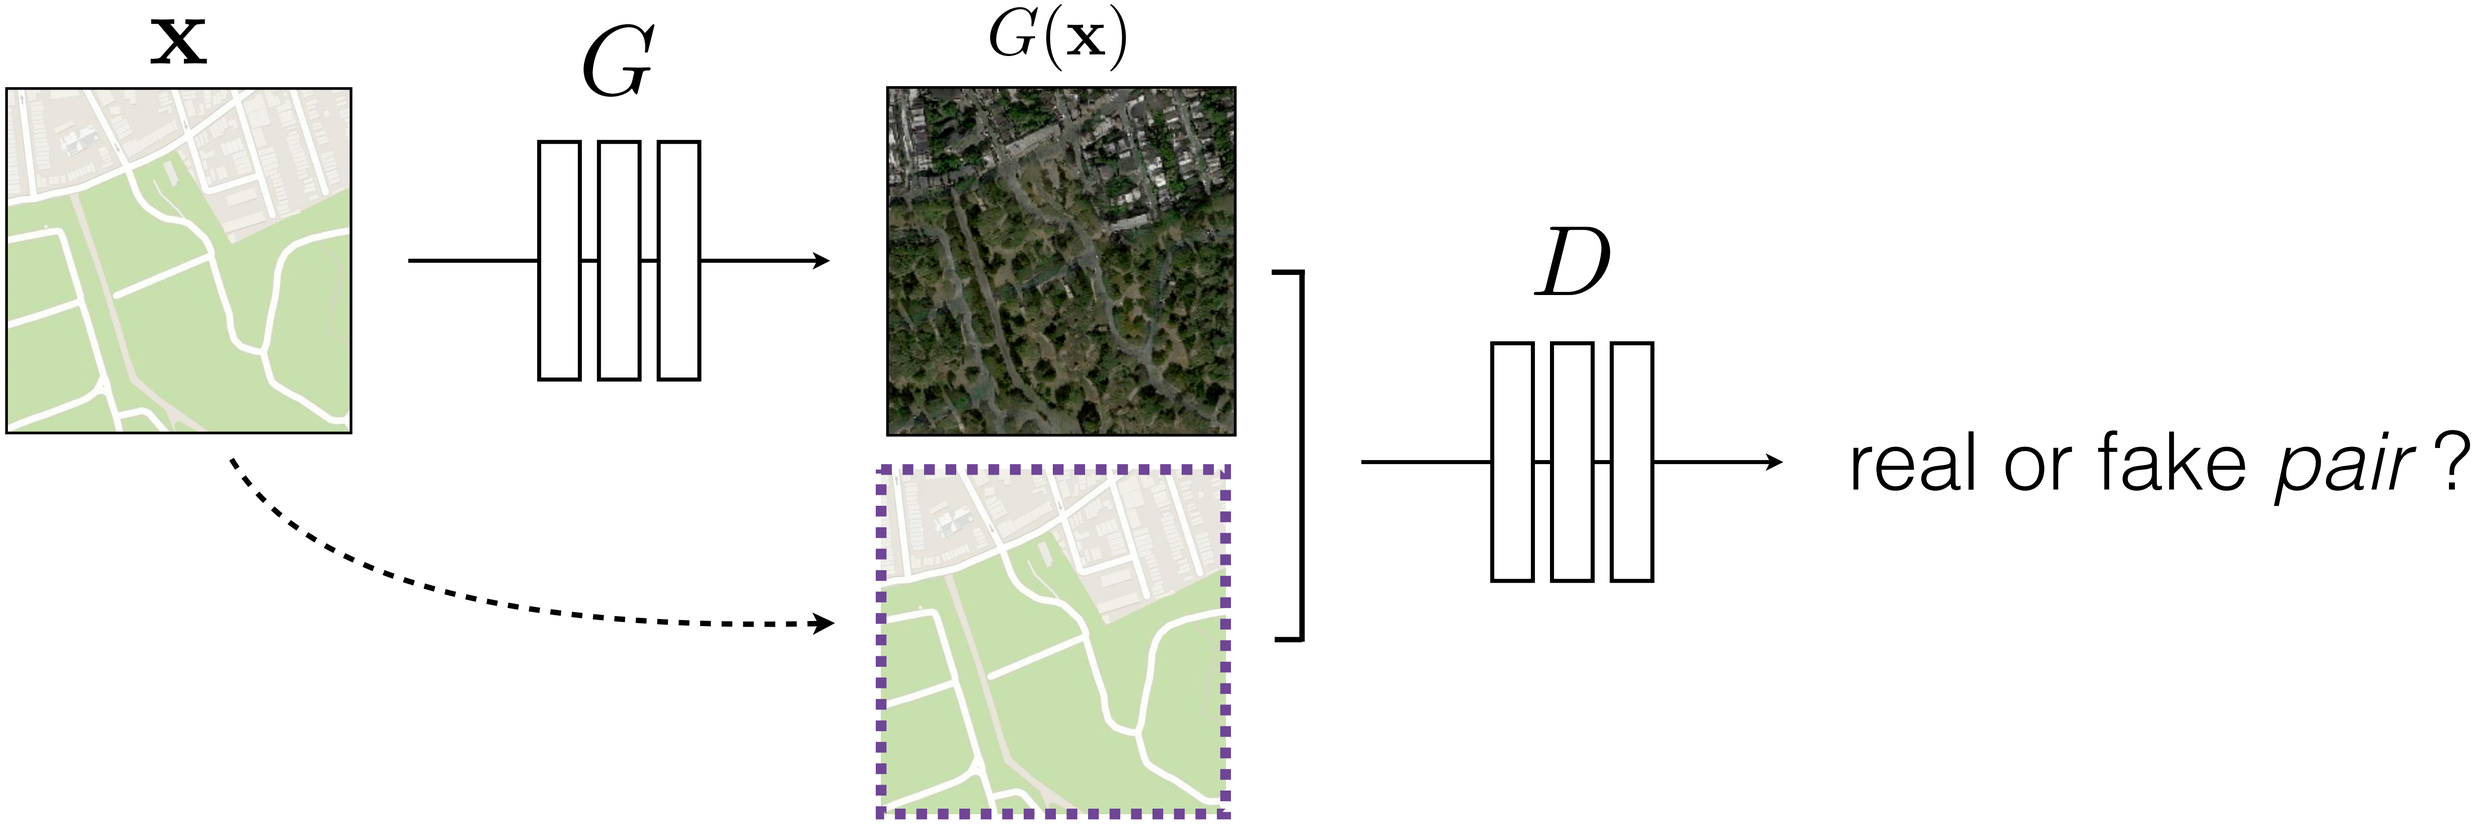

$$\arg\min_G \max_D \; \mathbb{E}_{\mathbf{x},\mathbf{y}}[\;\; \log D(\mathbf{x}, G(\mathbf{x})) + \log(1 - D(\mathbf{x}, \mathbf{y}))\;\;]$$

Training Details: Loss function

Conditional GAN

$$G^* = \arg\min_G \max_D \mathcal{L}_{cGAN}(G, D) + \lambda \mathcal{L}_{L1}(G).$$

Training Details: Loss function

Conditional GAN

$$G^* = \arg\min_G \max_D \mathcal{L}_{cGAN}(G, D) + \lambda \mathcal{L}_{L1}(G).$$

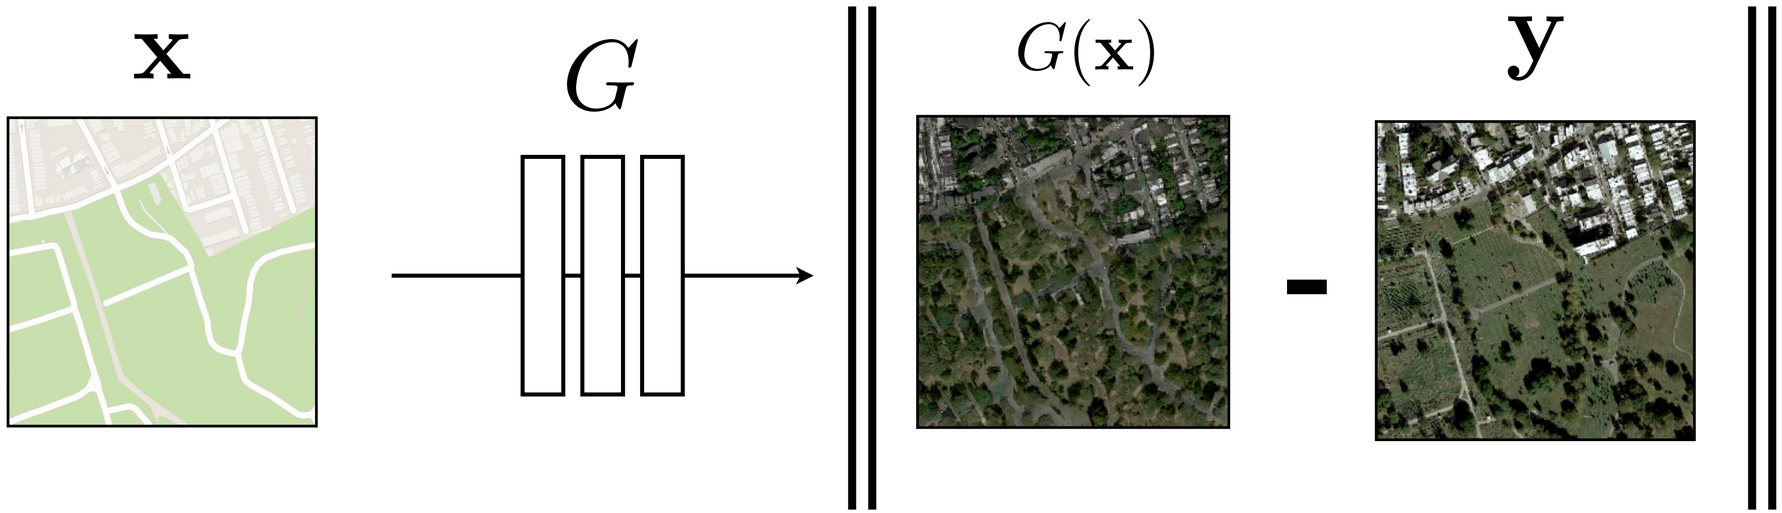

Stable training + fast convergence

[c.f. Pathak et al. CVPR 2016]

Input

Output

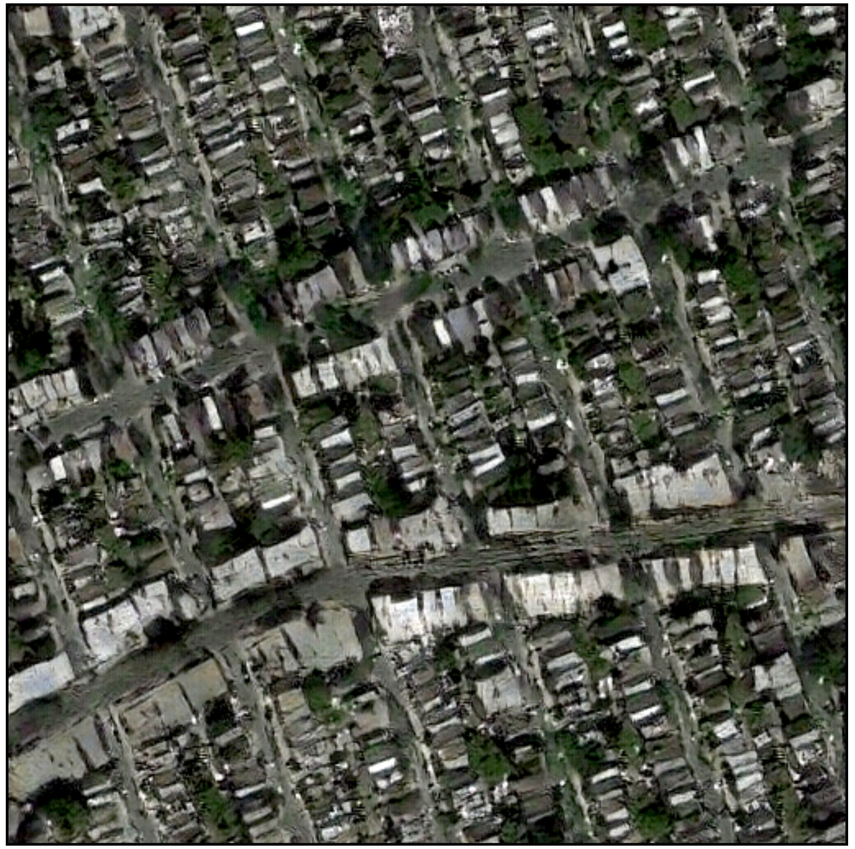

Groundtruth

Data from
[maps.google.com]

Input

Output

Groundtruth

Data from [maps.google.com]

Why deep learning

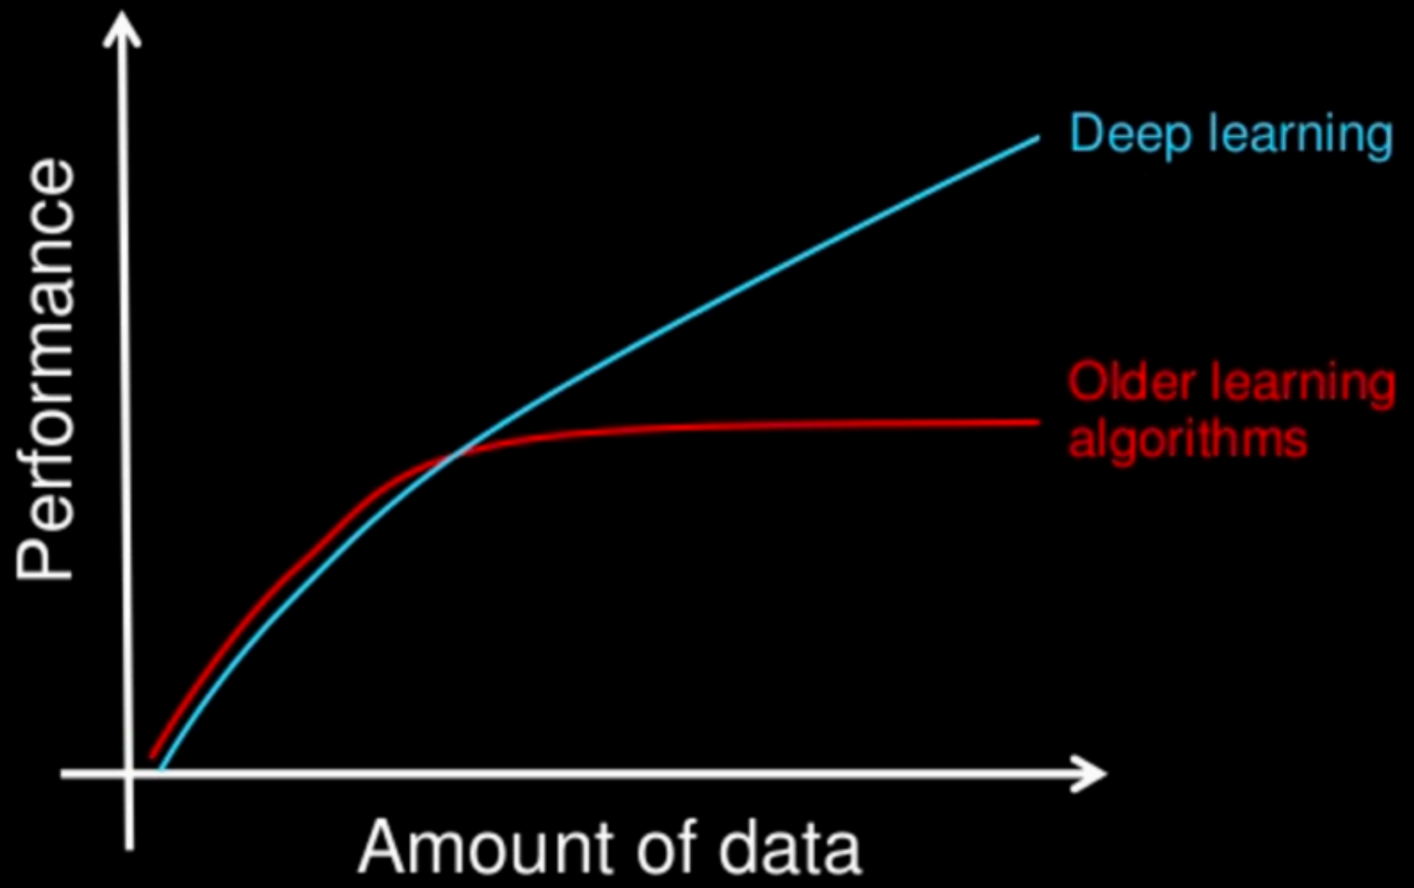

How do data science techniques scale with amount of data?

[Slide credit: Andrew Ng]

Why structured objectives

(cartoon)

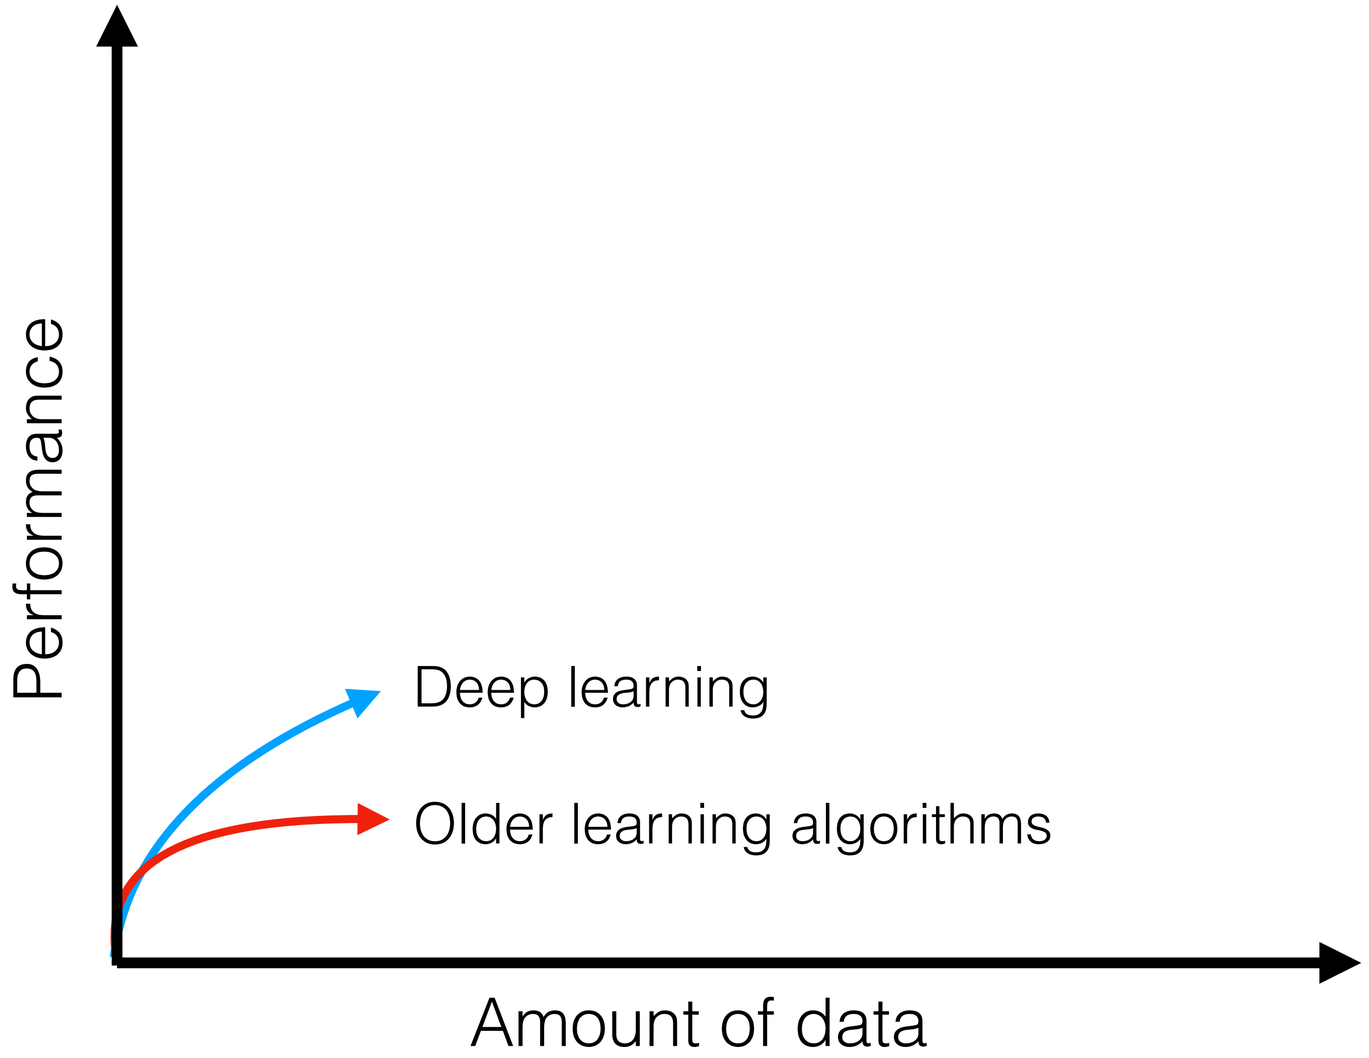

Why structured objectives

(cartoon)

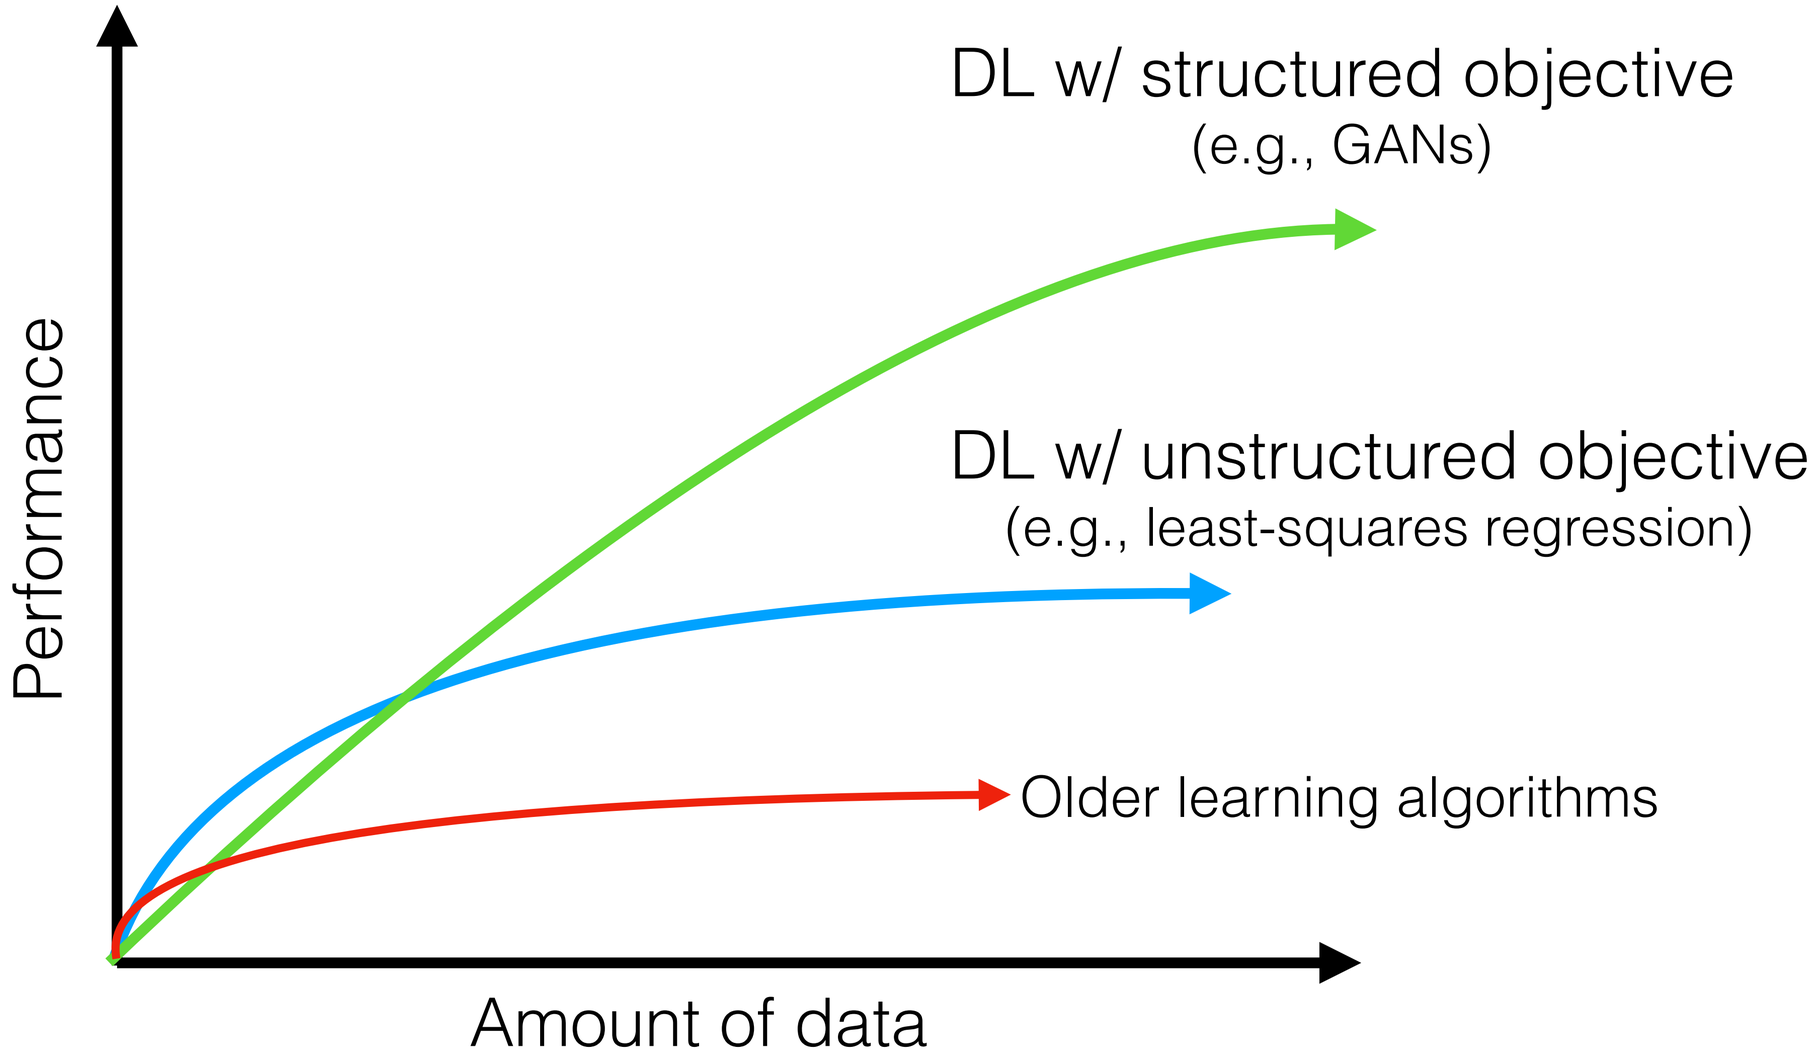

Input

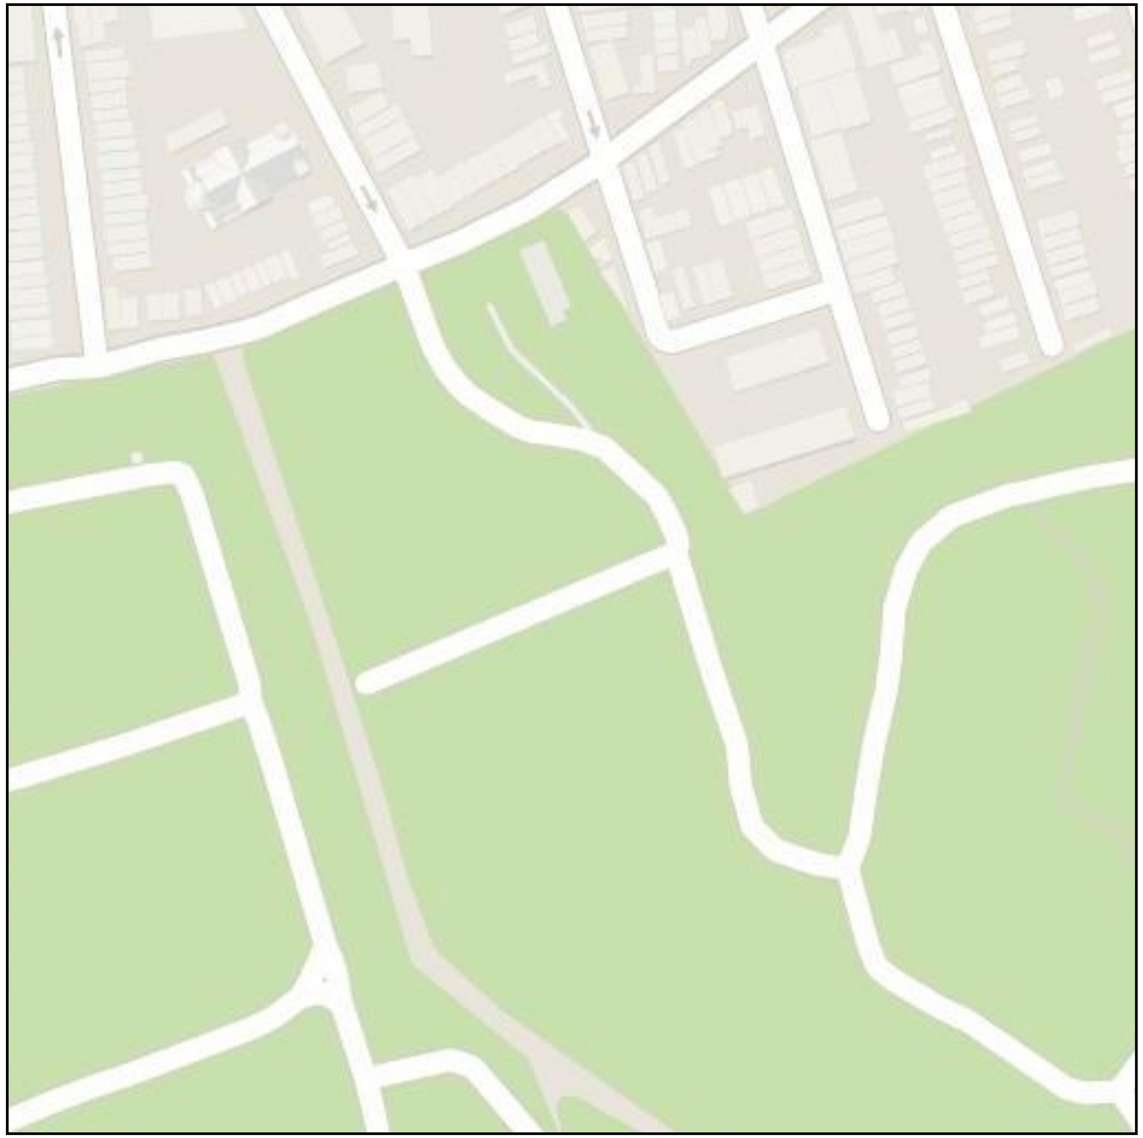

Unstructured prediction (L1)

Input

Structured Prediction (cGAN)

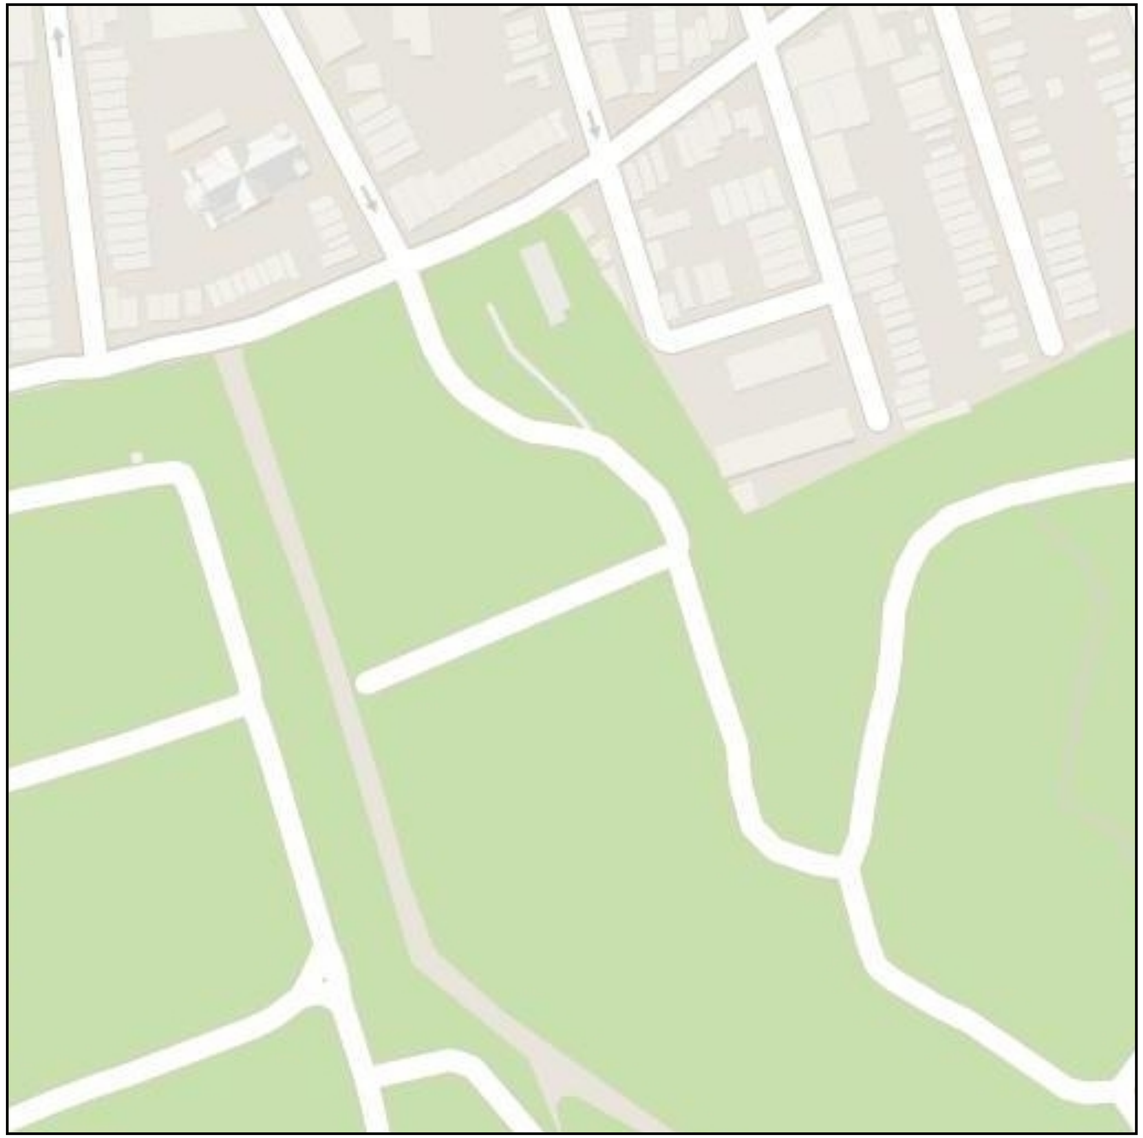

Training data

[HED, Xie & Tu, 2015]

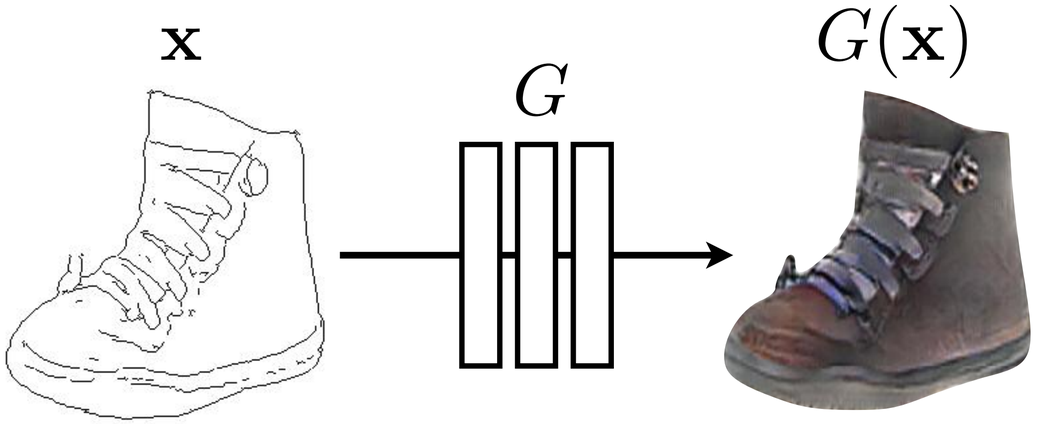

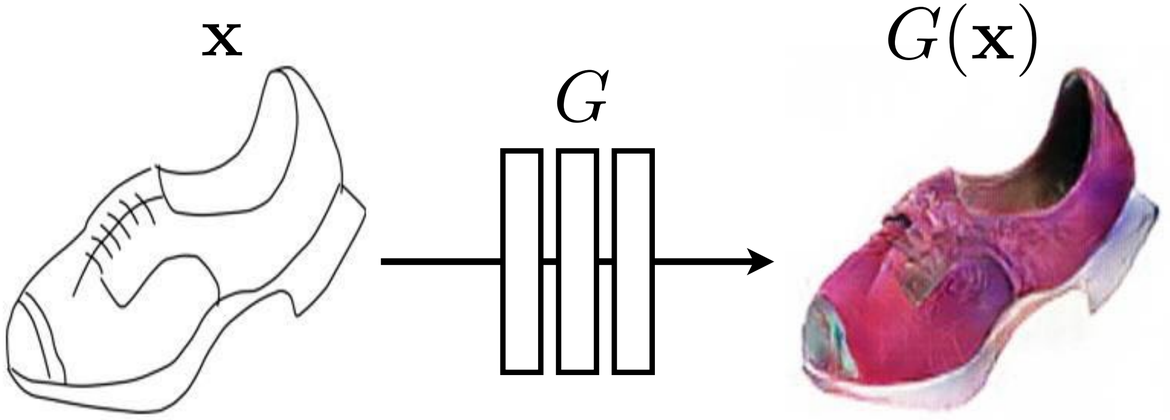

#edges2cats [Chris Hesse]

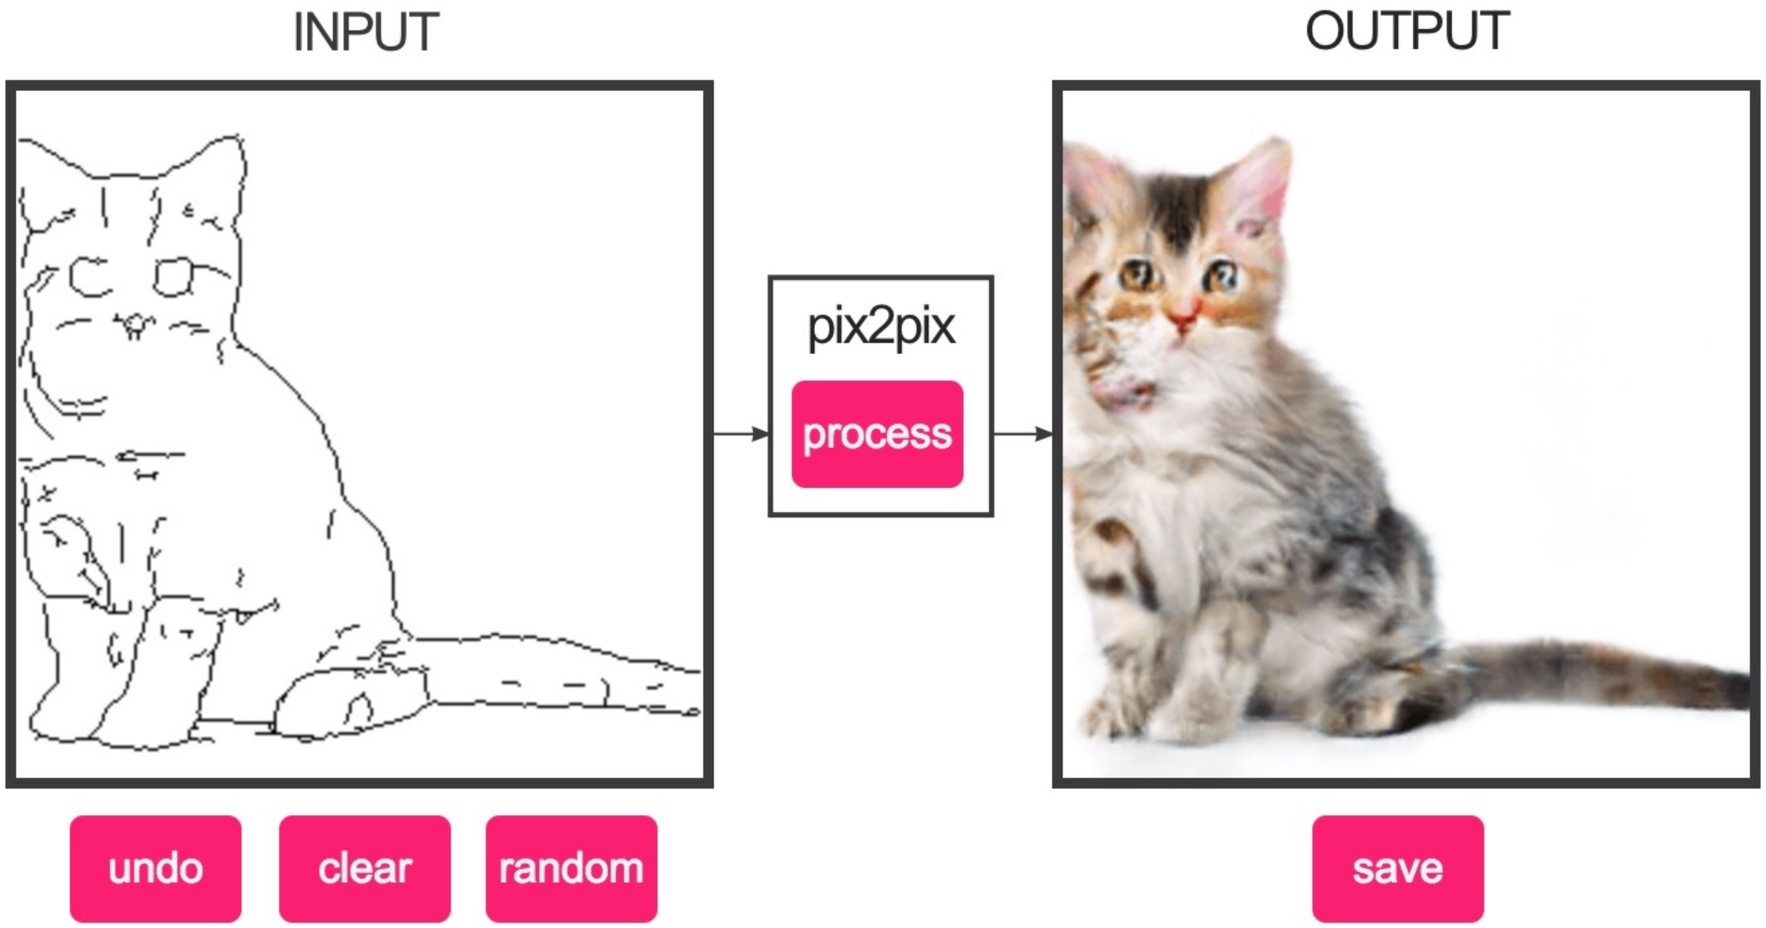

edges2cats

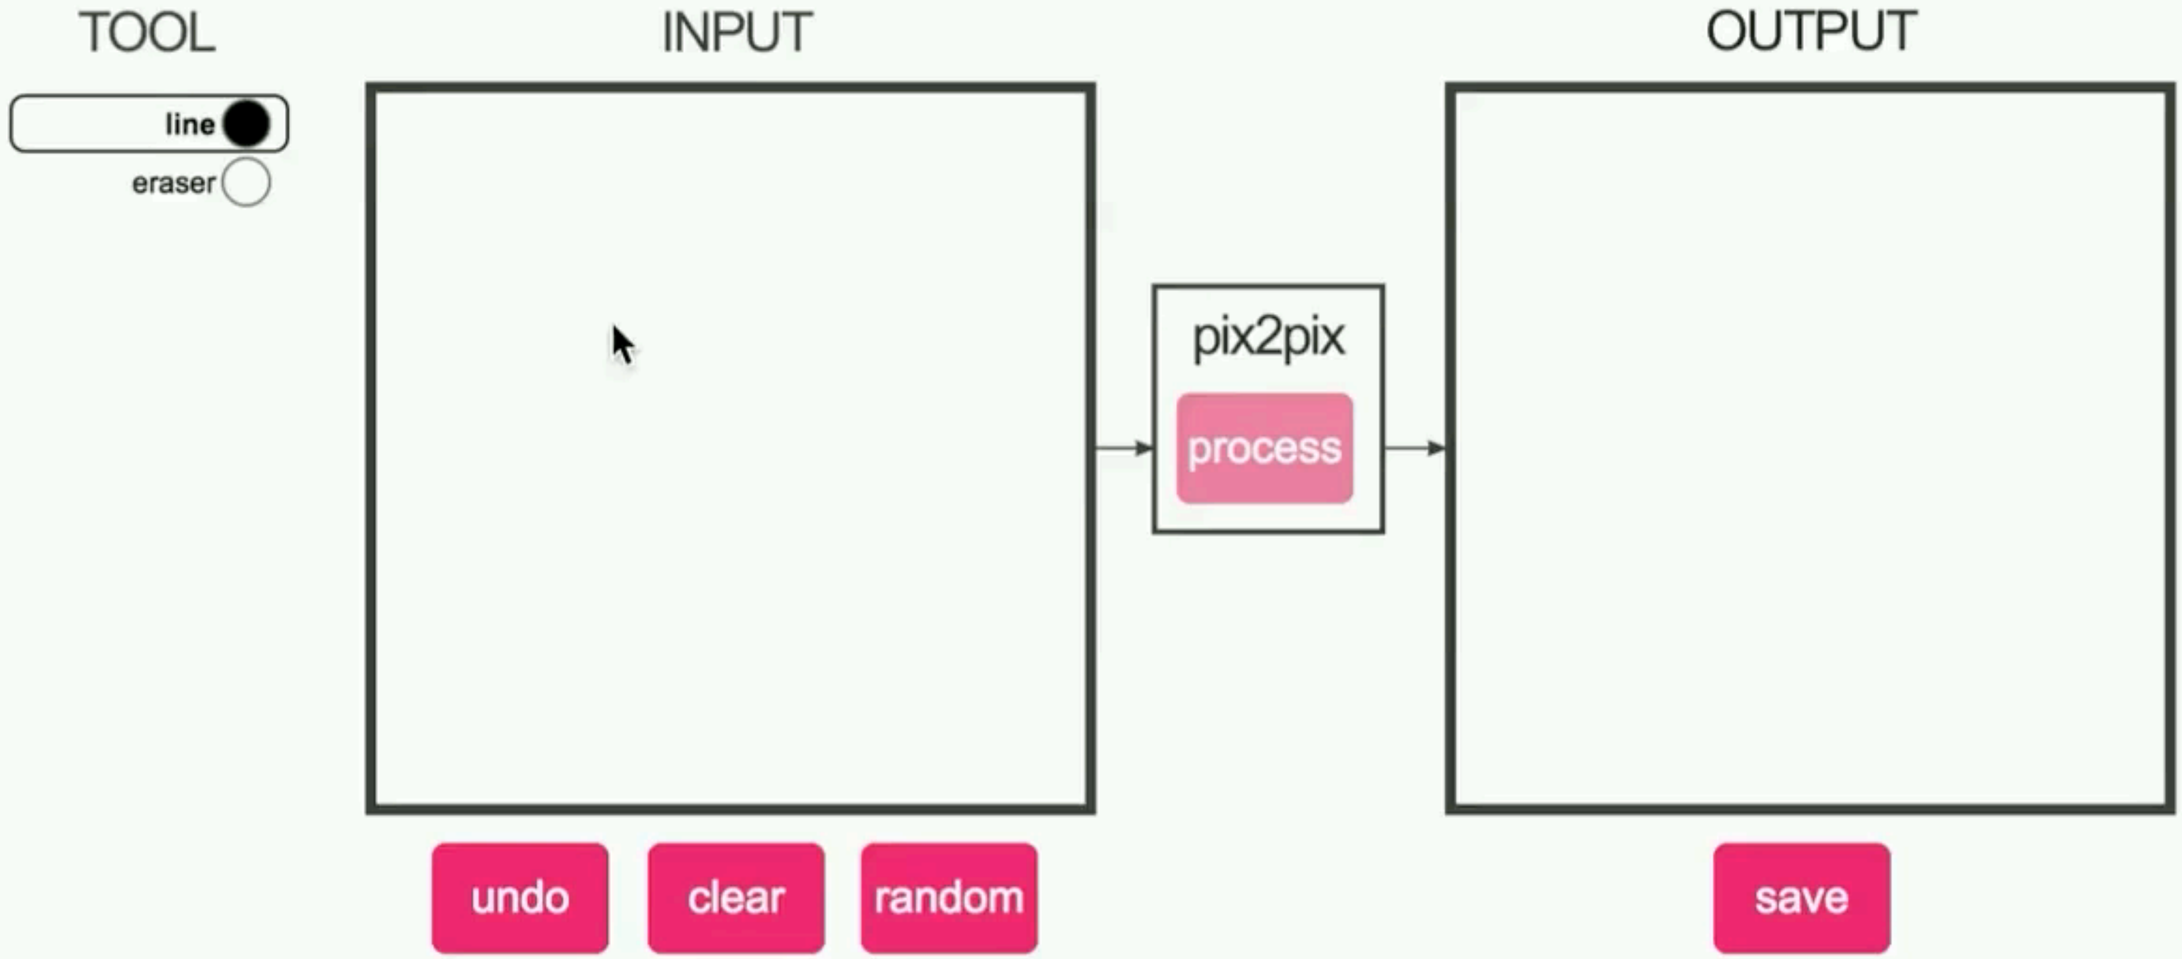
TOOL
line
eraser
INPUT
pix2pix
process
OUTPUT
undo
clear
random
save

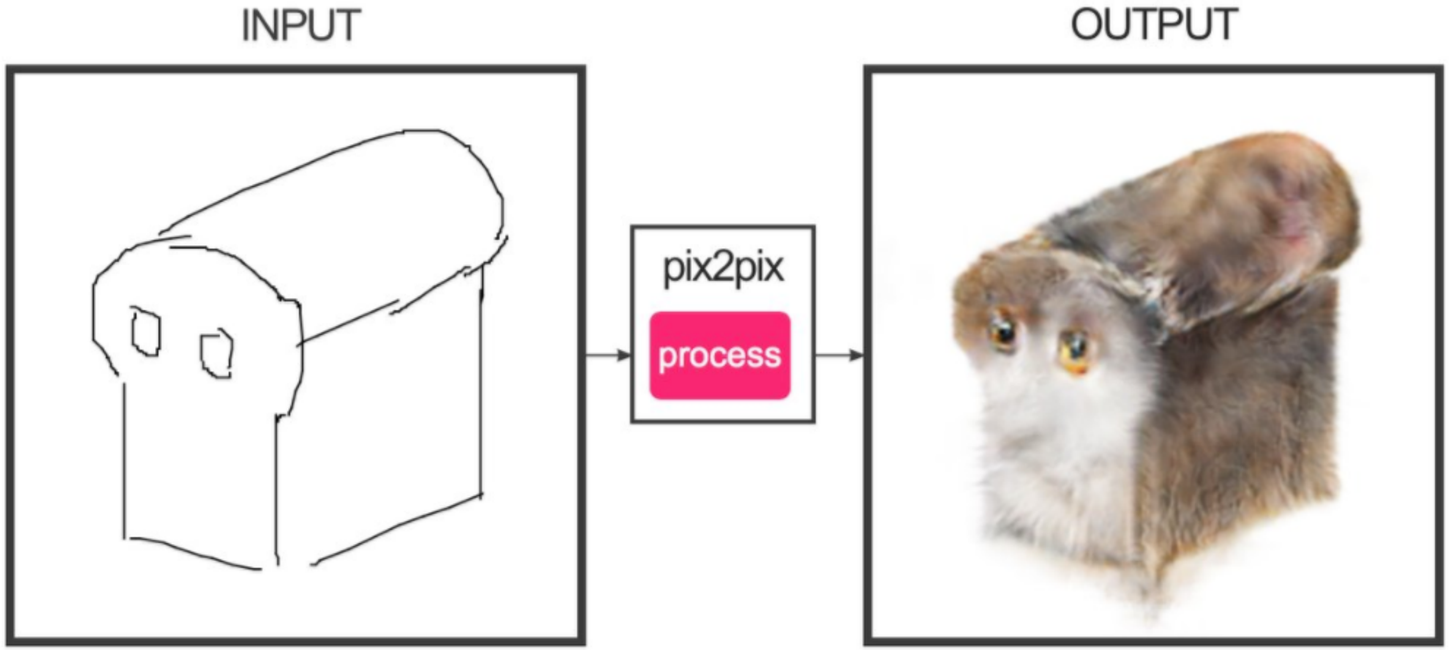

Ivy Tasi @ivymyt

Vitaly Vidmirov @vvid

1. Image synthesis
2. Structured prediction
3. **Domain mapping**

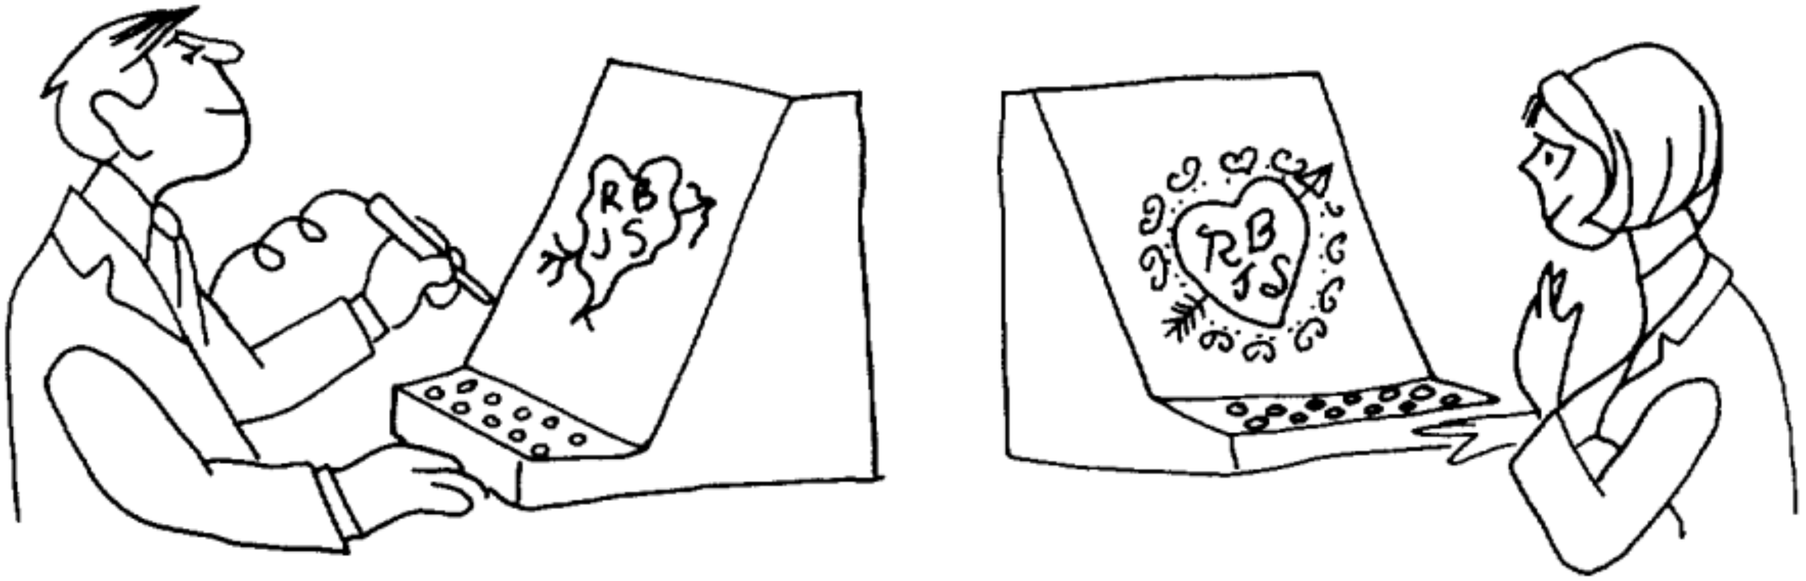

Domain mapping

[Includes slides from Jun-Yan Zhu, Taesung Park]

[Cartoon: The Computer as a Communication Device, Licklider & Taylor 1968]

Paired data

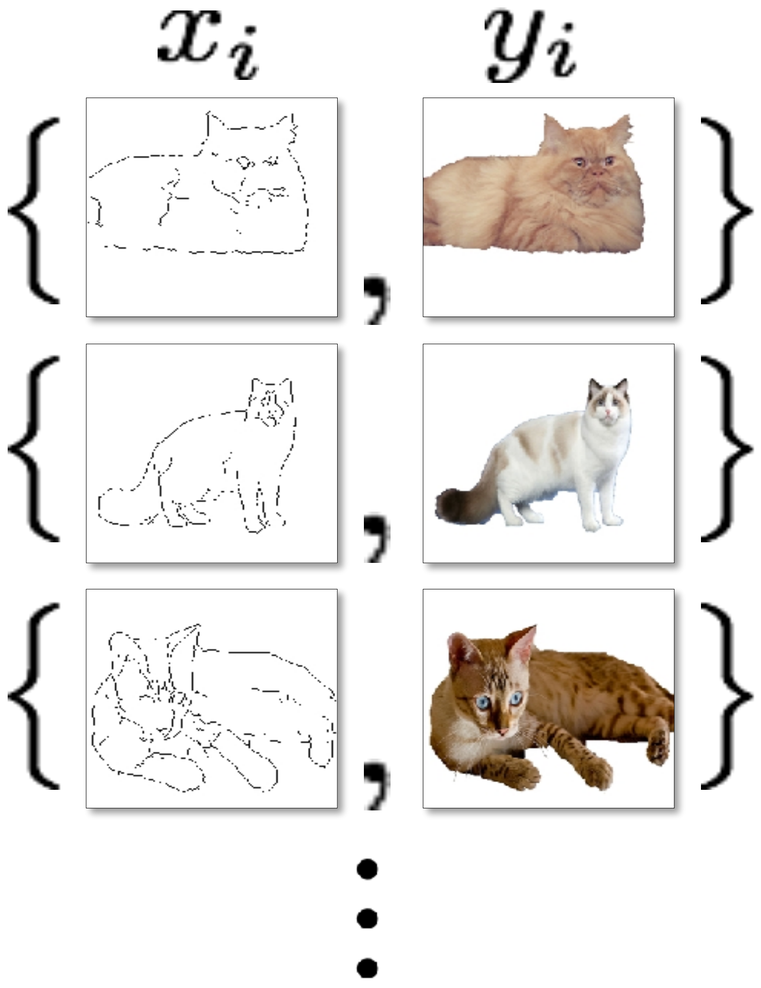

Unpaired data

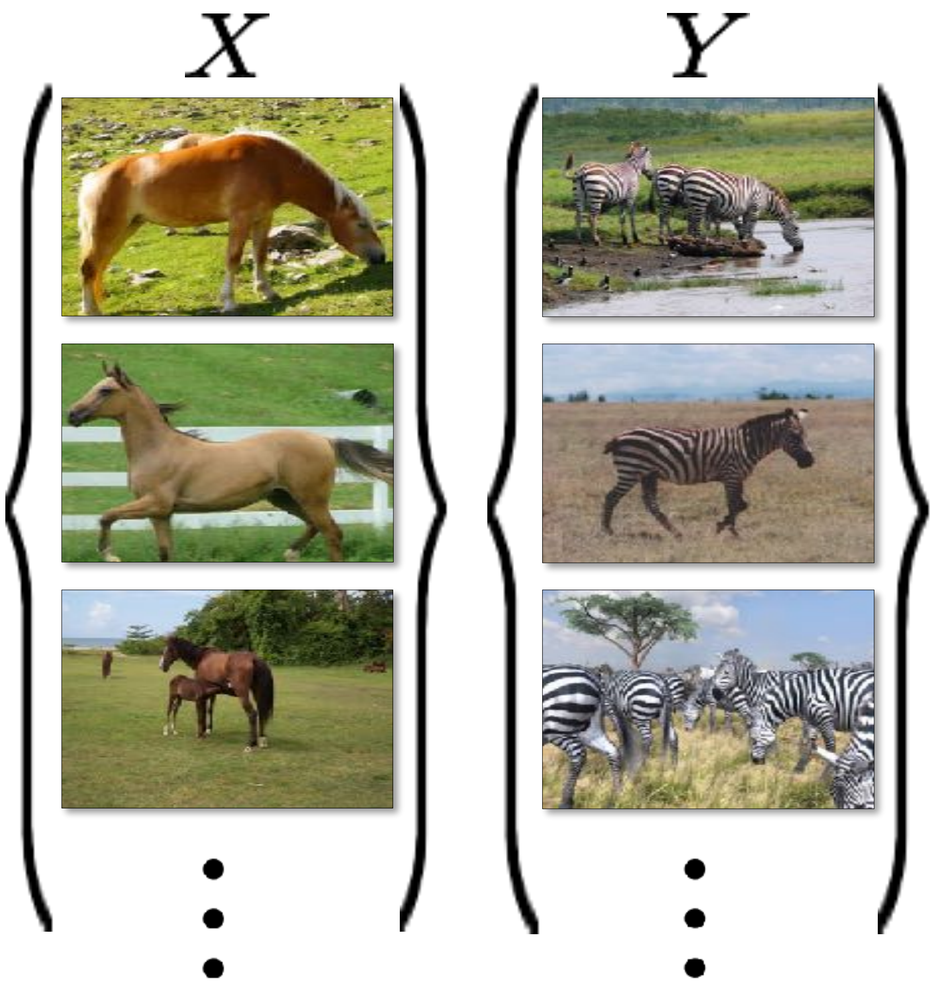

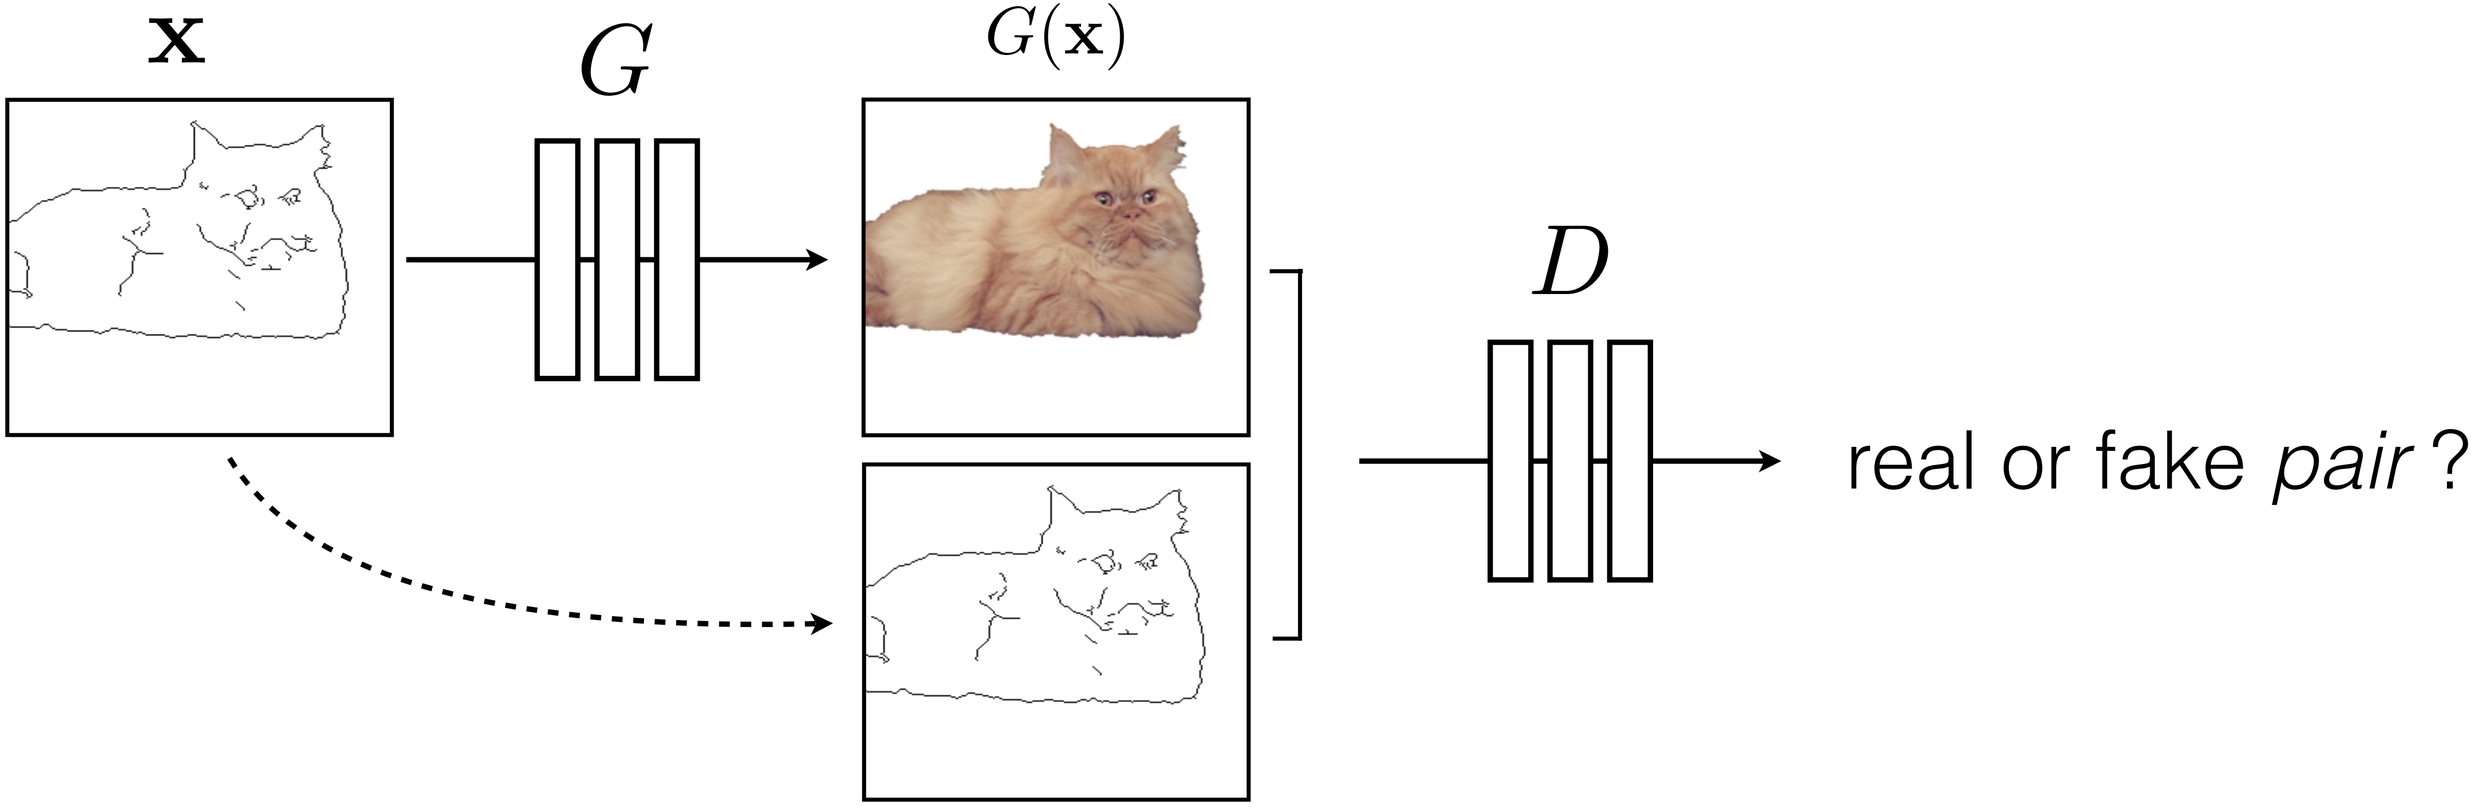

$$\arg\min_G \max_D \; \mathbb{E}_{\mathbf{x},\mathbf{y}}[\;\; \log D(\mathbf{x}, G(\mathbf{x})) + \log(1 - D(\mathbf{x}, \mathbf{y}))\;\;]$$

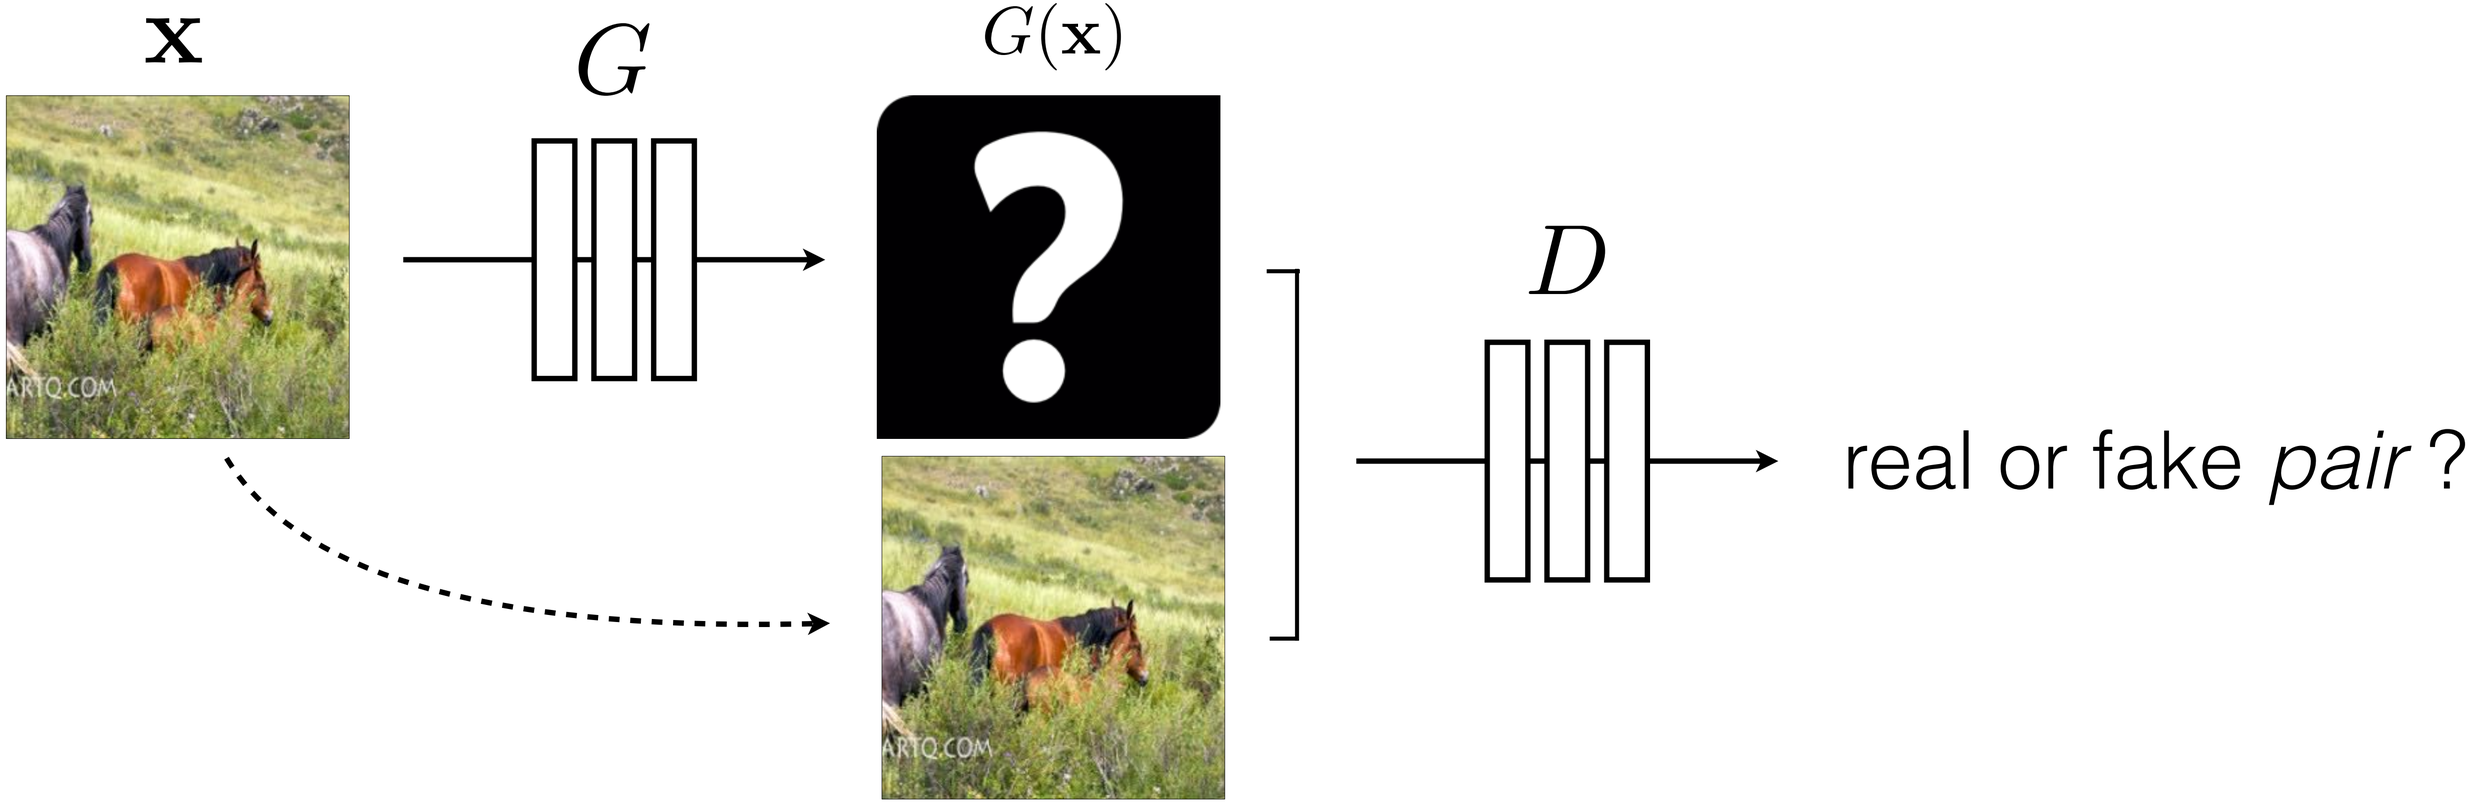

$$\arg\min_G \max_D \; \mathbb{E}_{\mathbf{x},\mathbf{y}}[\;\; \log D(\mathbf{x}, G(\mathbf{x})) + \log(1 - D(\mathbf{x}, \mathbf{y}))\;\;]$$

No input-output pairs!

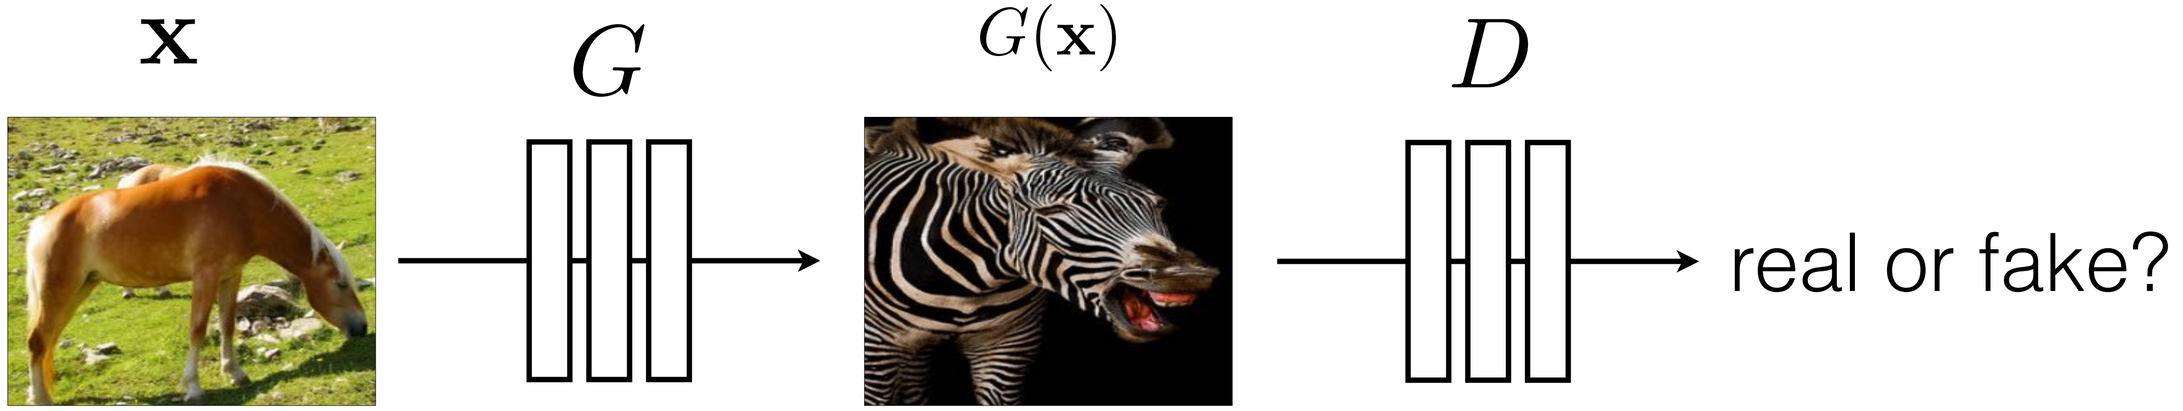

$$\arg\min_G \max_D \; \mathbb{E}_{\mathbf{x},\mathbf{y}}[\; \log D(G(\mathbf{x})) \quad + \quad \log(1 - D(\mathbf{y}))\;]$$

Usually loss functions check if output matches a target *instance*

GAN loss checks if output is part of an admissible *set*

$\mathbf{x}$

G

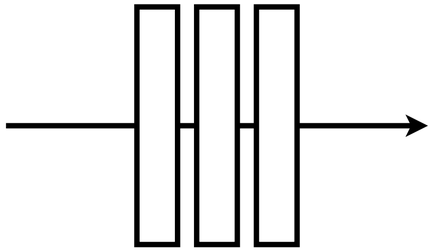

$G(\mathbf{x})$

D

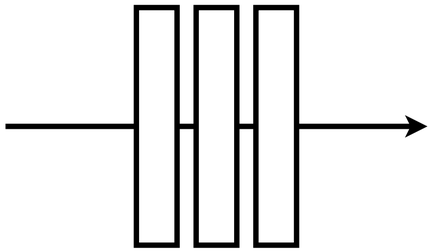

Real!

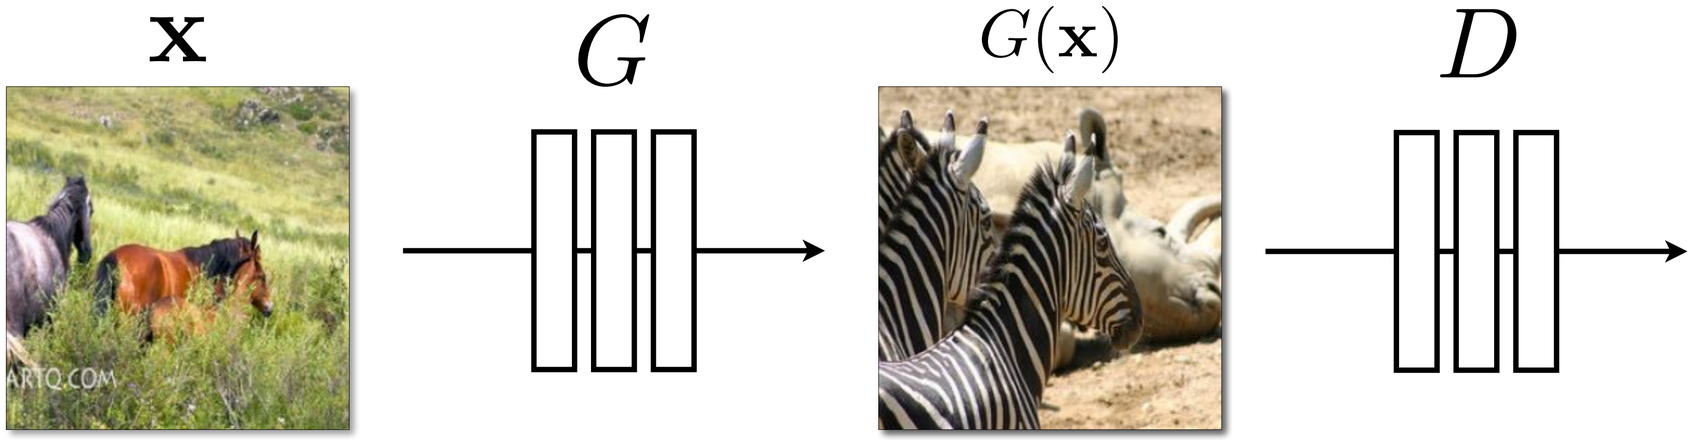

Nothing to force output to correspond to input

CycleGAN, or there and back aGAN

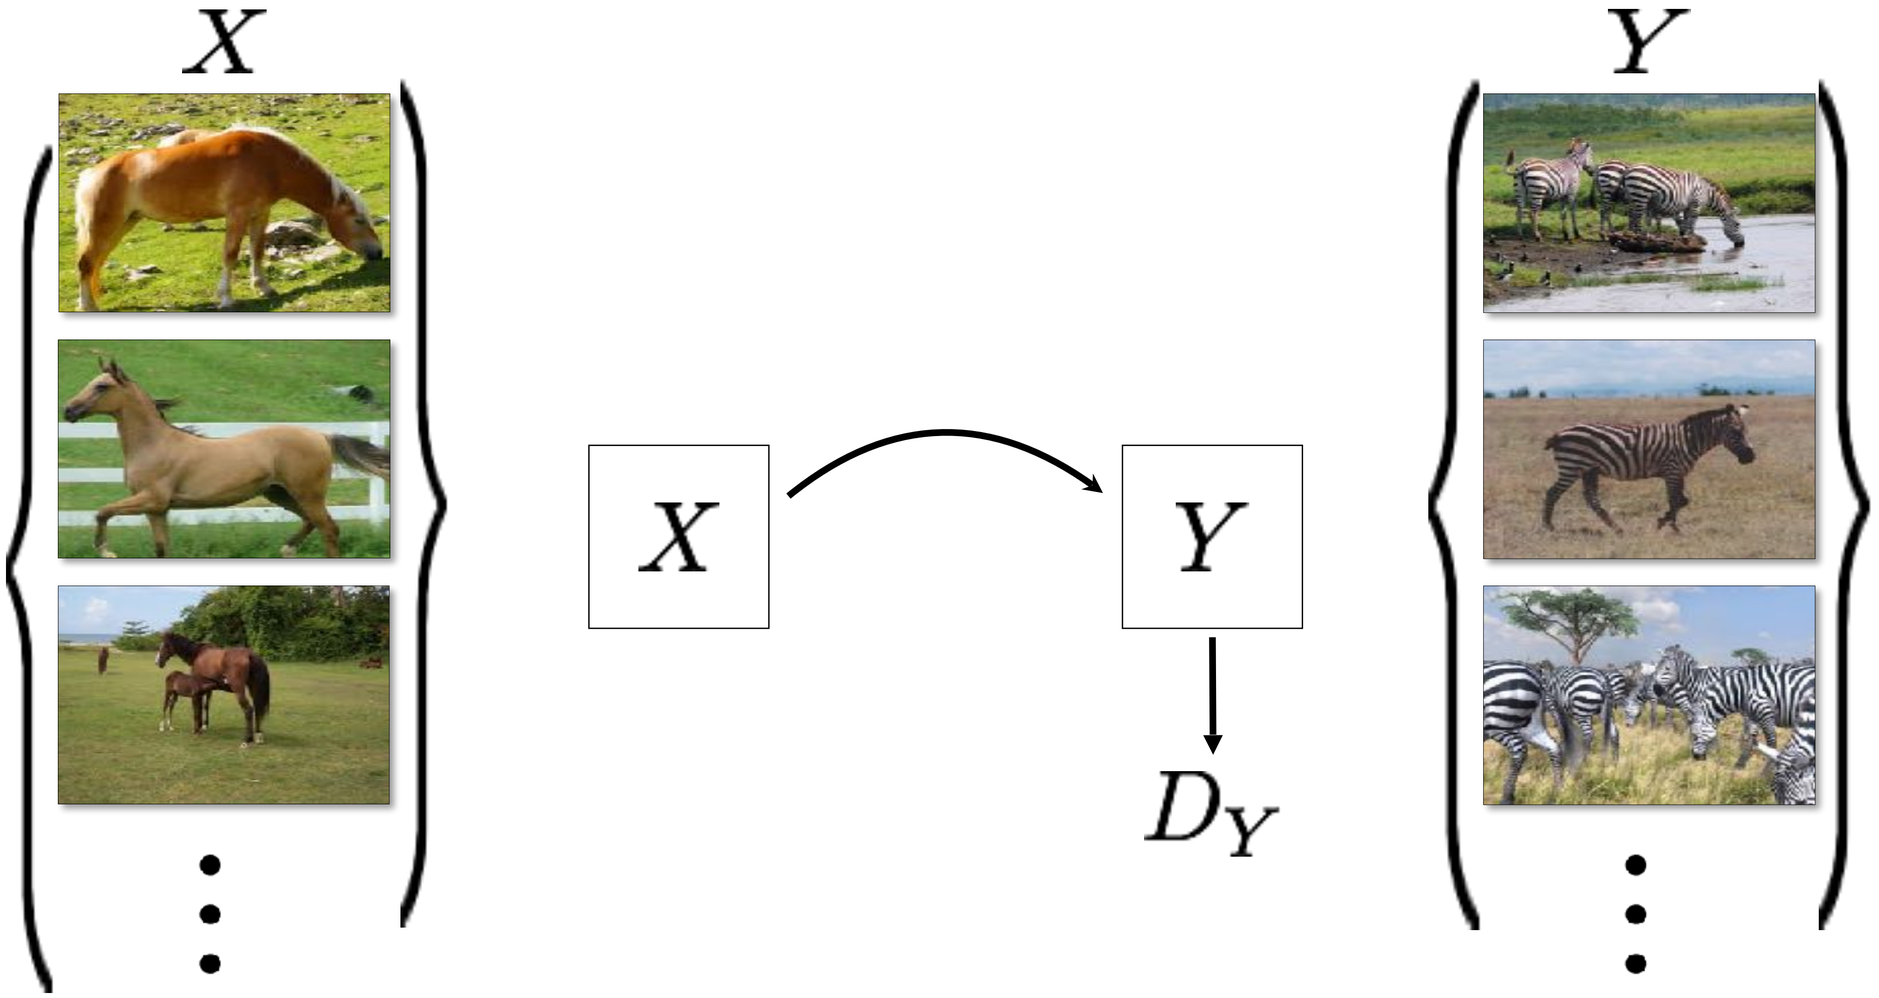

[Zhu*, Park* et al. 2017], [Yi et al. 2017], [Kim et al. 2017]

CycleGAN, or there and back aGAN

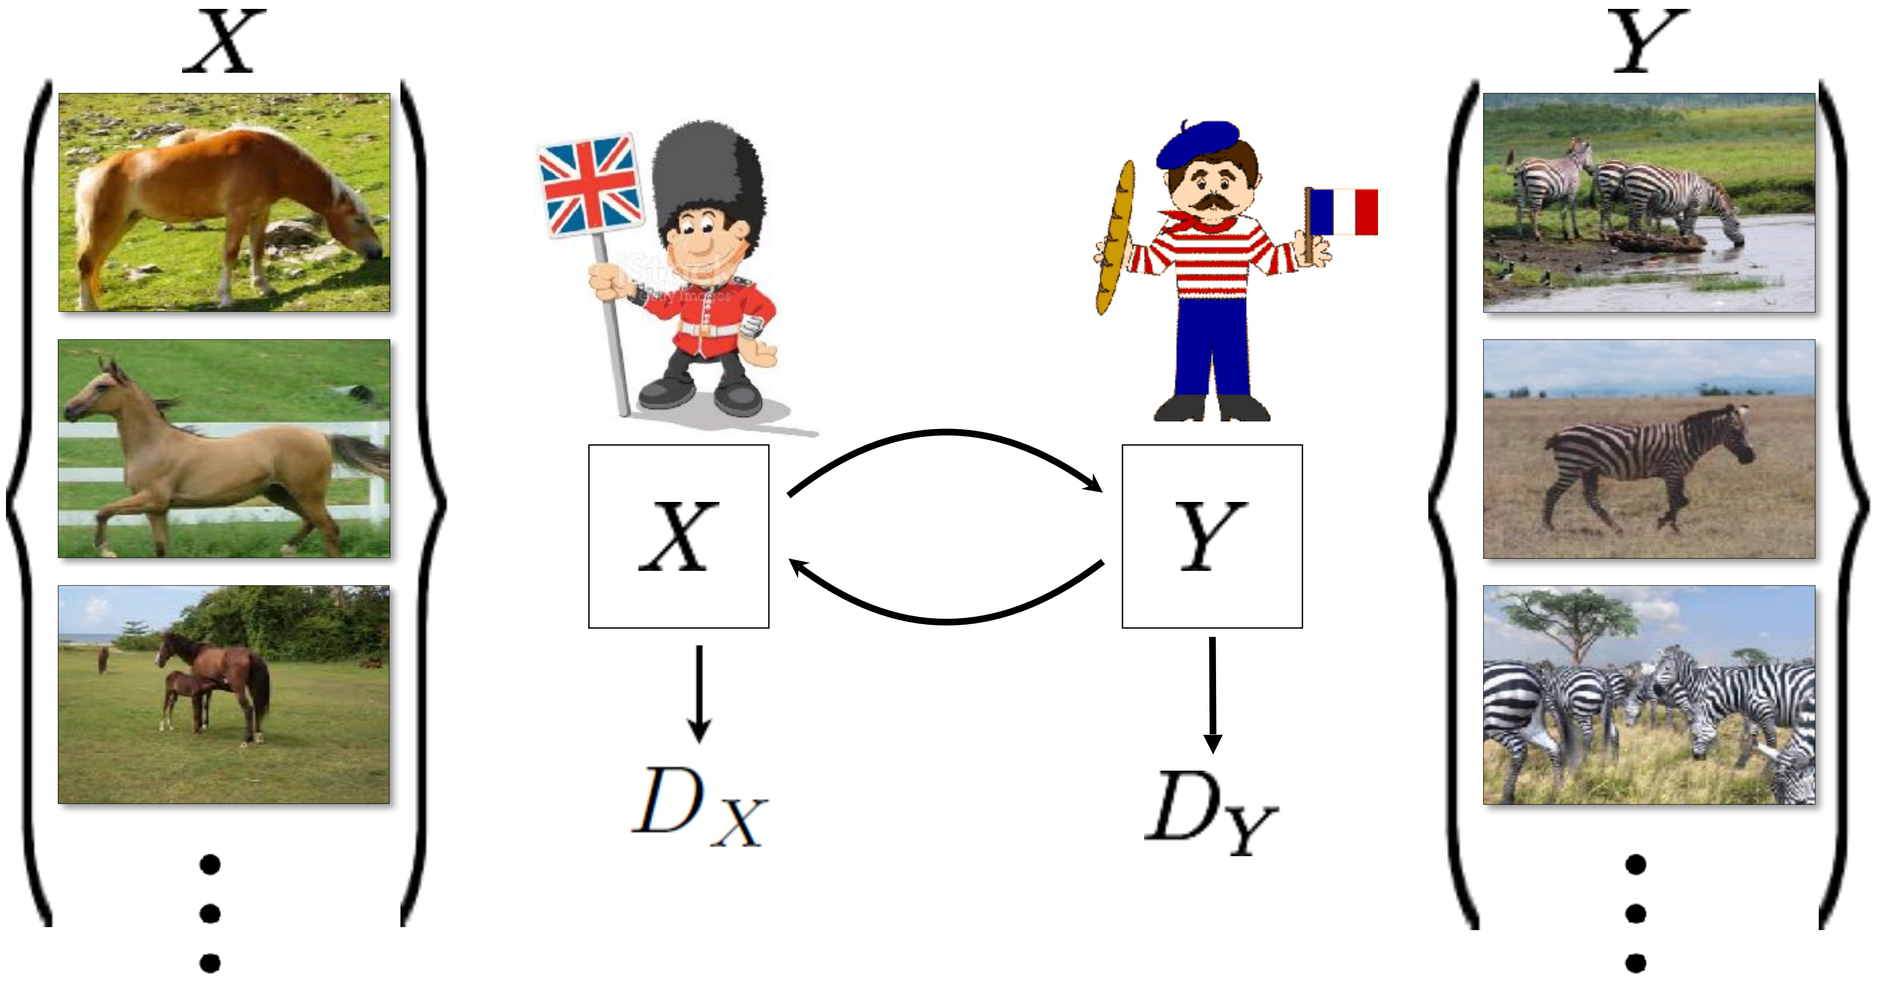

Cycle Consistency Loss

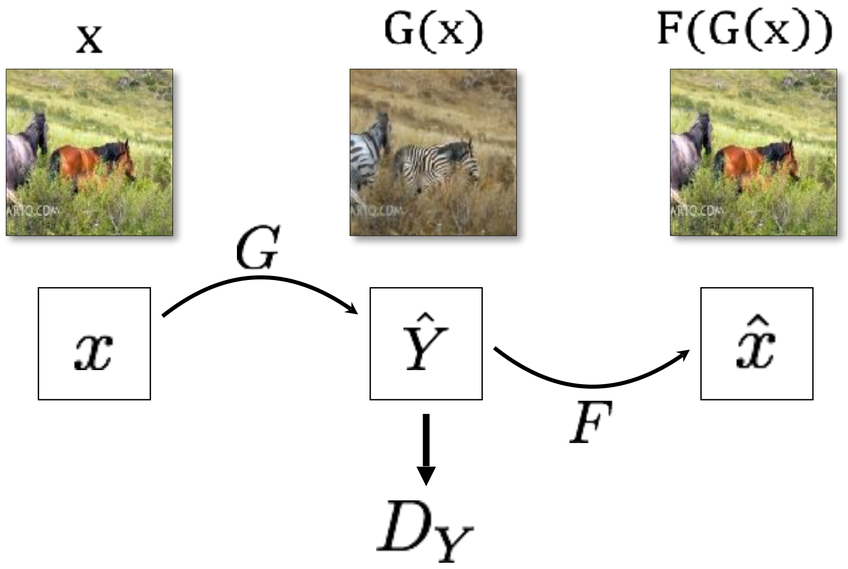

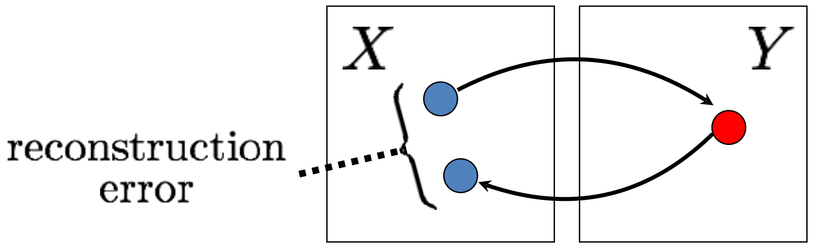

$$\|F(G(x)) - x\|_1$$

Cycle Consistency Loss

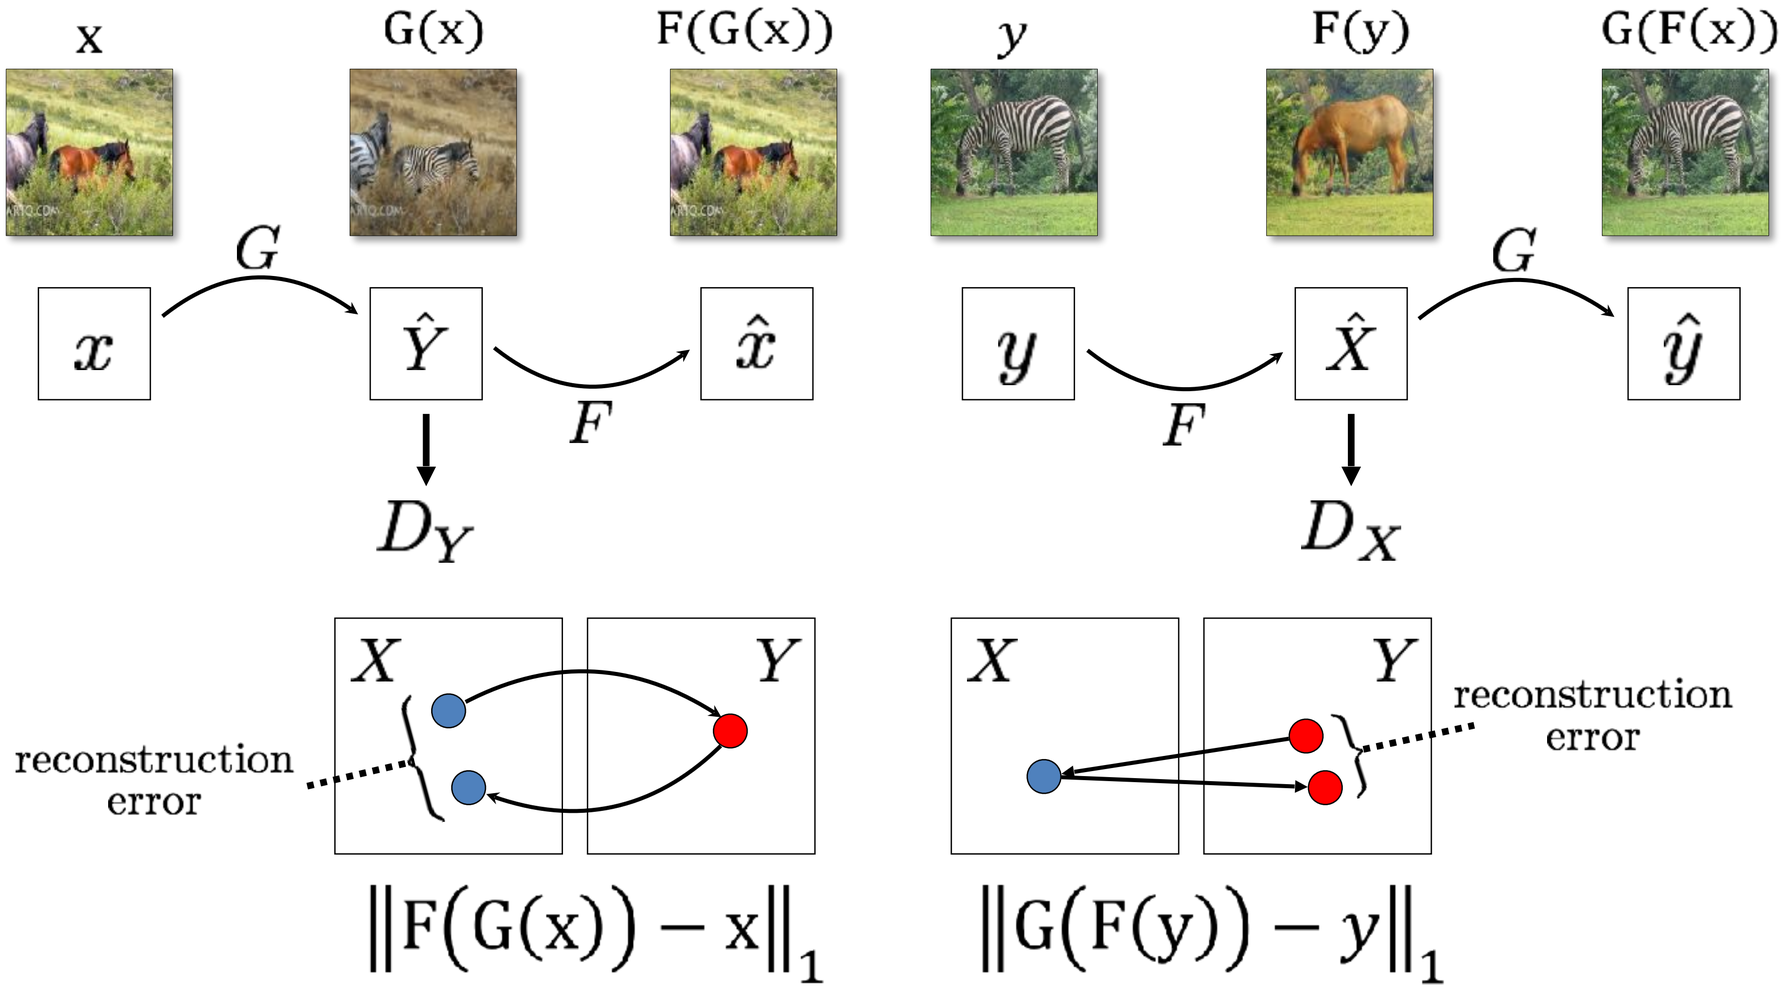

Paired translation

Training data

Objective

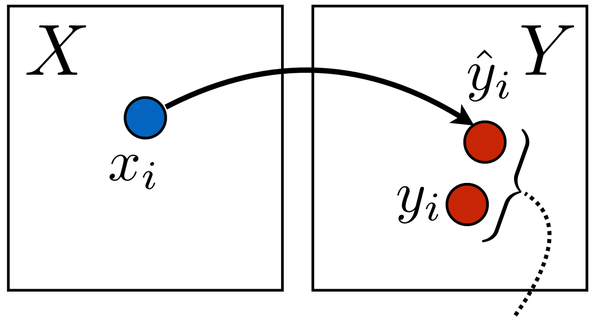

regression error

Input

Result

[“pix2pix”, Isola, Zhu, Zhou, Efros, 2017]

Unpaired translation

Training data

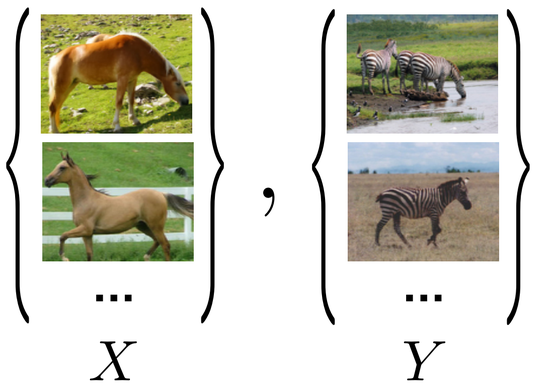

Objective

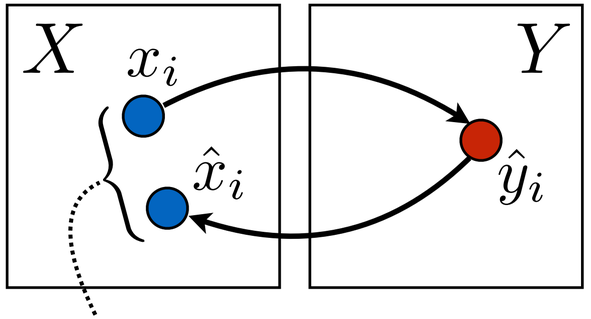

cycle-consistency error

Input

Result

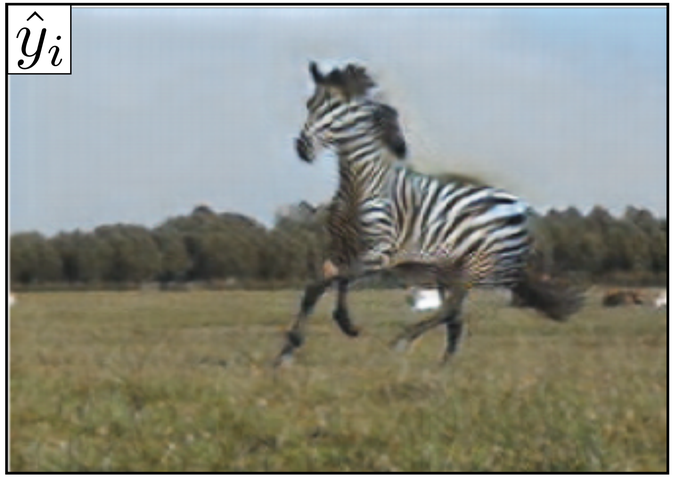

[“CycleGAN”, Zhu*, Park*, Isola, Efros, 2017]

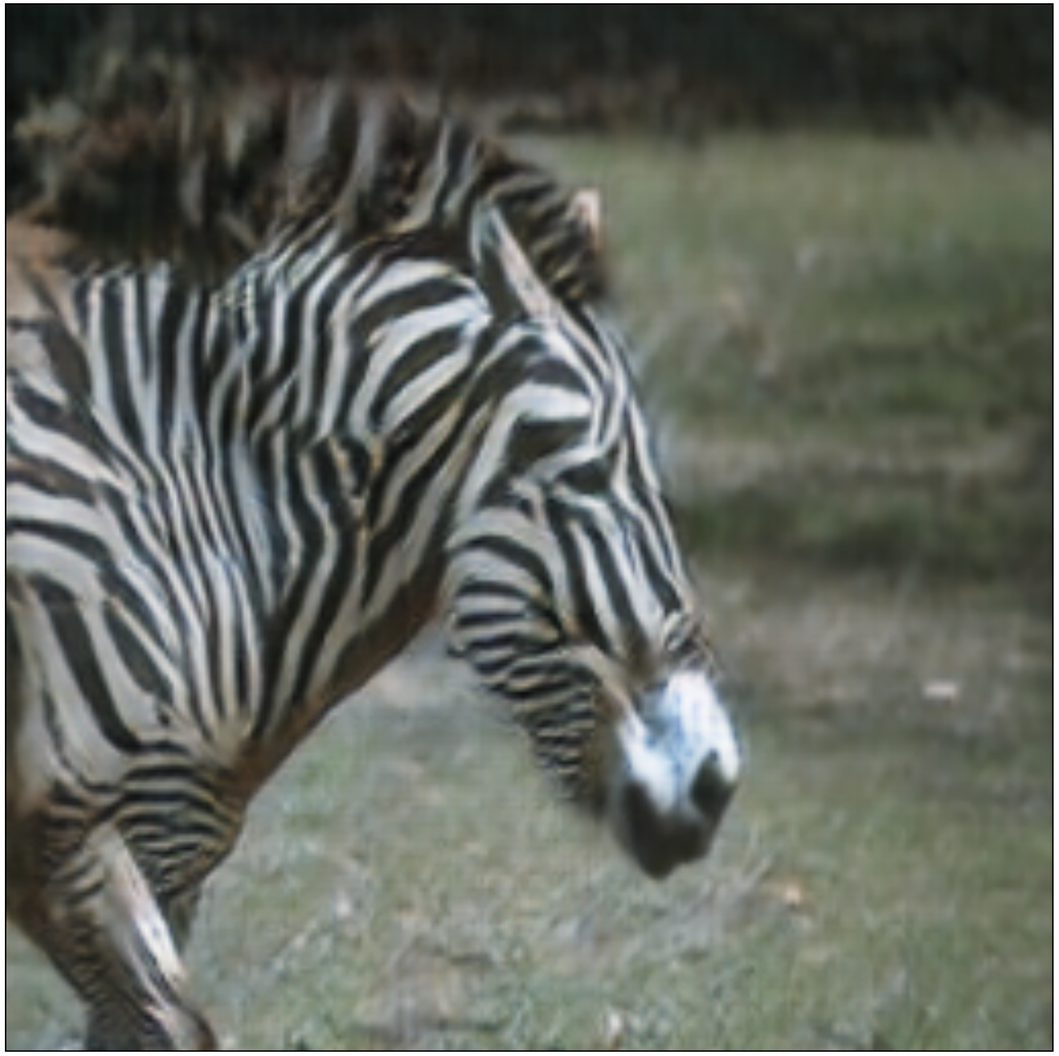


Input
Monet
Van Gogh
Cezanne
Ukiyo-e

GANs

Gaussian

Target distribution

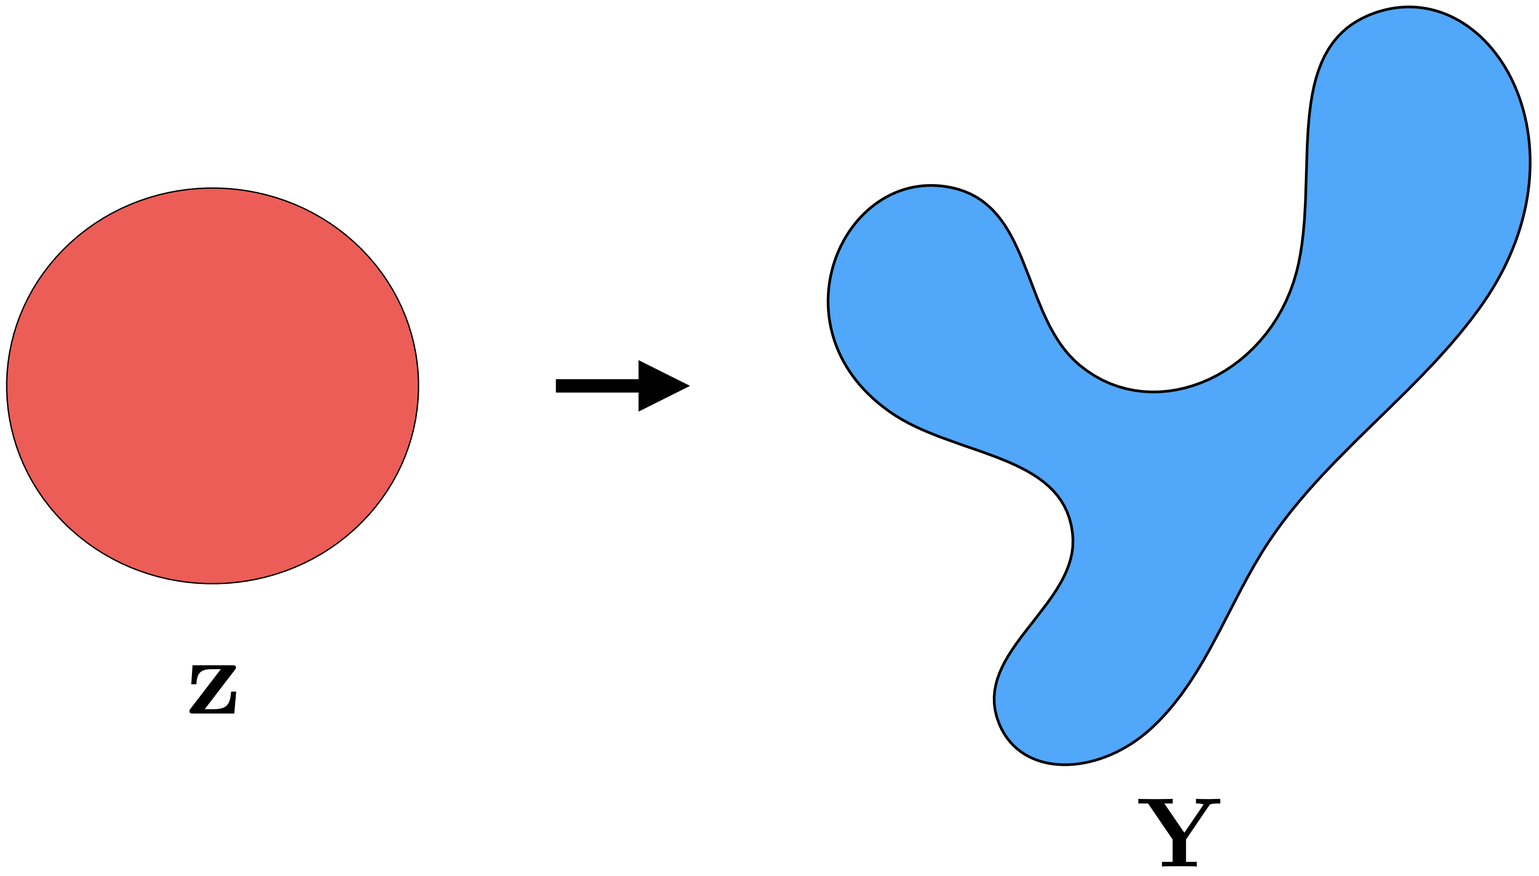

CycleGAN

Horses

Zebras

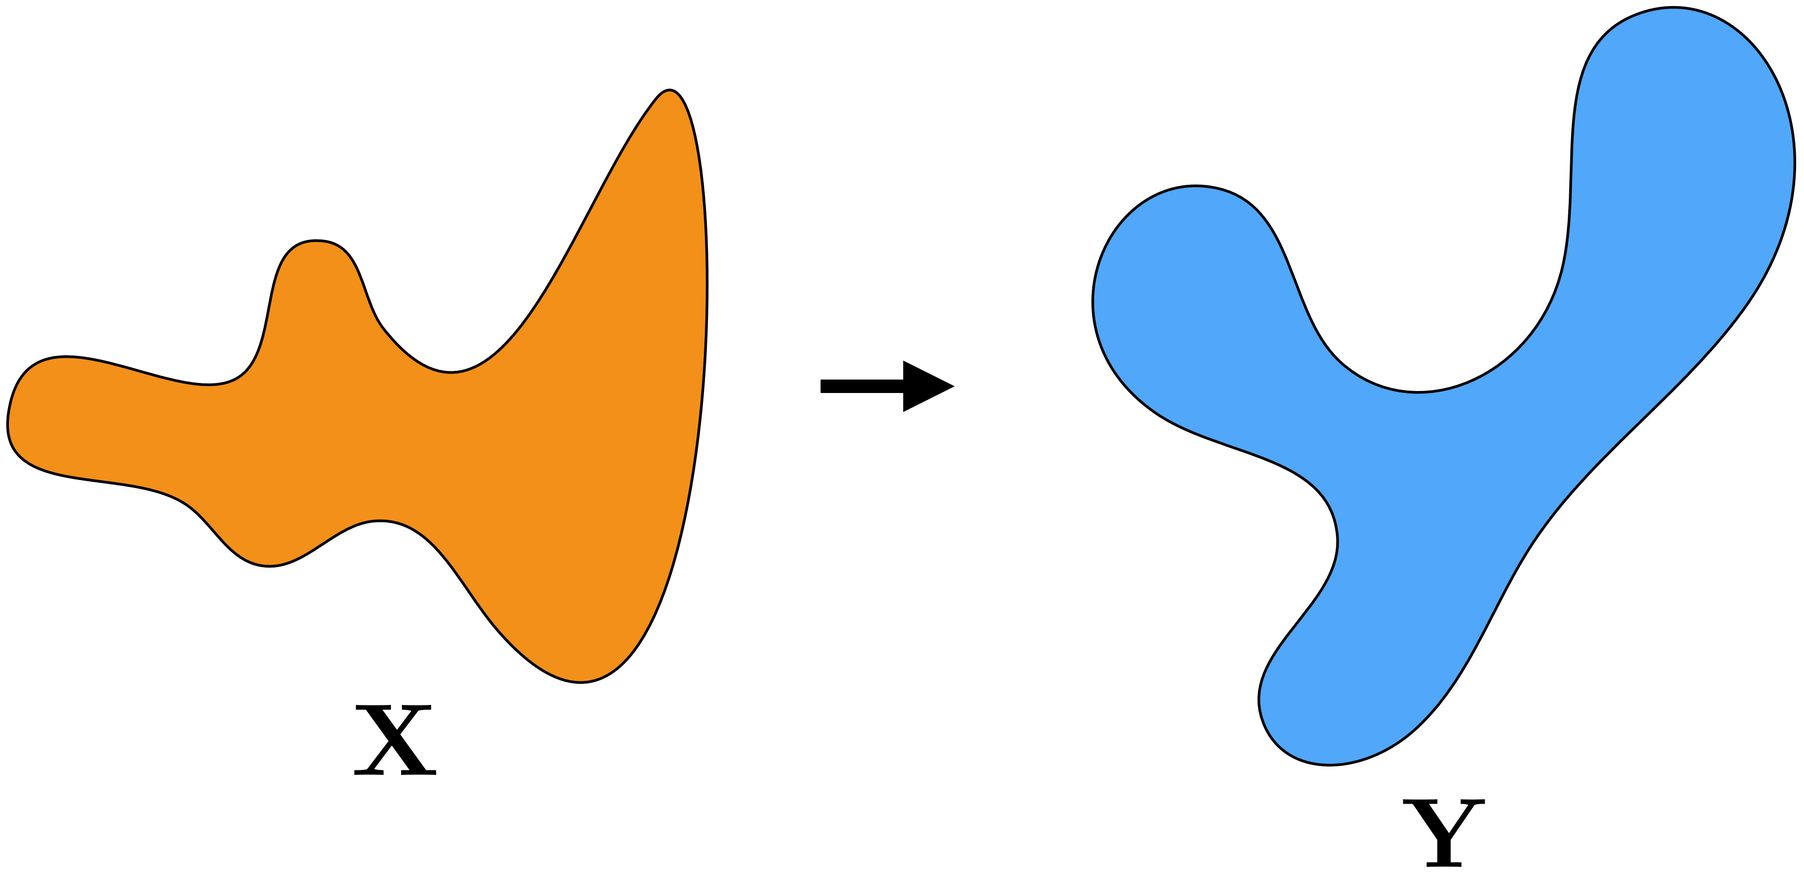

What would it look like if…?

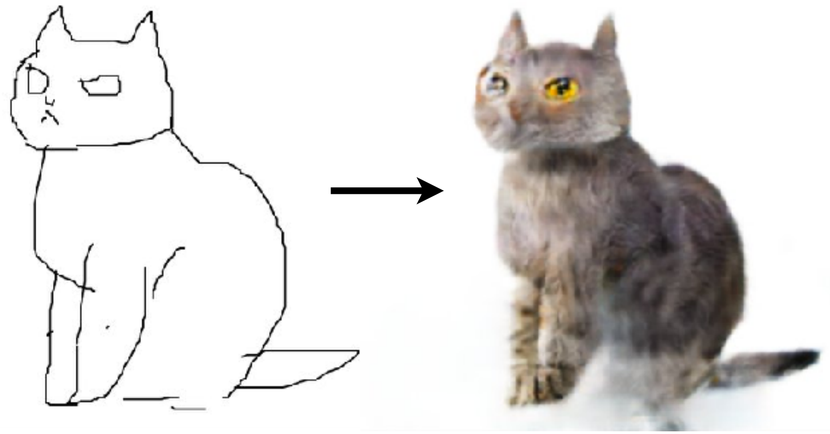

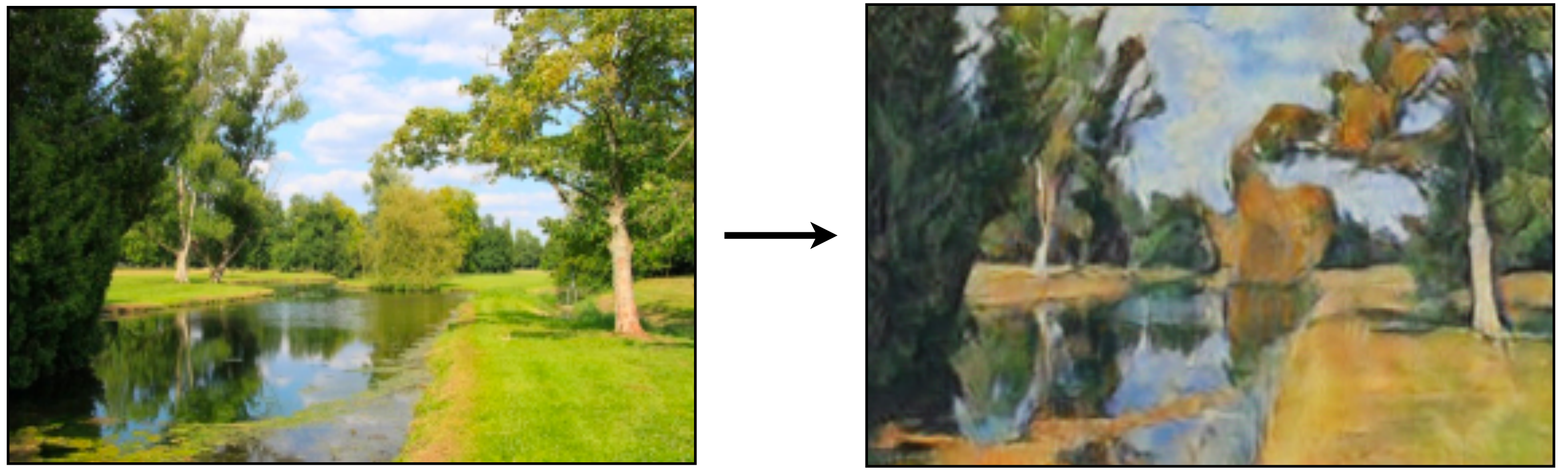

What would it look like if…?

MRI

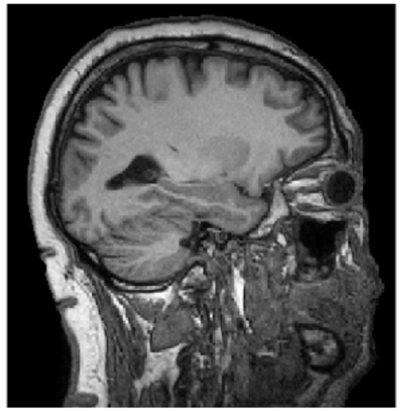

CT

[Wolterink et al, 2017]

Sim

“Real”

[Hoffman et al, 2018]