

Artificial Intelligence II (CS4442 & CS9542)

Classification: Logistic Regression

Boyu Wang
Department of Computer Science
University of Western Ontario

Recall: tumor example

- Thirty real-valued variables per tumor.
- Two variables that can be predicted:
 - Outcome (R=recurrence, N=non-recurrence)
 - Time (until recurrence, for R, time healthy, for N).

tumor size	texture	perimeter	...	outcome	time
18.02	27.6	117.5		N	31
17.99	10.38	122.8		N	61
20.29	14.34	135.1		R	27
...					

Recall: supervised learning

- ▶ A training example i has the form: (x_i, y_i) , where $x_i \in \mathbb{R}^n$ is the number of features (feature dimension). If $y \in \mathbb{R}$, this is the **regression** problem.
- ▶ If $y \in \{0, 1\}$, this is the **binary classification** problem.
- ▶ If $y \in \{1, \dots, C\}$ (i.e., y can take more than two values), this is the **multi-class classification** problem.
- ▶ Most binary classification algorithms can be extended to multi-class classification algorithms.

Linear model for classification

- ▶ As in linear regression, we consider a linear model h_w for classification:

$$h_w(x) = w^\top x$$

where we already augmented the data: $x \rightarrow [x; 1]$.

- ▶ Rules for binary classification: $h_w(x) \geq 0 \Rightarrow y = 1$;
 $h_w(x) < 0 \Rightarrow y = 0$
- ▶ How to choose w ?

Error (cost) function for classification

- Recall: for regression, we use the **sum-of-squared errors**:

$$J(w) = \frac{1}{2} \sum_{i=1}^m (h_w(x_i) - y_i)^2$$

Error (cost) function for classification

- Recall: for regression, we use the **sum-of-squared errors**:

$$J(w) = \frac{1}{2} \sum_{i=1}^m (h_w(x_i) - y_i)^2$$

- Can we use it for classification?

Error (cost) function for classification

- Recall: for regression, we use the **sum-of-squared errors**:

$$J(w) = \frac{1}{2} \sum_{i=1}^m (h_w(x_i) - y_i)^2$$

- Can we use it for classification?
- We could, but it is not the best choice
 - If $y = 1$, we want $h_w(x) > 0$ as much as possible, which reflects our confidence of classification
 - recall the connection between linear regression and maximum likelihood with Gaussian assumption

Probabilistic view for classification

- ▶ We would like a way of learning that is more suitable for the problem

Probabilistic view for classification

- ▶ We would like a way of learning that is more suitable for the problem
- ▶ Classification can be formulated as the following question: given a data point x , what is the probability $p(y|x)$?

Probabilistic view for classification

- ▶ We would like a way of learning that is more suitable for the problem
- ▶ Classification can be formulated as the following question: **given a data point x , what is the probability $p(y|x)$?**
- ▶ **Intuition:** we want find the conditional probability $p(y|x; w)$ parameterized by w in a way such that

$$h_w(x) = w^\top x \rightarrow +\infty \Rightarrow p(y = 1|x) \rightarrow 1$$

$$h_w(x) = w^\top x \rightarrow -\infty \Rightarrow p(y = 1|x) \rightarrow 0 \quad (\text{i.e., } p(y = 0|x) \rightarrow 1)$$

$$h_w(x) = w^\top x = 0 \Rightarrow p(y = 1|x) = 0.5$$

Sigmoid function

Consider the following function:

$$\sigma(a) \triangleq \frac{1}{1 + e^{-a}} = \frac{e^a}{1 + e^a}$$


- $a \rightarrow +\infty \Rightarrow \sigma(a) \rightarrow 1$
- $a \rightarrow -\infty \Rightarrow \sigma(a) \rightarrow 0$
- $a = 0 \Rightarrow \sigma(a) = 0.5$

Plug $h_w(x)$ into $\sigma(\cdot)$:


$$p(y = 1|x; w) \triangleq \sigma(h_w(x)) = \frac{1}{1 + e^{-w^\top x}},$$

which is exactly what we are looking for!

1D sigmoid function



(a)



(b)

Figure: The sigmoid function and the predicted probabilities.

Figure credit: Kevin Murphy

2D sigmoid function

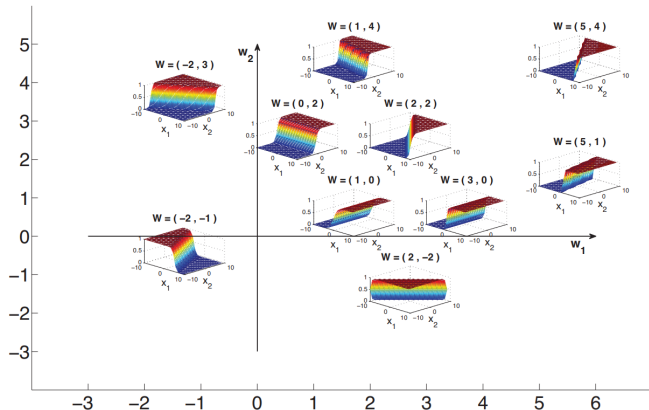


Figure: Plots of $\sigma(w_1x_1 + w_2x_2)$

Figure credit: Kevin Murphy

Error (cost) function for classification – revisited

- ▶ Maximum likelihood classification assumes that we will find the hypothesis that maximizes the (log) likelihood of the training data:

$$\arg \max_h \log p(\{x_i, y_i\}_{i=1}^m | h) = \arg \max_h \sum_{i=1}^m \log p(x_i, y_i | h)$$

(using the i.i.d. assumption)

- ▶ If we ignore the marginal distribution $p(x)$, we may maximize the **conditional probability** of the labels, given the inputs and the hypothesis h :

$$\arg \max_h \sum_{i=1}^m \log p(y_i | x_i; h)$$

The cross-entropy loss function for logistic regression

- ▶ Recall that for any data point (x_i, y_i) , the conditional probability can be represented by the sigmoid function:

$$p(y_i = 1 | x_i; h) = \sigma(h_w(x_i))$$

- ▶ Then the log-likelihood of a hypothesis h_w is

$$\begin{aligned}\log L(w) &= \sum_{i=1}^m \log p(y_i | x_i; h_w) = \sum_{i=1}^m \begin{cases} \log \sigma(h_w(x_i)), & \text{if } y_i = 1 \\ \log(1 - \sigma(h_w(x_i))), & \text{if } y_i = 0 \end{cases} \\ &= \sum_{i=1}^m (y_i \log t_i + (1 - y_i) \log(1 - t_i)),\end{aligned}$$

where $t_i = \sigma(h_w(x_i))$.

- ▶ The **cross-entropy loss function** is the negative of $\log L(w)$.

Linear regression vs. logistic regression

Both use linear model: $h_w(x) = w^\top x$

► Conditional probability:

- linear regression: $p(y_i|x_i; w) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2}\left(\frac{y_i - w^\top x_i}{\sigma}\right)^2}$
- logistic regression: $p(y_i|x_i; w) = \begin{cases} \frac{1}{1+e^{-w^\top x_i}}, & \text{if } y_i = 1 \\ 1 - \frac{1}{1+e^{-w^\top x_i}}, & \text{if } y_i = 0 \end{cases}$

► Log-likelihood function:

- linear regression: $\log L(w) = \sum_{i=1}^m \log \left(\frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2}\left(\frac{y_i - w^\top x_i}{\sigma}\right)^2} \right)$
- logistic regression:
 $\log L(w) = \sum_{i=1}^m (y_i \log t_i + (1 - y_i) \log(1 - t_i))$, where
 $t_i = \sigma(h_w(x_i))$.

► Solution:

- linear regression: $w = (x^\top X)^{-1} X^\top y$
- logistic regression: No analytical solution


Optimization procedure: gradient descent

Objective: **minimize** a loss function $J(w)$


Gradient descent for function minimization

- 1: **Input:** number of iterations: N , **learning rate:** α
- 2: Initialize $w_{(0)}$
- 3: **for** $n = 1$ to N **do**
- 4: Compute the gradient: $g_{(n)} = \nabla J(w_{(n)})$
- 5: $w_{(n+1)} = w_{(n)} - \alpha g_{(n)}$
- 6: **if** converges (e.g, $|w_{(n+1)} - w_{(n)}| \leq \epsilon$) **then**
- 7: Stop
- 8: **end if**
- 9: **end for**

Effect of the learning rate



(a)



(b)

Figure: Gradient descent on a simple function, starting from $(0, 0)$, for 20 steps using a fixed learning rate α . (a) $\alpha = 0.1$. (b) $\alpha = 0.6$

Figure credit: Kevin Murphy

Optimization procedure: Newton's method

Objective: **minimize** a loss function $J(w)$

Newton's method for function minimization

- 1: **Input:** number of iterations: N
- 2: Initialize $w_{(0)}$
- 3: **for** $n = 1$ to N **do**
- 4: Compute the gradient: $g_{(n)} = \nabla J(w_{(n)})$
- 5: Compute the **Hessian**: $H_{(n)} = \nabla^2 J(w_{(n)})$
- 6: $w_{(n+1)} = w_{(n)} - H_{(n)}^{-1} g_{(n)}$
- 7: **if** converges (e.g, $|w_{(n+1)} - w_{(n)}| \leq \epsilon$) **then**
- 8: Stop
- 9: **end if**
- 10: **end for**

Gradient descent for logistic regression

Objective: **minimize** the cross-entropy loss function ($-\log L(w)$):

$$J(w) \triangleq -\log L(w) = -\sum_{i=1}^m (y_i \log t_i + (1 - y_i) \log(1 - t_i))$$

$$\nabla J(w) = \sum_{i=1}^m (t_i - y_i) x_i, \quad t_i = \sigma(h_w(x_i)) = \frac{1}{1 + e^{-w^\top x_i}}$$

Gradient descent for logistic regression

- 1: **Input:** number of iterations: N , learning rate: α
- 2: Initialize $w_{(0)}$
- 3: **for** $n = 1$ to N **do**
- 4: Compute the gradient: $g_{(n)} = \nabla J(w_{(n)}) = \sum_{i=1}^m (t_i - y_i) x_i$
- 5: $w_{(n+1)} = w_{(n)} - \alpha g_{(n)}$
- 6: **if** converges (e.g., $|w_{(n+1)} - w_{(n)}| \leq \epsilon$) **then**
- 7: Stop
- 8: **end if**
- 9: **end for**

Newton's method for logistic regression

Objective: **minimize** the cross-entropy loss function ($-\log L(w)$):

$$J(w) \triangleq -\log L(w) = -\sum_{i=1}^m (y_i \log t_i + (1 - y_i) \log(1 - t_i))$$

$$\nabla J(w) = \sum_{i=1}^m (t_i - y_i) x_i, \quad H(w) = \nabla^2 J(w) = \sum_{i=1}^m t_i(1 - t_i) x_i x_i^\top = X^\top S X$$

where $S \in \mathbb{R}^{m \times m}$ is a diagonal matrix with elements $S_{ii} = t_i(1 - t_i)$,
 $t_i = \sigma(h_w(x_i)) = \frac{1}{1 + e^{-w^\top x_i}}$.

Then, the update rule becomes:

$$w_{(n+1)} = w_{(n)} - H(w_{(n)})^{-1} \nabla J(w_{(n)})$$

Gradient descent vs. Newton's method

- ▶ Newton's method usually requires significantly fewer iterations than gradient descent
- ▶ Computing the Hessian and its inverse is expensive
- ▶ Approximation algorithms exist which help to compute the product of the inverse Hessian with gradient without explicitly computing H

Cross-entropy loss vs. squared loss




Figure: Decision boundaries obtained by minimizing squared loss (magenta line) and cross-entropy loss (green line)

Figure credit: Christopher Bishop

Regularized logistic regression

One can do regularization for logistic regression just like in the case of linear regression.

- ▶ ℓ_2 -regularized logistic regression

$$J(w) \triangleq -\log L(w) + \frac{\lambda}{2} \|w\|_2^2$$

- 2. ℓ_1 -regularized logistic regression

$$J(w) \triangleq -\log L(w) + \lambda \|w\|_1$$

Multi-class logistic regression

- For 2 classes:

$$p(y = 1|x; w) = \frac{1}{1 + e^{-w^\top x}} = \frac{e^{w^\top x}}{1 + e^{w^\top x}}$$

- For C classes $\{1, \dots, C\}$:

$$p(y = c|x; w_1, \dots, w_C) = \frac{e^{w_c^\top x}}{\sum_{c=1}^C e^{w_c^\top x}}$$

– called the **softmax** function

Multi-class logistic regression

Gradient descent for multi-class logistic regression

```
1: Input: number of iterations:  $N$ , learning rate:  $\alpha$ 
2: Initialize  $C$  vectors:  $w_{1,(0)}, \dots, w_{C,(0)}$ 
3: for  $n = 1$  to  $N$  do
4:   for  $c = 1$  to  $C$  do
5:     Compute the gradient with respect to  $w_{c,(n)}$ :  $g_{c,(n)}$ 
6:      $w_{c,(n+1)} = w_{c,(n)} - \alpha g_{c,(n)}$ 
7:   end for
8:   if converges (e.g,  $|w_{c,(n+1)} - w_{c,(n)}| \leq \epsilon$  for all classes) then
9:     Stop
10:  end if
11: end for
```

Multi-class logistic regression

- ▶ After training, the probability of $y = c$ is given by

$$p(y = c|x; w_1, \dots, w_C) \triangleq \sigma_c(x) = \frac{e^{w_c^\top x}}{\sum_{c=1}^C e^{w_c^\top x}}$$

- ▶ Predict class label as the most probable label:

$$y = \arg \max_c \sigma_c(x)$$

Summary

- ▶ Logistic regression for classification
- ▶ Cross-entropy loss for logistic regression
- ▶ No analytical solution that minimizes the cross-entropy loss/maximizes the log-likelihood
- ▶ Use gradient descent to find a solution
- ▶ Multi-class logistic regression